

본 수업자료는 2025년도 과학기술 정보통신부 및 정보통신기획평가원의 'sw중심대학사업' 지원을 받아 제작 되었습니다.

ComputerVision

Week2

2025-2

Mobile Systems Engineering
Dankook University

Learning Objectives

- In this lecture, you will learn
 - (1) What LeNet-5 is and why it is historically important
 - (2) How a CNN processes images step-by-step.
 - (3) The roles of
 - Convolutional Layers: learn patterns like edges and textures.
 - Subsampling (Pooling) Layers: reduce the size of feature maps.
 - Fully Connected Layers: make the final decision.
 - (4) How to calculate the number of trainable parameters in each layer.
 - (5) Why LeNet-5 is considered the foundation of modern deep learning models.


LeNet-5: A Landmark in CNN History

■ LeNet-5

- Developed by Yann LeCun's team in 1998
- One of the **first convolutional neural networks**
- Designed to **recognize handwritten digits** (e.g., MNIST)
- **Introduced in the paper**
"Gradient-Based Learning Applied to Document Recognition"
- Paved the way for modern deep learning models



LeNet and MNIST handwritten digit recognition



LeNet-5: A Landmark in CNN History

■ Historical Importance of LeNet-5

- **1. Real-world Application**


- LeNet-5 was originally developed for digit recognition tasks in real documents such as **postal codes** and **bank checks**, addressing real industrial needs in the 1990s.

- **2. Automatic Feature Extraction**

- It replaced manually engineered features with a **trainable, end-to-end pipeline**, enabling the model to learn **edge, shape, and texture representations** directly from raw pixel data.

- **3. Efficient Design Philosophy**

- LeNet-5 used **shared weights**, **local receptive fields**, and **subsampling (average pooling)** to significantly **reduce the number of parameters**, allowing it to generalize well with limited computational resources and data.



LeNet-5: A Landmark in CNN History

■ Historical Importance of LeNet-5


• 4. Designing for Generalization

- LeNet-5 demonstrated the importance of keeping model complexity low when training data is limited.

- $E_{test} - E_{train} = k \left(\frac{h}{P}\right)^\alpha$; k, α : constants (empirically determined)

✓ This equation explains how the **generalization gap** ($E_{test} - E_{train}$) grows with model complexity h and shrinks with more data P .

✓ LeNet-5 was carefully designed to keep h small by using **shared weights**, **local connections**, and **subsampling**.




Architecture of LeNet-5

■ Architecture Overview

- **LeNet-5 is made of the following layers**

- **Input layer:** Receives a 32×32 grayscale image (1 channel)
- **Three convolutional layers:** Extract patterns like **edges and textures** → **C1, C3, C5**
- **Two subsampling layers** (average pooling): Reduce the size of feature maps to make the model faster and simpler → **S2, S4**
- **One fully connected layer:** Combines all features and makes a decision → **F6**
- **Output layer:** Gives the final prediction using **10 class scores** (digits 0 to 9)



Architecture of LeNet-5

■ C1 – First Convolution Layer

- Input: 32×32 grayscale image
- 6 filters of size 5×5
→ produces 6 feature maps of size 28×28 (i.e., 6@ 28×28)

- No padding, stride = 1

- Activation: tanh

- Number of Parameters in C1


- $(5 \times 5 \times 1 + 1) \times 6 = 156$ parameters

- ✓ The size of each filter is 5×5 .
- ✓ Input has 1 channel
- ✓ Plus 1 bias, which is also a trainable parameter
- ✓ 6 filters (i.e., 6 output channels)

INPUT
 32×32

C1: feature maps
6@ 28×28

S2:
6@1





Architecture of LeNet-5

■ Convolution in Deep Learning

• What does a Convolution Layer do in CNNs?

- In deep learning, a **convolution layer** is used to
 - ✓ Extract **local features** from an input image
 - ✓ Slide a small filter (kernel) across the input
 - ✓ At each location, **multiply** and **sum** the overlapping values
 - ✓ Create a **feature map** that highlights patterns (e.g., edges, textures)

- “Unlike traditional signal processing,
the filters are **learned** from data during training.”



Example of convolution layer with 3×3 filter and 1 of stride ⁸

Architecture of LeNet-5

■ Convolution in Deep Learning

• How to Compute Output (Size)?

- Let's define

- ✓ W_{in} , H_{in} : Input width and height

- ✓ K : Kernel (filter) size

- ✓ P : Padding

- ✓ S : Stride

- Then the output size W_{out} , H_{out} is

$$\checkmark W_{out} = \left\lceil \frac{W_{in}-K+2P}{S} \right\rceil + 1$$

$$\checkmark H_{out} = \left\lceil \frac{H_{in}-K+2P}{S} \right\rceil + 1$$

Input image

9	4	1	2	2
1	1	1	0	4
1	2	1	0	6
1	0	0	2	4
9	6	7	4	1

Filter

0	2	1
4	1	0
1	0	1

Output array

$$\begin{aligned} \text{Output } [0][0] &= (9*0) + (4*2) + (1*4) \\ &+ (1*1) + (1*0) + (1*1) + (2*0) + (1*1) \\ &= 0 + 8 + 1 + 4 + 1 + 0 + 1 + 0 + 1 \\ &= 16 \end{aligned}$$



I see a filter and an output...


But what happens if we use **padding** or change the **stride**?

Architecture of LeNet-5


■ Convolution in Deep Learning

• What is Stride?


- **Stride** is the number of steps the filter moves each time.
- Default: stride = 1 (slide the filter **1 pixel** at a time)
- If stride > 1 → output becomes **smaller** (you skip some positions)




Stride 1



Feature Map




Stride 2



Feature Map

• What is Padding?


- **Padding** adds extra pixels (usually zeros) around the input image.
- Purpose
 - ✓ Keep output size the same as input (called "**same**" padding)
 - ✓ Prevent information loss at the edges



Architecture of LeNet-5

■ What is Convolution?

- “Convolution is a mathematical operation used **to combine two functions into a third one.**”



- “It captures how one function modifies or responds to another function as it shifts across time or space.”
- “Convolution is used to measure how similar two functions are at different alignments.”

Architecture of LeNet-5

■ What is Convolution?

- Let's define two continuous functions for convolution.

- $(f * g)(t) = \int_{-\infty}^{\infty} f(\tau) g(t - \tau) d\tau$



- Step-by-step Explanation


- 1. Reverse one function

✓ To compute the convolution, we first **flip** one of the two functions.

- Original function: $g(\tau)$
- Flipped version: $g(-\tau)$

✓ Why do we flip the function?

- A. To properly measure this similarity, one function (usually g) must be **reversed** so that it aligns correctly with the other.
- B. This is especially important when **matching a known pattern** (like in signal processing).
- C. Convolution is used to measure **how similar two functions are** at different alignments.



Architecture of LeNet-5

- What is Convolution?

- Step-by-step Explanation

- 1. Reverse one function

- ✓ Why do we flip the function?

- f : the signal (i.e., input data)

- g : the pattern or template (i.e., filter)

- What we want is To find where f and g match best — i.e., how similar they are at different positions.

- However, if we don't flip g , they might slide in opposite directions, and the result can be misleading.

- ✓ Let's Take a Simple Example

- Let $f = [1, 2, 3]$, $g = [3, 2, 1]$

- If we don't flip, and just multiply and sum:

- $1 \cdot 3 + 2 \cdot 2 + 3 \cdot 1 = 10$

- Now, flip g to get $g' = [1, 2, 3]$

- $1 \cdot 1 + 2 \cdot 2 + 3 \cdot 3 = 14$



The result is **higher and more meaningful** when we flip.

Architecture of LeNet-5

■ What is Convolution?

- Let's define two continuous functions for convolution.

- $(f * g)(t) = (f * g)(t) = \int_{-\infty}^{\infty} f(\tau) g(t - \tau) d\tau$

- Step-by-step Explanation

- 2. Shift the flipped function


✓ The function g is then **shifted** by t .

➤ Flipped version: $g(-\tau)$

➤ Shifted version: $g(-\tau + t) = g(t - \tau)$

✓ Why Do We Shift by t ?

- The variable t tells us **where** we are checking for similarity
- For each position t , we compare g to a **different part** of f .
- We are **sliding** g across the entire input f .



Think of It Like This

At $t = 0$: g starts at the beginning of f

At $t = 1$: g moves one step right

At $t = 2$: move again, and so on...

We compute a value at each t that tells us **how well** g matches f at that position.

Architecture of LeNet-5

■ What is Convolution?

- Let's define two continuous functions for convolution.

- $$(f * g)(t) = \int_{-\infty}^{\infty} f(\tau)g(t - \tau) d\tau$$

- Step-by-step Explanation

- 3. Multiply and accumulate

- ✓ At each shift, we compute the product $f(\tau)g(t - \tau)$ and then integrate $\int_{-\infty}^{\infty} \dots d\tau$ (i.e., sum) over all τ to produce a single output value.
 - ✓ This is like measuring how well the two functions overlap at each position.
→ "The Similarity of Two Functions"

- Intuitive Interpretation

- ✓ "How similar is the input f to the flipped and shifted version of g (i.e., filter)?"
 - Sometimes the overlap is large → high output.
Sometimes there's no overlap → output is zero.

Architecture of LeNet-5

■ What is Convolution?

- In deep learning, we use the **discrete version**.

- $$(f * g)(t) = \sum_{-\infty}^{\infty} f[m] \cdot g[n - m]$$

- **Why Is Convolution Discrete in CNNs?**

- ✓ 1. Input Data Is Discrete – CNNs process **digital images**, which are made of **pixels**. So, a digital image is a **2D array of values**, not a continuous function.
- ✓ 2. Efficient for Computation – **Continuous convolution** involves integration, which is slow and hard to compute.
- ✓ 3. Learnable Filters in Deep Learning – CNN filters are **learnable parameters**. Furthermore, continuous filters would be too complex to learn efficiently

- **Why No Flipping in CNNs?**

- ✓ In signal processing, convolution flips g to match known patterns.
- ✓ But in CNNs, we don't have a known pattern — we **learn the filter weights g !**
- ✓ So **flipping isn't needed**, and cross-correlation is used instead.

Why No Flipping in CNNs?



Architecture of LeNet-5

- **S2 – First Subsampling (Average Pooling) Layer**

- Input: **6 feature maps of size 28×28 (i.e., $6@28 \times 28$)**

- **Pooling operation:**

- Uses **2x2 average pooling**

- **Stride = 2** → reduces each map to **14x14**

- **Activation function:** tanh after pooling

- **Number of Parameters in S2**

- $(1 \text{ weight} + 1 \text{ bias}) \times 6 \text{ feature maps} = 12 \text{ parameters}$

- **Purpose of Pooling**


- Reduces spatial size

- Introduces **translation invariance**

- Keeps computational cost low

- Retains important structure before next convolution

- Provides **non-linearity** based on activation functions




Architecture of LeNet-5

■ S2 – First Subsampling (Average Pooling) Layer

- What Is Translation Invariance?

- A feature is still recognized even if it slightly moves in the input image.



- How Does Pooling Help?

- ✓ Pooling layers help build **shift invariance** in convolutional networks.
 - ✓ Shift invariance means that *the same maximum value will be found under the pooling kernel even if the image is shifted slightly.*
 - ✓ Focuses on “what exists”, not exactly “where”
 - ✓ *However, this shift invariance is only locally true and may not hold if the image is shifted too much.*

Architecture of LeNet-5

- S2 – First Subsampling (Average Pooling) Layer


- Activation Functions Add Non-Linearity After Pooling

- Why Is Non-Linearity Important?

- ✓ Without non-linearity, a neural network would just be a **linear function**
 - ✓ Linear layers stacked together still behave like a single linear layer
 - ✓ We need **non-linear functions** to model **complex patterns**

- Activation After Pooling

- ✓ In CNNs, pooling layers are typically followed by an **activation function** (like ReLU, tanh, etc.)
 - This introduces **non-linearity** into the model.



Architecture of LeNet-5

- S2 – First Subsampling (Average Pooling) Layer

- Activation Functions Add Non-Linearity After Pooling

- What Does "Linear" Mean?

- ✓ In mathematics, a function is **linear** if it satisfies

- 1. Additivity – $f(x + y) = f(x) + f(y)$

- 2. Homogeneity (scaling) – $f(a \cdot x) = a \cdot f(x)$

- ✓ Why Is That a Problem?

- If a neural network uses only **linear layers**, stacking multiple layers is still just one big linear transformation.

- “It cannot model complex patterns, curves, or decisions.”

- How Do Activation Functions Add Non-Linearity?

- ✓ Activation functions like: ReLu $f(x) = \max(0, x)$, Sigmoid $f(x) = \frac{1}{1+e^{-x}}$, tanh $f(x) = \tanh(x)$

- ✓ These functions break linearity: $f(a \cdot x) \neq a \cdot f(x)$


- “This allows the network to **learn complex decision boundaries and features.**”

Architecture of LeNet-5


■ S2 – First Subsampling (Average Pooling) Layer

- **Types of Pooling in CNNs**

- Average Pooling – (1) Computes the **mean value** in each window and (2) smooths out feature maps



- Max Pooling – (1) Selects the **maximum value** in the window and (2) Preserves strong activations, highlights key features



Architecture of LeNet-5

S2 – First Subsampling (Average Pooling) Layer

Types of Pooling in CNNs

- Global Average (or Max) Pooling


- ✓ GAP computes the **average value** of each feature map (channel)

- ✓ Converts a feature map of size $H \times W$ to a **single value** per channel

- ✓ Output becomes a $1 \times 1 \times C$ tensor ($C = \text{number of channels}$)

- ✓ For each channel, simply compute

$$\triangleright \text{GAP}(X_c) = \frac{1}{H \times W} \sum_{i=1}^H \sum_{j=1}^W X_c(i, j)$$



Why Was GAP Introduced?

- To replace the **fully connected (FC) layer** at the end of CNNs

- FC layers have: **(1) too many parameters** → high computational cost, **(2) High risk of overfitting**, **(3) Fixed input size** → limits flexibility


- GAP was proposed to solve these problems by: **(1) Reducing the number of parameters to zero**, **(2) Making the model lighter and more flexible**, **(3) Still allowing for end-to-end learning**

Architecture of LeNet-5

■ C3 – Second Convolution Layer

• Overview of C3 Layer

- Input: 6 feature maps from S2 layer (size 14×14)
- Filter: 5×5 size
- Output: 16 feature maps (size 10×10)
- Stride: 1, No padding
- **"But not all input maps are connected to all output maps!"**



• Why Not Full Connection?

- The original LeNet-5 paper gives two reasons
 - ✓ **(1) Reduce the number of connections** (parameter efficiency)
 - ✓ **(2) Encourage diversity:** Each output feature map extracts **different combinations of features**, avoiding redundancy

Architecture of LeNet-5

■ C3 – Second Convolution Layer

- Why Not Full Connection?

- Grouped Connections

	0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
0	X			X	X	X			X	X	X	X	X	X	X	
1	X	X			X	X	X	X		X	X	X	X	X	X	
2	X	X	X				X	X	X		X		X	X	X	
3		X	X	X			X	X	X	X		X	X	X	X	
4			X	X	X		X	X	X	X	X		X	X	X	
5				X	X	X		X	X	X	X	X	X	X	X	

TABLE I
EACH COLUMN INDICATES WHICH FEATURE MAP IN S2 ARE COMBINED
BY THE UNITS IN A PARTICULAR FEATURE MAP OF C3.

Output Maps (C3)	Input Maps (S2) Used	Number of Outputs
Group 1	3 sequential input maps	6
Group 2	4 sequential input maps	6
Group 3	4 non-sequential input maps	3
Group 4	all 6 input maps	1

✓ These groupings are manually chosen hyperparameters

Architecture of LeNet-5

■ C3 – Second Convolution Layer

- Parameter Calculation

- Each convolution filter has: *Filter size × input maps + bias*

Group	Formula	Parameters
Group 1	$(5 \times 5 \times 3 + 1) \times 6$	456
Group 2	$(5 \times 5 \times 4 + 1) \times 6$	606
Group 3	$(5 \times 5 \times 4 + 1) \times 3$	303
Group 4	$(5 \times 5 \times 6 + 1) \times 1$	151

✓ Total trainable parameters: $456+606+303+151=1,516$

- Why This Design Is Smart

- Reduces parameter count without sacrificing performance
- Prevents overfitting by reducing redundancy
- Promotes specialization in different output maps

Architecture of LeNet-5

S4 – Second Subsampling (Average Pooling) Layer

Overview of S4 Layer


- Input: 16 feature maps of size **10×10** (from C3)

- Pooling
 - ✓ Average Pooling
 - ✓ **Filter size: 2×2, Stride: 2**
 - ✓ **Output: 16 feature maps of size 5×5**

- Activation Function: tanh

- Learnable Parameters
 - ✓ Each feature map uses **1 learnable weight** and **1 bias**

$$\checkmark (\text{1 learnable weight} + \text{1 bias}) \times 16 = 32 \text{ parameters}$$



Architecture of LeNet-5

■ C5 – Third Convolution Layer

• Overview of C5 Layer

- Input: 16 feature maps of size 5×5 (from S4)

- Output: 120 feature maps of size 1×1

- How?

- ✓ Each filter covers all 16 input maps

- ✓ Filter size: $5 \times 5 \times 16$ (i.e., 5x5 per input map, across 16 maps)


- Activation Function: tanh

- Learnable Parameters

- ✓ Each output feature map has $5 \times 5 \times 16 = 400$ weights

- ✓ Plus 1 bias

- ✓ $(5 \times 5 \times 16 + 1) \times 120 = 48,120$ parameters



Architecture of LeNet-5

■ F6 – Fully Connected Layer

• Overview of F6 Layer

- **Input:** 120 values from previous C5 layer (1×1 feature maps)

- **Output:** 84 neurons

- **Connection**


- ✓ Fully connected to all 120 inputs

- ✓ Each output neuron also has a **bias**

- **Activation Function:** tanh

- **Parameters**

- ✓ Each output neuron \rightarrow 120 weights + 1 bias



✓ $(120+1) \times 84 = 10,164$ *parameters*

Architecture of LeNet-5

■ Output Layer

- **Overview of Output Layer**

- **Output:** 10 neurons
(for digit classification: 0 through 9)

- **Key Characteristics**


- ✓ Uses **Radial Basis Function (RBF)** units instead of softmax
 - ✓ Each neuron is connected to **all 84 units** from F6

- **Learning**

- ✓ **Backpropagation** is used to update the weights
 - ✓ The output neuron with the **highest activation** is selected as the predicted class

- **Why RBF?**

- ✓ The original LeNet-5 paper used **Euclidean distance-based RBF units**.
 - ✓ These act like **prototypes** for each digit class and measure **similarity** to input.



Summary – What We Learned about LeNet-5

■ What is LeNet-5?

- One of the **first CNNs**, developed by **Yann LeCun** in 1998 for **digit recognition** (MNIST).
- Pioneered **end-to-end learning** from raw pixels using learned filters.
- Introduced key design ideas still used in modern CNNs.
- Architecture Breakdown

Layer	Type	Output Shape	Key Details
C1	Convolution	6@28×28	5×5 filters, stride=1
S2	Avg Pooling	6@14×14	2×2 pooling, tanh activation
C3	Convolution	16@10×10	Not fully connected – grouped connections
S4	Avg Pooling	16@5×5	2×2 pooling, learnable weights
C5	Convolution	120@1×1	5×5×16 filters (full connection)
F6	Fully Connected	84	Each neuron gets all 120 inputs
Output	RBF Layer	10	Radial Basis Function units

Assignment 1 – Dive into LeNet-5

■ What you will explore

- Please revisit the following questions based on your reading and implementation

- Q1. Why did LeNet-5 use RBF in the output layer instead of softmax?
- Q2. How is LeNet-5 different from modern CNN architectures?
- Q3. Why is C3 only partially connected to S2?
- Q4. Why is LeNet-5 still important today?

- Your Task

- Assignment 1

- ✓ Read the original paper “*Gradient-Based Learning Applied to Document Recognition*”
 - ✓ Understand the full design and reasoning behind LeNet-5
 - ✓ Answer the above four questions
 - ✓ Implement LeNet-5 and Train LeNet-5 on MNIST Dataset