

본 수업자료는 2025년도 과학기술 정보통신부 및 정보통신기획평가원의 'sw중심대학사업' 지원을 받아 제작 되었습니다.

# **ComputerVision**

## **Week3**

2025-2

Mobile Systems Engineering  
Dankook University

# Recap – What We Learned So Far

---

- Key Topics Reviewed

- **What is a Convolutional Neural Network (CNN)?**

- A type of deep neural network designed to process data with grid-like topology (e.g., images). It automatically learns spatial hierarchies of features.

- **Basic Structure of LeNet-5**

- LeNet-5 was one of the earliest CNNs designed for digit recognition (MNIST). It consists of alternating **convolution**, **pooling**, and **fully connected layers**.

- **Feature Extraction in CNNs**

- Through convolution and pooling, CNNs transform raw pixel values into **high-level features**, such as edges, textures, or object parts.

# Recap – What We Learned So Far

---

## ■ Key Topics Reviewed

- Typical CNN Layer Structure (from LeNet-5)

Layer	Type	Output Shape	Key Details
C1	Convolution	6@28×28	5×5 filters, stride=1
S2	Avg Pooling	6@14×14	2×2 pooling, tanh activation
C3	Convolution	16@10×10	Not fully connected – grouped connections
S4	Avg Pooling	16@5×5	2×2 pooling, learnable weights
C5	Convolution	120@1×1	5×5×16 filters (full connection)
F6	Fully Connected	84	Each neuron gets all 120 inputs
Output	RBF Layer	10	Radial Basis Function units

- **Convolution Layer:** Detects local patterns using filters (kernels)
- **Pooling Layer:** Reduces spatial dimensions and helps generalization
- **Fully Connected Layer:** Performs classification based on extracted features

# Motivation for Going Deeper

---

- Why Do We Need Deeper Networks?

- Shallow Networks

- Work well for **simple tasks** like digit recognition (e.g., MNIST)
    - Can only capture **low-level features** like edges or simple textures
    - Quickly **saturate** in performance as task complexity increases


- Complex Datasets Need Abstraction

- Real-world images (e.g., ImageNet) include complex variations:
      - ✓ Background clutter
      - ✓ Different lighting, poses, textures
    - Models need to learn **hierarchies of features**:
      - ✓ Low-level: edges, corners
      - ✓ Mid-level: textures, patterns
      - ✓ High-level: object parts, semantic meaning

# Motivation for Going Deeper

---


- Why Do We Need Deeper Networks?
  - Deeper Networks Can Build Feature Hierarchies



- Each additional layer can capture more abstract concepts
- Final layers make decisions based on **high-level understanding**
- This leads to better generalization and classification accuracy

# Introduction to AlexNet

## ■ The Turning Point — ImageNet 2012



- AlexNet won the ImageNet Large Scale Visual Recognition Challenge (ILSVRC) in 2012.
- Reduced top-5 error from ~26% to 16.4% — a groundbreaking result.
- *“Marked the start of the deep learning era in computer vision.”*

# Introduction to AlexNet

---

## ■ Who Built AlexNet?



- Developed by **Alex Krizhevsky**, a student of **Geoffrey Hinton**, with support from **Ilya Sutskever**.
- Hinton motivated Alex with an unusual bet: "**If you improve the accuracy by 1%, I'll let you skip the qualifying exam.**"
- AlexNet exceeded expectations and made history.

# Introduction to AlexNet

---

## ■ What is ImageNet?





- Created by Fei-Fei Li in 2006.
- Contains **14 million+ labeled images**, classified into **20,000+ categories**.
- Enabled large-scale training of deep models.
- Used for the annual ILSVRC competition (2010–2017).

# Introduction to AlexNet

## ■ What Made AlexNet Special?

- Architecture: 5 convolutional layers + 3 fully connected layers
- ReLU instead of sigmoid or tanh → **faster training**
- Dropout → **reduced overfitting**
- Trained on **2 NVIDIA GTX 580 GPUs (3GB)** for scalability

## ■ Architectural Diagram of AlexNet



- Add visual comparison of LeNet and AlexNet (already attached by user)
- Highlight how AlexNet expands LeNet's ideas with deeper layers and more filters

# Introduction to AlexNet

---

## ■ Technical Innovations

- **ReLU** non-linearity made training faster
- **Split training over two GPUs** to fit in memory (model parallelism)
- **Data augmentation**: random crops, flips to prevent overfitting
- **Local Response Normalization (LRN)** to encourage competition between neurons

## ■ Why AlexNet Was a Breakthrough


- Demonstrated that deep CNNs could beat traditional hand-crafted features (e.g., SIFT, HOG)
- Proved that **data + computing power** unlocks deep learning's potential
- Inspired new models: VGG, GoogLeNet, ResNet

### Legacy of AlexNet

1. Opened the era of modern deep learning
2. Changed how AI research and industry approached vision tasks
3. GPUs became essential hardware for training deep models

# AlexNet Overall Architecture

## ■ What Does the Architecture of AlexNet Look Like?




- 8 layers with trainable parameters: **5 Convolutional layers and 3 Fully Connected layers**
- **Input:** RGB image,  $227 \times 227 \times 3$  channels (Red, Blue, Green)
- **Output:** 1000-class softmax
- Uses **two GPUs** to split the computation

# AlexNet

## ■ Layer 1 — Convolution + ReLU + Pooling + LRN

- **Convolution layer**

- Filter (i.e., kernel): 96 filters of size  $11 \times 11 \times 3$   
→ 96 output channels




- **Stride = 4, No padding**

- **Output size:  $55 \times 55 \times 96$  channels**

- **ReLU activation**

- **Overlapping max pooling:  $3 \times 3$  window, stride 2**

- **Local Response Normalization (LRN)**




# AlexNet

---

- Layer 1 — Convolution + ReLU + Pooling + LRN

- What is ReLU and Why Was It a Game-Changer?

- What is ReLU?



- ✓ ReLU stands for **Rectified Linear Unit**

- ✓ Formula:  $\text{ReLU}(x) = \max(0, x)$

- ✓ Graph:

- Negative input  $\rightarrow 0$


- Positive input  $\rightarrow$  linear increase

# AlexNet

## ■ Layer 1 — Convolution + ReLU + Pooling + LRN

- What is ReLU and Why Was It a Game-Changer?

- Why ReLU Helped AlexNet Succeed



Property	Sigmoid / tanh (Saturating)	ReLU (Non-saturating)
Output range		
Gradient near 0		
Computation		
Convergence speed		

- ✓ Non-saturating nonlinearity → Avoids the vanishing gradient problem
- ✓ Faster training → ReLU allowed AlexNet to converge **6x faster** than tanh
- ✓ Helped AlexNet handle a **very deep architecture** with **60M** parameters
- ✓ Without ReLU, large-scale training on ImageNet would have been **impractical**

- Layer 1 — Convolution + ReLU + Pooling + LRN

- What is Local Response Normalization (LRN) and Why Was It a Game-Changer?

- What is Local Response Normalization?

- ✓ LRN is a normalization technique applied **across channels** at the same spatial location  $(x, y)$
      - ✓ Inspired by **lateral inhibition** in neuroscience → strong activations suppress weaker neighbors
      - ✓ Applied **after ReLU** to encourage feature diversity

- Mathematical Formula of LRN

$$b_{x,y}^i =$$


- ✓  $a_{x,y}^i$ : ReLU output at location  $(x, y)$  in the  **$i$ -th** channel
      - ✓  $N$ : total number of feature maps
      - ✓  $n$ : local window size (typically 5)
      - ✓  $k = 2, \alpha = 10^{-4}, \beta = 0.75$  (as used in the AlexNet paper)

# AlexNet

## ■ Layer 1 — Convolution + ReLU + Pooling + LRN

- What is Local Response Normalization (LRN) and Why Was It a Game-Changer?

- Visual Explanation of LRN



- ✓  $n$ : local window size (typically 5)
- ✓  $k = 2, \alpha = 10^{-4}, \beta = 0.75$
- ✓  $a^i_{x,y}$ : ReLU output at location  $(x, y)$  in the  $i$ -th channel
- ✓  $N$ : total number of feature maps

# AlexNet

---

## ■ Layer 1 — Convolution + ReLU + Pooling + LRN

- What is Local Response Normalization (LRN) and Why Was It a Game-Changer?

- Why Use LRN?

- ✓ After ReLU, some neurons have **very large outputs**, which can dominate learning
    - ✓ LRN **suppresses extreme activations** by normalizing each response with nearby channels
    - ✓ This promotes **competition** among neurons → more diverse and selective features

- Where Was LRN Used?

- ✓ Not applied in every layer
    - ✓ Used only in **Conv1 and Conv2** in AlexNet
    - ✓ Helped reduce
      - Top-1 error by 1.4%
      - Top-5 error by 1.2%

# AlexNet

## ■ Layer 1 — Convolution + ReLU + Pooling + LRN

- What is Overlapping Max Pooling and Why Was It a Game-Changer?

- What is Pooling in CNNs?

- ✓ Pooling reduces the **spatial size** of feature maps
    - ✓ Helps make representations **more compact** and **robust to translation**
    - ✓ Most common type

- **Max Pooling** — selects the largest value in a region

- ✓ Traditionally used with **non-overlapping filters** (e.g., kernel=2, stride=2)

- What is Overlapping Pooling?

- ✓ Uses a **smaller stride** than the kernel size
    - ✓ Example: kernel = 3x3, **stride = 2**
    - ✓ Pooling windows **overlap** instead of being disjoint

1	1	2	4
5	6	7	8
3	2	1	0
1	2	3	4

max pool with 2x2 filters and stride 2

6	8
3	4

1	1	2	4
5	6	7	8
3	2	1	0
1	2	3	4

max pool with 2x2 filters and stride 1

6	7	8
6	7	8
3	3	4

# AlexNet

---

## ■ Layer 1 — Convolution + ReLU + Pooling + LRN

- What is Overlapping Max Pooling and Why Was It a Game-Changer?

- Why Use Overlapping Pooling?

Metric	Non-Overlapping (Stride = 2)	Overlapping (Stride = 1)
Information		
Receptive field usage		
Accuracy (AlexNet paper)		
Convergence speed		


- When to Use Overlapping Pooling?

- ✓ Useful in **early layers** to retain more spatial detail
    - ✓ Helpful when the input is small, and information loss matters
    - ✓ More computation than non-overlapping pooling
    - ✓ Less commonly used today (often replaced by other techniques like strided convolutions or attention)

# AlexNet

## ■ Layer 2 — Parallel Convolution + Pooling + LRN

- Parallel GPU Processing
  - 128 filters (i.e., kernels) of size  $5 \times 5 \times 48$  (2 sets due to 2 GPUs)
  - Stride = 1, Padding = 2
  - Output:  $27 \times 27 \times 256$  ( $128 \times 2$  GPUs)
  - Overlapping max pooling:  $3 \times 3$ , stride 2  $\rightarrow 13 \times 13 \times 256$
  - LRN applied again



# AlexNet

## ■ Layer 3 to 5 — Depth & Inter-GPU Communication

- Deeper Layers for Rich Features

- Layer 3


- ✓ 384 filters of size  $3 \times 3 \times 256$

- Layer 4

- ✓ 192 filters of size  $3 \times 3 \times 192$   
(separate on each GPU)

- Layer 5

- ✓ 128 filters of size  $3 \times 3 \times 192$ ,  
followed by max pooling  $\rightarrow 6 \times 6 \times 256$



# AlexNet

## ■ Layers 6–8 — Fully Connected & Classification

- **Fully Connected Layers**

- Layer 6: FC(9216 → 4096), ReLU, Dropout
- Layer 7: FC(4096 → 4096), ReLU, Dropout
- Layer 8: FC(4096 → 1000), Softmax

- **What Are Fully Connected Layers?**

- A **Fully Connected (FC)** layer connects **every neuron** from the previous layer to **every neuron** in the current layer.
- Equivalent to a **dense matrix multiplication**



$$\checkmark \mathbf{y} =$$

- What Do FC Layers Do?

- ✓ 1.

- ✓ 2.


- ✓ 3.



# From AlexNet to VGGNet (2014)

## ■ Introduction to VGGNet (2014)

- Developed by **Simonyan** and Zisserman at Oxford's Visual Geometry Group (VGG)
- Runner-up at **ImageNet Challenge 2014**
- Known for simplicity and depth: VGG16 and VGG19
- Uses **only 3x3 convolution** filters throughout the network



# VGGNet

---

- Motivation and Design of VGGNet
  - Research Goal
    - Investigate how **network depth** affects image recognition accuracy.
  - Fixed Design Choice
    - All convolutional layers use **3×3 filters** (small receptive field), with **stride = 1**, and **padding = 1**.  
This choice allows deeper stacking while keeping spatial resolution stable.
  - Progressive Deepening
    - VGG models are built with **increasing depth** — from 11, 13, 16 to 19 weight layers.  
(These are referred to as configurations A–E in the paper.)
  - Observation
    - As the depth increased, **classification error consistently decreased** on ImageNet.  
This confirmed that deeper networks can capture more complex patterns and improve accuracy.

# VGGNet

## ■ Comparison of VGG16 and VGG19 Architectures

- Common Traits


- Input: Both take a **224×224 RGB** image as input
- Conv Filter: Use only **3×3 convolution filters** throughout
- Pooling: **Include 5 max pooling layers**
- FC: **End with 3 fully connected layers**: 4096 → 4096 → 1000
- Activation: **ReLU activation** used after each conv layer
- No Local Response Normalization (LRN)

- **VGG16**

- **13 convolutional layers + 3 fully connected layers**
- Total 16 weight layers
- Conv blocks: 2-2-3-3-3 pattern

- **VGG19**

- **16 convolutional layers + 3 fully connected layers**
- Total 19 weight layers
- Conv blocks: 2-2-4-4-4 pattern



VGG16



VGG19

# VGGNet


## ■ Comparison of VGG16 and VGG19 Architectures

### • Difference from AlexNet


- Much deeper (8 layers in AlexNet vs. 16/19 in VGGs)
- Simpler and more uniform design (only  $3 \times 3$  filters)
- Better performance despite higher depth due to small filters and regular structure



AlexNet



VGG16



VGG19

# VGGNet

## ■ VGG16

- **Architecture Overview**

- **Input:** 224×224 RGB image
- **Layers:** 13 convolutional layers + 3 fully connected layers
- **Conv Filter:** Convolution layers use 3×3 filters, stride = 1, padding = 1
- **Pooling:** Max pooling: 2×2 with stride 2 after certain conv blocks
- **Activation:** ReLU used after each conv layer
- **No Local Response Normalization (LRN) in deeper models**

- **VGG16 Layer-by-Layer Breakdown**

- **Conv1\_1, Conv1\_2:** filters → MaxPooling →
- **Conv2\_1, Conv2\_2:** filters → MaxPooling →
- **Conv3\_1 to Conv3\_3:** filters → MaxPooling →
- **Conv4\_1 to Conv4\_3:** filters → MaxPooling →
- **Conv5\_1 to Conv5\_3:** filters → MaxPooling →
- **FC1, FC2:** neurons each
- **FC3:** neurons → Softmax



# VGGNet

---

## ■ VGG16

- Why Stack  $3 \times 3$  Convolutions?

- Key Idea
  - ✓ Instead of using large filters (e.g.,  $5 \times 5$  or  $7 \times 7$ ), VGGNet stacks multiple  $3 \times 3$  convolution layers to achieve the same receptive field but with **more non-linearity** and **fewer parameters**.
- Stacking two  $3 \times 3$  convs =  $5 \times 5$  receptive field
- Three  $3 \times 3$  convs =  $7 \times 7$  receptive field

- Comparison of Receptive Fields


Filter Strategy	Effective Receptive Field	Nonlinear Layers	Parameters (per channel)
One $5 \times 5$ convolution			
Two stacked $3 \times 3$ convolutions			
One $7 \times 7$ convolution			
Three stacked $3 \times 3$ convolutions			

- Stacking  $3 \times 3$ s achieves **same receptive field** with:
- **More non-linearity (ReLU)** → better learning capacity
- **Fewer parameters** → less overfitting risk and faster training

# VGGNet

## ■ VGG16

- Benefits of  $3 \times 3$  Stacking



- **(1) Increased Nonlinearity:** Each convolution is followed by ReLU, increasing the model's capacity to learn complex patterns.
- **(2) Reduced Parameters:** For same output depth  $C$ , using  $3 \times 3 \times C \times C$  multiple times requires fewer weights than one large  $5 \times 5$  or  $7 \times 7$  convolution.
- **(3) Faster Training:** Fewer parameters mean faster convergence and less memory usage.
- **(4) Modular Design:** Allows deeper networks without exploding parameter count.

# VGGNet

---

## ■ Fully Connected Layers in VGGNet

- **Structure**

- Flattened output from conv layers:  $7 \times 7 \times 512 = 25,088$
- FC1: 4096 units + **ReLU** + **Dropout (p = 0.5)**
- FC2: 4096 units + **ReLU** + **Dropout (p = 0.5)**
- FC3: 1000 units (for classification) + **Softmax**

- **Additional Notes**


- The FC layers contain the majority of the model's parameters.
- **ReLU** enables **faster training** and introduces **non-linearity**.
- **Dropout** helps reduce overfitting.
- These FC layers can later be converted into convolutional layers to form a **Fully Convolutional Network**.

## ■ What is Dropout?

- Concept

- Dropout is a regularization technique used in neural networks.
- It works by randomly “dropping out” neurons during training with a certain probability, called the **dropout rate**.
- Dropped neurons are temporarily removed along with their connections, meaning they do not participate in **forward propagation or backpropagation** for that training step.

- Example (with Dropout Rate = 0.5)



- In the figure, the left side shows a fully connected layer where all neurons are active.
- On the right side, with dropout rate = 0.5, about **half of the neurons are removed at random**.
- For instance, in this example, 2 out of 4 neurons are dropped. The specific neurons that are dropped can change in each training iteration.

## ■ What is Dropout?

- **Key Notes**

- 1. The **dropout rate** is a **hyperparameter**.
  - ✓ A common choice is **0.5** for fully connected layers.
- 2. During training, dropout makes the network less dependent on specific neurons and forces it to learn more **robust and general features**.
- 3. At test time, however, dropout is **NOT applied**.
  - ✓ Instead, all neurons are used with scaling to keep outputs consistent.
- 4. Dropout helps **prevent overfitting** by reducing reliance on certain strong features, and it also creates an **ensemble effect**, since each random dropout configuration can be seen as training a different sub-network.
  - ✓ At inference, combining all neurons is like averaging many models together, improving **generalization**.

## ■ Why Do We Use Dropout?

- **Preventing Overfitting**


- Without dropout, a network can easily **overfit**, meaning it relies too heavily on certain strong features.
- Overfitting reduces the model's ability to generalize to unseen data.
- Dropout forces the model to learn **robust representations**, since it cannot depend on the presence of specific neurons every time.

- **Ensemble Effect**

- Each random dropout configuration creates a **different sub-network**, called a **realization (or instance)**.
- During training, the network optimizes many different sub-networks in parallel.
- At test time, all neurons are used together, which is similar to averaging the outputs of many networks.
- This process acts like an **ensemble of models**, which improves accuracy and reduces bias.

## ■ Why Do We Use Dropout?

- Key Idea on Ensemble Effect




- On the left, we see multiple realizations where different neurons are dropped in each instance.
- On the right, all of these sub-networks are combined into a single larger network.
- This shows that dropout not only prevents overfitting but also gives the benefits of an **ensemble method** without the need to train many separate models.

## ■ Dropout in Mini-Batch Training



### • How Dropout Works with Mini-Batches

- Dropout is applied **independently to each mini-batch** during training.
- With a dropout rate of **0.5**, each neuron has a 50% chance of being dropped in every batch.
- The neurons that are dropped can vary across different batches.

### • Example



**Drop-out Rate = 0.5**




- In **Batch 1**, neurons 2 and 3 are dropped.
- In **Batch 2**, only neuron 3 is dropped.
- In **Batch 3**, neurons 1, 2, and 3 are dropped together.
- This randomness ensures that the network does not rely on specific neurons across all training data.

## ■ Dropout in Mini-Batch Training

### • Why This is Important

- By changing the dropped neurons in every batch, dropout forces the model to learn **multiple redundant representations**.
- This leads to stronger **generalization** since no single pathway can dominate the learning process.




- Because each mini-batch trains a slightly **different sub-network**, dropout creates the effect of training an **ensemble of networks**.
- ✓ At test time, when all neurons are used together, the model behaves like an **averaged ensemble**, which improves stability and accuracy.

## ■ Dropout During Testing

- Key Difference from Training

- During **training**, some neurons are randomly dropped with probability  $\alpha$  (dropout rate).
- During **testing (inference)**, all neurons are active, but their outputs are **scaled** to match the expected values from training.

- Scaling Formula



$$h_n =$$

- Here,  $a$  is the activation function, and  $\alpha$  is the dropout rate.
- The factor  $(1 - \alpha)$  adjusts the output so that the average activation at test time is consistent with training.

## ■ Dropout During Testing

- Why Scaling is Needed

- In training, fewer neurons are active on average due to dropout.
- Without scaling, test outputs would be **systematically larger**, since all neurons are active.
- Scaling ensures **fair comparison** between training and testing phases, keeping the same magnitude of outputs.

- Key Takeaway

- At **training time**: Dropout improves generalization by dropping neurons.
- At **test time**: No neurons are dropped, but **scaling compensates for dropout**, ensuring consistent behavior and stable predictions.

## ■ Data Augmentation and Scale Jittering

### • Data Augmentation Techniques

- Resize each training image so that the **shortest side  $\geq 256$**
- Random crop:  $224 \times 224$  window
- Random horizontal flip
- Color jittering

### • Scale Jittering

- Single-scale training: fixed  $S = 256$
- Multi-scale training:  $S$  randomly sampled from  $[256, 512]$



224x224      256x256      512x512




224x224      224x224



256x256



224x224      224x224



512x512


224x224

224x224

## ■ Data Augmentation and Scale Jittering

- Advantages

- Increases dataset diversity
- Captures different **object contexts**
  - ✓ Small crops → entire object
  - ✓ Large crops → object parts
- Helps **reduce overfitting**
- Experiments showed **multi-scale training improved classification accuracy** over single-scale



224x224



256x256



512x512

isotropically-rescaled training image

