

# **ComputerVision**

**Week4 – 5**

2025-2

Mobile Systems Engineering  
Dankook University

# Why Go Deeper?

---

- The Promise of Depth in Deep Neural Networks
  - Deep neural networks have revolutionized visual recognition.
  - More layers = Better representations:
    - Capture low-level (edges), mid-level (textures), and high-level (objects) features.
  - Empirically proven
    - VGG, GoogLeNet: depth correlates with improved performance on ImageNet and COCO.



**BUT... deeper networks are hard to train!**

# The Optimization Challenges

---

- When More Layers Make Things Worse

- 1. Vanishing/Exploding Gradients

- As gradients are backpropagated, they can diminish (vanish) or blow up (explode).
    - Leads to unstable or slow training.

- 2. Saturation of Accuracy

- Adding layers sometimes leads to **higher training error**, not just test error.
    - Known as the “**degradation problem**”.


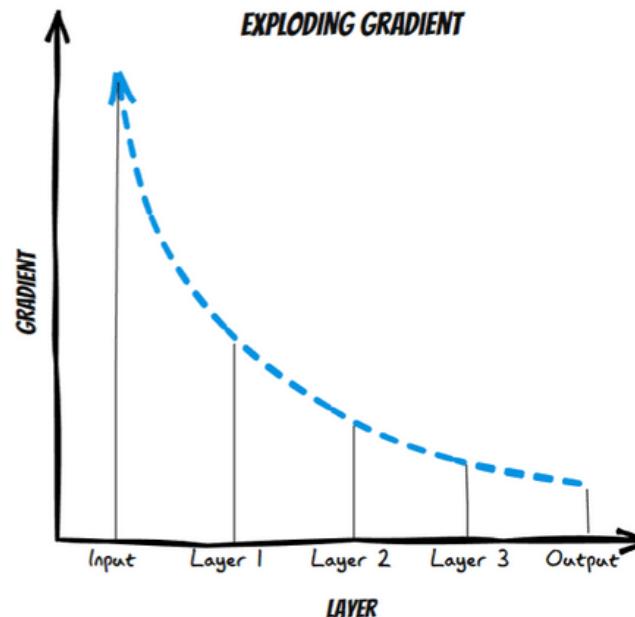
- 3. Intuition

- If we take a shallower network and add more layers, the deeper one should perform at least as well.
    - But in reality: deeper plain networks perform worse.

# What Are Vanishing and Exploding Gradients?

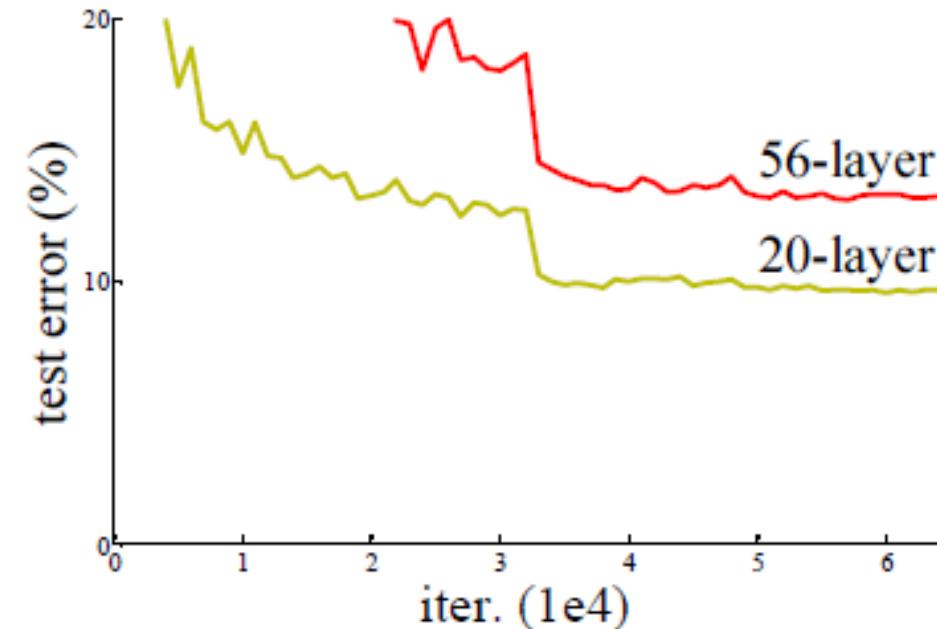
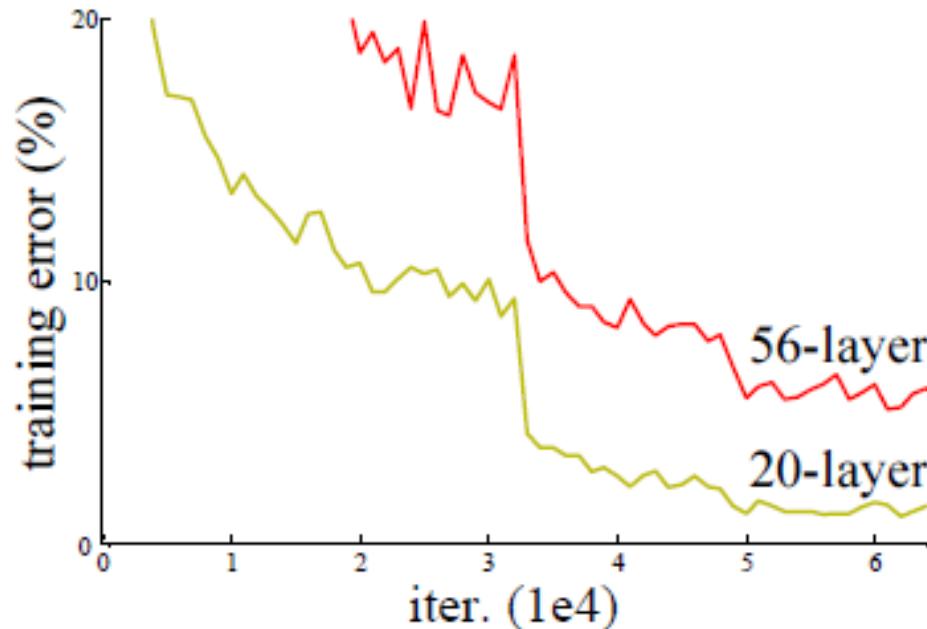
## ■ Vanishing and Exploding Gradients

- As we train deep neural networks using **backpropagation**, gradients are propagated backward through layers.
- During this process
  - Gradients may **shrink toward zero** (vanish)
  - Or **grow excessively large** (explode)
- Both scenarios make training **unstable or slow**.



# Degradation Problem (Empirical Evidence)

- Deeper Plain Networks Have Higher Training Error
  - Observation from Deep Residual Learning for Image Recognition (ResNet) – He et al. (2015)



- Train plain networks of 20 and 56 layers on CIFAR-10
- Despite more capacity, **56-layer network performs worse** (training + test error ↑)  
→ *“Indicates optimization failure, not overfitting!”*

# Understanding the Degradation Problem

---

- There Exists a Simple Solution — But It's Hard to Learn

- The deeper model has a **constructed solution**
  - Suppose you have a well-trained shallow network.
  - You can always create a deeper version by:
    - ✓ Copying the original layers.
    - ✓ Adding extra layers that perform **identity mapping** (i.e., output = input).
  - In theory, this should produce at least equal performance.
- So why does performance degrade in practice?
  - It's not due to **overfitting** (since training error increases).
  - It's not due to **vanishing gradients** (since techniques like BatchNorm are applied).

# Understanding the Degradation Problem

---

- There Exists a Simple Solution — But It's Hard to Learn

- The real problem

- Stochastic Gradient Descent (SGD) fails to find the identity mapping.
    - The network **struggles to learn a perfect "pass-through" behavior through non-linear transformations.**
    - Even a seemingly trivial task — like learning to output the same input — becomes difficult for deep non-linear layers.

- Key Insight

- Depth alone is not the problem — **learning identity mappings through standard layers is.**

# Need for Better Formulation

## ▪ Residual Learning Reformulation From Learning Mappings to Learning Residuals

- Key Insight

- Rather than learn the full mapping  $H(x)$ , let's learn only what is missing from the input.
- Define residual function:  $F(x) = H(x) - x \Rightarrow H(x) = F(x) + x$ .


- Residual Learning Reformulation

- Original mapping:  $H(x)$
- Reformulated:  $H(x) = F(x) + x$

- Why this helps

- If the optimal mapping is close to identity, then
  - ✓  $F(x)$  is close to zero  $\rightarrow$  easier to learn.
- Even if  $H(x)$  is complex, it may still be easier to express the difference from the input than to learn the entire function from scratch.

*"This concept is like preconditioning in numerical optimization — reshaping the problem to make it easier for solvers to converge."*



# Analogy: Learning Perturbations

---


- **Residuals as Minor Adjustments to Known Inputs**

- Think of  $F(x)$  as a **small correction or delta** to the input
    - For example, if the correct output is close to the input, the network only needs to learn the difference.


- **Identity mapping as the “default path”**

- The skip connection directly carries the input forward.
    - If the network doesn't need to change it much, the residual function learns small tweaks.
    - If the input needs major changes, the residual function takes full control.

- **Example – Writing**



VS



# Summary of Motivation

---

- Why Residual Learning Became a Breakthrough
  - Deep networks are **theoretically powerful**, but practically **difficult to optimize**.
  - As depth increases, **plain networks**
    - Encounter vanishing gradients, even with tricks like BatchNorm.
    - Show **increased training error** — a clear sign of optimization issues.
  - **Residual learning provides a solution**
    - Skip connections allow gradients to flow unimpeded.
    - Identity mappings are **easy to learn** with this architecture.
    - Residual blocks make it easier for the optimizer to converge.
  - **Impact**
    - Enables networks with **>100 layers** to be trained effectively.
    - Became the foundation for
      - ✓ ResNet, ResNeXt
      - ✓ Faster R-CNN (backbone)
      - ✓ Mask R-CNN, and more.

# What Are Vanishing and Exploding Gradients?

---

## ■ What is a Computational Graph?

- A computational graph is a **directed acyclic graph** representing a function.
  - **Nodes:** operations (e.g.,  $+$ ,  $\times$ , ReLU, etc.)
  - **Edges:** flow of values (scalars, vectors)
- Enables automatic differentiation via backpropagation
- **Example**
  - If  $z = x + y$ , and  $L = \sin(z)$ , then:
  - **Forward pass:** compute  $z = x + y$ , then  $L = \sin(z)$
  - **Backward pass:** compute  $\frac{dL}{dz}, \frac{dL}{dx}, \frac{dL}{dy}$

# What Are Vanishing and Exploding Gradients?

## ■ Computational Graph Example

- Function:  $L = (x + y) \cdot z$

- Graph

- Inputs:  $x, y, z$
  - Ops: addition  $x + y$ , multiplication with  $z$
  - Final node: loss  $L$


- → Forward Pass: compute each value

- Intermediate variable:  $q = x + y = -2 + 5 = 3$
  - Final output:  $f = q \cdot z = 3 \cdot (-4) = -12$

Backpropagation: a simple example

$$f(x, y, z) = (x + y)z$$

e.g.  $x = -2, y = 5, z = -4$



# What Are Vanishing and Exploding Gradients?

## ■ Scalar Backpropagation

- Function:  $L = (x + y) \cdot z$

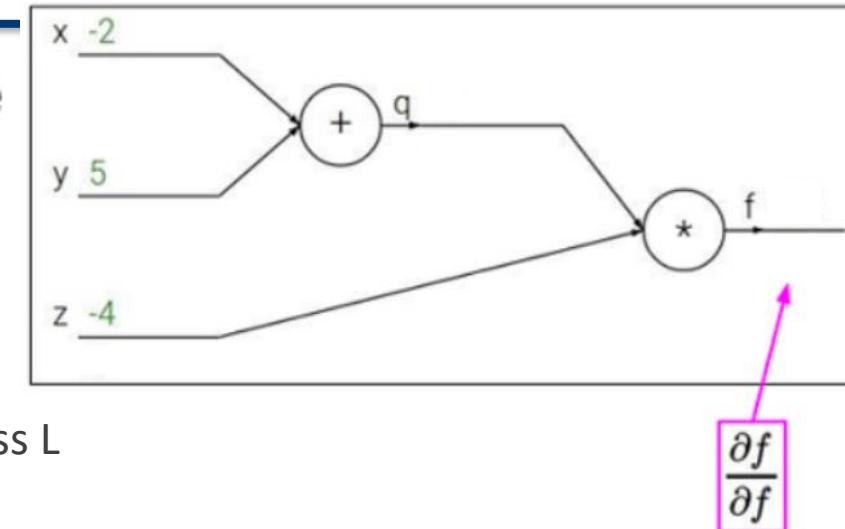
Backpropagation: a simple example

$$f(x, y, z) = (x + y)z$$

e.g.  $x = -2, y = 5, z = -4$

- Graph

- Inputs:  $x, y, z$  / Ops: addition  $x + y$ , multiplication with  $z$  / Final node: loss  $L$



## • ← Backward Pass (Backpropagation)

- We want –  $\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y}, \frac{\partial f}{\partial z}$

- Step 1 – From output to  $q$  and  $z$ :  $\frac{\partial f}{\partial q} = z = -4$  ,  $\frac{\partial f}{\partial z} = q = 3$

- Step 2 – From  $q$  to  $x$  and  $y$ :  $\frac{\partial q}{\partial x} = 1$  ,  $\frac{\partial q}{\partial y} = 1$

✓ Apply chain rule

$$\triangleright \frac{\partial f}{\partial x} = \frac{\partial f}{\partial q} \cdot \frac{\partial q}{\partial x} = -4 \cdot 1 = -4$$

$$\triangleright \frac{\partial f}{\partial y} = \frac{\partial f}{\partial q} \cdot \frac{\partial q}{\partial y} = -4 \cdot 1 = -4$$

# What Are Vanishing and Exploding Gradients?


## ■ Vector Backpropagation

### • Extending Backpropagation to Matrix Operations

- Goal: Given a vector input  $x$  and a weight matrix  $W$ , we aim to compute the gradient of the function

$$f(x, W) = \|Wx\|^2 = \sum_{i=1}^n (Wx)_i^2$$

- ✓  $x \in \mathbb{R}^d$ : input vector (column)
- ✓  $W \in \mathbb{R}^{n \times d}$ : weight matrix (trainable parameters)
- ✓  $q = Wx \in \mathbb{R}^n$ : intermediate vector
- ✓  $f = \|q\|^2 = q^T q$ : scalar output



# What Are Vanishing and Exploding Gradients?

---


- Vector Backpropagation

- Forward and Backward Pass (Step-by-Step)

- Step 1: Forward Pass

$$\checkmark x = \begin{bmatrix} 0.2 \\ 0.4 \end{bmatrix}, W = \begin{bmatrix} 0.1 & 0.5 \\ -0.3 & 0.8 \end{bmatrix}, q = Wx = \begin{bmatrix} 0.1 \cdot 0.2 + 0.5 \cdot 0.4 \\ -0.3 \cdot 0.2 + 0.8 \cdot 0.4 \end{bmatrix} = \begin{bmatrix} 0.22 \\ 0.26 \end{bmatrix}$$

$$\checkmark f = \| q \|^2 = 0.22^2 + 0.26^2 = 0.116$$



# What Are Vanishing and Exploding Gradients?


## ■ Vector Backpropagation

### • Forward and Backward Pass (Step-by-Step)

- Step 2.1: Backward Pass – Gradient w.r.t.  $q$

$$\checkmark f(x, W) = \|Wx\|^2 = \|q\|^2 = q^T q$$

$$\checkmark \frac{\partial f}{\partial q} = \nabla_q f = 2q = \begin{bmatrix} 0.44 \\ 0.52 \end{bmatrix}$$



- ✓  $\nabla$ : It represents a **gradient operator** in multivariable calculus.

➤ If you have a function:  $f(x_1, x_2, \dots, x_n)$ , then the **gradient** of  $f(\nabla f)$  is denoted by

$$\nabla f = \begin{bmatrix} \frac{\partial f}{\partial x_1} \\ \frac{\partial f}{\partial x_2} \\ \vdots \\ \frac{\partial f}{\partial x_n} \end{bmatrix}$$

➤ The gradient is a **vector** that points in the direction of **steepest increase** of the function  $f$ .

### ✓ In Backpropagation

➤  $\nabla_x f$ : How the scalar output  $f$  changes with respect to vector input  $x$ .

➤  $\nabla_W f$ : The matrix of partial derivatives showing how each weight affects the loss.

➤ **Example:** If  $f(x, y) = x^2 + y^2$ , then:

$$\nabla f = \begin{bmatrix} \frac{\partial f}{\partial x} \\ \frac{\partial f}{\partial y} \end{bmatrix} = \begin{bmatrix} 2x \\ 2y \end{bmatrix}$$

# What Are Vanishing and Exploding Gradients?

## ■ Vector Backpropagation

- Forward and Backward Pass (Step-by-Step)

- Step 2.1: Backward Pass – Gradient w.r.t.  $q$

✓  $f(x, W) = \|Wx\|^2 = \|q\|^2 = q^T q$

✓  $\frac{\partial f}{\partial q} = \nabla_q f = 2q = \begin{bmatrix} 0.44 \\ 0.52 \end{bmatrix}$

- Step 2.1: Step-by-Step Derivation


✓ Step 1: Expand the Function

➤  $f(q) = q^T q = \sum_{i=1}^n q_i^2$

✓ Step 2: Differentiate Each Component

➤ We compute  $\frac{\partial f}{\partial q_i} = \frac{\partial}{\partial q_i} \sum_{j=1}^n q_j^2 = \frac{\partial}{\partial q_i} (q_1^2 + q_2^2 + \dots + q_n^2)$

➤ Since  $q_i^2$  is the only term that depends on  $q_i$ , the derivative is  $\frac{\partial f}{\partial q_i} = 2q_i$



# What Are Vanishing and Exploding Gradients?

## ■ Vector Backpropagation

- Forward and Backward Pass (Step-by-Step)

- Step 2.1: Step-by-Step Derivation

- ✓ Step 1: Expand the Function

- $f(q) = q^T q = \sum_{i=1}^n q_i^2$

- ✓ Step 2: Differentiate Each Component

- We compute  $\frac{\partial f}{\partial q_i} = \frac{\partial}{\partial q_i} \sum_{j=1}^n q_j^2 = \frac{\partial}{\partial q_i} (q_1^2 + q_2^2 + \dots + q_n^2)$


- Since  $q_i^2$  is the only term that depends on  $q_i$ , the derivative is  $\frac{\partial f}{\partial q_i} = 2q_i$

- ✓ Step 3: Write the Gradient as a Vector

- Putting all partial derivatives together →

- ✓ Final Result

- $\frac{\partial}{\partial q} (q^T q) = 2q$



$$\nabla_q f = \begin{bmatrix} \frac{\partial f}{\partial q_1} \\ \frac{\partial f}{\partial q_2} \\ \vdots \\ \frac{\partial f}{\partial q_n} \end{bmatrix} = \begin{bmatrix} 2q_1 \\ 2q_2 \\ \vdots \\ 2q_n \end{bmatrix} = 2q$$

# What Are Vanishing and Exploding Gradients?

## ■ Vector Backpropagation

- Forward and Backward Pass (Step-by-Step)

- Step 2.2: Backward Pass – Gradient w.r.t.  $W$

✓ Using chain rule and matrix calculus:  $\nabla_W f = \frac{\partial f}{\partial W} = 2q \cdot x^T$

✓ Step-by-Step – Derivation of the Gradient

➤ Given Function:  $f(x, W) = \|Wx\|^2 = (Wx)^T(Wx)$

➤ Step1. Gradient w.r.t.  $q$ :

We first rewrite  $f$  in terms of  $q$ :  $f = q^T q$

Then take the gradient of  $f$  with respect to the vector  $q$ :  $\frac{\partial f}{\partial q} = \nabla_q f = 2q$

➤ Step2. Apply the Chain Rule – To get the gradient of  $f$  with respect to  $W$ , we apply the chain rule

$$\frac{\partial f}{\partial W} = \frac{\partial f}{\partial q} \cdot \frac{\partial q}{\partial W} = 2q \cdot \frac{\partial q}{\partial W}$$

- Step 2.1: Backward Pass – Gradient w.r.t.  $q$

✓  $f(x, W) = \|Wx\|^2 = \|q\|^2 = q^T q$

✓  $\frac{\partial f}{\partial q} = \nabla_q f = 2q = \begin{bmatrix} 0.44 \\ 0.52 \end{bmatrix}$

# What Are Vanishing and Exploding Gradients?

## ■ Vector Backpropagation

- Forward and Backward Pass (Step-by-Step)

- Step 2.2: Backward Pass – Gradient w.r.t.  $W$

- ✓ Step-by-Step – Derivation of the Gradient

- Step3. Derivative of  $q = Wx$

➤ Given Function:  $f(x, W) = ||Wx||^2 = (Wx)^T(Wx)$

➤ Step1. Gradient w.r.t.  $q$ :  $\frac{\partial f}{\partial q} = \nabla_q f = \nabla_q q^T q = 2q$

➤ Step2.  $\frac{\partial f}{\partial W} = \frac{\partial f}{\partial q} \cdot \frac{\partial q}{\partial W} = 2q \cdot \frac{\partial q}{\partial W}$

Each component of  $q$  is  $q_i = \sum_{j=1}^d W_{ij} x_j$

So the partial derivative is  $\frac{\partial q_i}{\partial W_{ij}} = x_j = \nabla_W q =$

$$\begin{bmatrix} \frac{\partial q}{\partial W_1} \\ \frac{\partial q}{\partial W_2} \\ \vdots \\ \frac{\partial q}{\partial W_n} \end{bmatrix} = \begin{bmatrix} x_1^T \\ x_2^T \\ \vdots \\ x_n^T \end{bmatrix} \Rightarrow \frac{\partial q}{\partial W} = x^T$$

$$\frac{\partial f}{\partial W_{ij}} = \frac{\partial f}{\partial q_i} \cdot \frac{\partial q_i}{\partial W_{ij}} = 2q_i \cdot \frac{\partial q_i}{\partial W_{ij}} \in \mathbb{R}^{n \times d}$$

$$= [2q_1 \quad 2q_2 \quad \dots \quad 2q_n] \cdot [x_1 \quad x_2 \quad \dots \quad x_d]^T = [2q_1 \quad 2q_2 \quad \dots \quad 2q_n] \cdot$$

$$\Rightarrow \frac{\partial f}{\partial W} = \frac{\partial f}{\partial q} \cdot \frac{\partial q}{\partial W} = 2q \cdot \frac{\partial q}{\partial W} = 2q \cdot x^T$$

$$\begin{bmatrix} x_1^T \\ x_2^T \\ \vdots \\ x_d^T \end{bmatrix} = \begin{bmatrix} 2q_1 x_1 & 2q_1 x_2 & \dots & 2q_1 x_d \\ 2q_2 x_1 & 2q_2 x_2 & \dots & 2q_2 x_d \\ \vdots & \vdots & \ddots & \vdots \\ 2q_n x_1 & 2q_n x_2 & \dots & 2q_n x_d \end{bmatrix}$$

# What Are Vanishing and Exploding Gradients?

## ■ Backpropagation in Fully Connected Layers

### • Gradient Computation Layer by Layer

- Let

- ✓  $\mathbf{z}^l = \mathbf{W}^l \mathbf{a}^{l-1} + \mathbf{b}^l$ : Linear transformation

- ✓  $\mathbf{a}^l = f(\mathbf{z}^l)$ : Activation (e.g., ReLU, sigmoid, tanh)

- Then

- ✓ Backpropagation equation for the gradient signal  $\rightarrow \delta^l = (\mathbf{W}^{l+1})^T \delta^{l+1} \circ f'(\mathbf{z}^l)$

- ✓ Gradient with respect to the weights  $\rightarrow \frac{\partial L}{\partial \mathbf{W}^l} = \delta^l (\mathbf{a}^{l-1})^T$

- ✓ Here,  $\delta^l$  is the **error signal** passed from layer  $l + 1$  back to layer  $l$ , scaled by the derivative of the activation.

- $\mathbf{z}^k \in \mathbb{R}^{d_k}$ : Pre-activation output of layer  $k$  (i.e., before applying the activation function)
- $\mathbf{W}^k$ : Weight matrix of the  $k$ -th layer
- $\mathbf{a}^{k-1}$ : Activation output from the previous layer
- $\mathbf{b}^k$ : Bias vector

Term	Meaning	Mathematical Definition
$\mathbf{z}^k$	Pre-activation of layer $k$	$\mathbf{z}^k = \mathbf{W}^k \mathbf{a}^{k-1} + \mathbf{b}^k$
$\mathbf{a}^k$	Output of layer $k$ (post-activation)	$\mathbf{a}^k = f(\mathbf{z}^k)$
$\delta^k$	Gradient signal at layer $k$	$\delta^k = \frac{\partial L}{\partial \mathbf{z}^k}$
$L$	Total number of layers (depth)	$k = 1, 2, \dots, L$

# What Are Vanishing and Exploding Gradients?

---

- Backpropagation in Fully Connected Layers

- Why Gradients Change Across Layers

- As we apply backpropagation repeatedly across many layers,

$$\delta^l = (W^{l+1})^T \delta^{l+1} \circ f'(\mathbf{z}^l)$$

- We are multiplying gradient vectors by weight matrices and activation derivatives at each step.
    - This means
      - ✓ If those values are **less than 1**, gradients **shrink**.
      - ✓ If they're **greater than 1**, gradients **grow**. This is the root cause of vanishing/exploding gradients.

# What Are Vanishing and Exploding Gradients?

---

## ■ Vanishing Gradient — Root Cause

- Why gradients shrink too much

- If

$$\| W^k \| < 1 \text{ and } | f'(z^k) | < 1$$

- then

$$\| \delta^\ell \| \rightarrow 0 \text{ as } L \rightarrow \infty$$

- This often happens when

- ✓ You use sigmoid or tanh
    - ✓ Weights are poorly initialized
    - ✓ The network is very deep

- Effect

- ✓ Gradients are too small to update early layers
    - ✓ Training becomes very slow or fails entirely

- $z^k \in \mathbb{R}^{d_k}$ : Pre-activation output of layer  $k$  (i.e., before applying the activation function)
  - $W^k$ : Weight matrix of the  $k$ -th layer
  - $a^k = f(z^k)$ : Output of layer  $k$  (post-activation)
  - $f'(z^k)$ : the derivative of the activation function  $f$  evaluated at  $z^k$

# What Are Vanishing and Exploding Gradients?

---

## ■ Theoretical View – Exponential Bounds

- If the Jacobian norms satisfy

- If

$$\left\| \frac{\partial a^k}{\partial a^{k-1}} \right\| \leq \lambda \text{ for all } k$$

- then

$$\left\| \frac{\partial L}{\partial W} \right\| \leq C \cdot \lambda^L$$

- $k$ : the index of layers
  - $L$ : the total number of layers

- Where

- ✓  $C$ : constant based on loss and inputs
    - ✓  $\lambda$ : upper bound on gradient propagation at each layer
    - ✓  $L$ : number of layers

- Interpretation

- ✓ If  $\lambda < 1 \rightarrow$  gradients vanish
    - ✓ If  $\lambda > 1 \rightarrow$  gradients explode
    - ✓ If  $\lambda = 1 \rightarrow$  gradients remain stable (ideal!)

# What Are Vanishing and Exploding Gradients?

---

- Think of **deep backpropagation** as simply applying the **chain rule multiple times**, just like in vector backpropagation—but layer by layer.

- Backpropagation

$$\delta^l = (W^{l+1})^T \delta^{l+1} \circ f'(z^l)$$

- Gradient

$$\frac{\partial L}{\partial W^l} = \delta^l (a^{l-1})^T$$

- This mirrors

$$\frac{\partial f}{\partial W} = \frac{\partial f}{\partial q} \cdot \frac{\partial q}{\partial W}$$

Error Signal

Input Influence



- **From Simple to Deep**

- Start with **vector form** to understand outer products and error  $\times$  input structures.
  - Extend to **deep nets** using recursive formulas and activation derivatives

# Recap – Summary of Motivation

---

- Why Residual Learning Became a Breakthrough
  - Deep networks are **theoretically powerful**, but practically **difficult to optimize**.
  - As depth increases, **plain networks**
    - Encounter vanishing gradients, even with tricks like BatchNorm.
    - Show **increased training error** — a clear sign of optimization issues.
  - **Residual learning provides a solution**
    - Skip connections allow gradients to flow unimpeded.
    - Identity mappings are **easy to learn** with this architecture.
    - Residual blocks make it easier for the optimizer to converge.
  - **Impact**
    - Enables networks with **>100 layers** to be trained effectively.
    - Became the foundation for:
      - ✓ ResNet, ResNeXt
      - ✓ Faster R-CNN (backbone)
      - ✓ Mask R-CNN, and more.

# What is ResNet?

---

- Deep Residual Networks – Introducing ResNet
  - A Breakthrough in Very Deep Neural Networks
    - Proposed by Kaiming He et al. in CVPR 2016
    - Paper: “Deep Residual Learning for Image Recognition”
    - Won ILSVRC 2015 with a top-5 error of 3.57%
    - Designed to enable very deep neural networks (e.g., 152 layers)
  - Problem Addressed
    - As depth increases, plain CNNs suffer from optimization difficulties and even higher training error → the degradation problem
  - ResNet’s solution
    - Introduce skip connections and learn residual functions instead of direct mappings

# Why Is ResNet Needed?

---

- Going Deeper Isn't Always Better
  - Deeper Networks Face Degradation Without the Right Design
    - Before ResNet: VGG, GoogLeNet made progress by going deeper
    - But: *simply adding layers doesn't guarantee better performance!*
    - Example
      - ✓ 56-layer plain CNN performs worse than 20-layer CNN on CIFAR-10
  - This is **not** due to overfitting
    - It's due to optimization issues — **gradient flow weakens**, training becomes harder
  - ResNet reformulates the learning task to address this!

# Why Is ResNet Needed?

## ▪ ResNet Architecture at a Glance

- High-Level Overview

- Main structure

- ✓ Initial Layer

- $7 \times 7$  convolution with stride 2
    - Followed by  $3 \times 3$  max pooling

- ✓ Four Stages of Residual Blocks


- Conv2\_x, Conv3\_x, Conv4\_x, Conv5\_x
    - Spatial resolution halves, channels double each stage

- ✓ Global Average Pooling (GAP)

- Aggregates features into a 1D vector

- ✓ Fully Connected Layer + Softmax

- Produces final classification scores



Example: ResNet-50 has 50 convolutional layers, mostly within residual blocks

# ResNet Variants – Different Depths

---

- ResNet Variants for Different Use Cases

- Deeper ≠ Slower When Designed Right

Model	Depth	Block Type	Parameters
ResNet-18	18	Basic ( $2 \times 3 \times 3$ convs)	~11M
ResNet-34	34	Basic	~21M
ResNet-50	50	Bottleneck	~25M
ResNet-101	101	Bottleneck	~44M
ResNet-152	152	Bottleneck	~60M

- Basic block is used in shallow networks (18/34)
  - Bottleneck block enables efficient deeper networks (50+)

# Key Components of ResNet

---

- Components That Make ResNet Work
  - What's Inside the Architecture?

Component	Role in the Network
Residual Block	Learns a residual function: $F(x)$ , and outputs $F(x) + x$
Skip Connection	Directly passes input through identity or projection mapping
Bottleneck Block	Reduces → processes → restores dimensions via $1 \times 1$ / $3 \times 3$ / $1 \times 1$ convs
BatchNorm + ReLU	Applied after each convolution layer


- These design choices ensure effective training of **very deep** models

# The first Key Components of ResNet

## ■ Introduction to the Residual Block

### • What is a Residual Block?

- A **Residual Block** is the core building unit of ResNet (Residual Network).
- It was designed to address the **degradation problem** in deep networks, where adding more layers leads to **worse performance**.
- Instead of directly learning the mapping  $H(x)$ , it learns a **residual function**  $F(x) = H(x) - x$ , and then reconstructs  $H(x)$  as  $H(x) = F(x) + x$



### • Basic Structure of a Residual Block

- The standard formulation:  $y = F(x, \{W_i\}) + x$ 
  - ✓ Where  $x$  = input feature map,  $F(x, \{W_i\})$  = residual function (usually 2–3 convolutional layers with weights  $\{W_i\}$ ),  $y$  = output feature map
- Key Concept
  - ✓ Instead of learning the full transformation, we let the stacked layers only learn the "difference"  $F(x)$  from the identity.

# The first Key Components of ResNet

---

## ■ Introduction to the Residual Block

- Why Add the Input Back?

- Key Idea

- ✓ Residual learning reformulates the desired mapping  $H(x)$  into:

$$H(x) = F(x) + x$$

- The network only needs to learn the residual function  $F(x)$ , not the full transformation  $H(x)$

- Why is this helpful? (Advantages)

- ✓ 1. Focus on what's new

- The layer learns only the part that needs to change (i.e., the residual difference  $H(x) - x$ ). This makes optimization easier, especially in deep networks.

- **Naming Origin:** The term “ResNet” comes from the word **residual**

- The network learns to minimize the residual  $H(x) - x$ . Hence, **Residual Network**.

# The first Key Components of ResNet

---

## ■ Introduction to the Residual Block

- Why Add the Input Back?

- Why is this helpful? (Advantages)

- ✓ 2. Convergence behavior

- As the depth increases and the network is well-trained, the input  $x$  becomes increasingly close to the output  $H(x)$ , so the residual  $F(x) \rightarrow 0$ .

- This stabilizes training and encourages minimal necessary updates.

- ✓ 3. Implementation simplicity

- (1) No major change in architecture is required.

- (2) Simply add a **shortcut connection** from input to output.

- (3) No extra parameters are introduced since  $x$  is reused.

- (4) Aside from the final addition operation, the shortcut adds almost no computational cost.

# The Second Key Components of ResNet

---

- Introduction to Skip (i.e., Shortcut) Connections

- Types of Shortcut (i.e., Skip) Connections

- When input and output dimensions do not match,


- ✓ 1. Identity Mapping with Padding

$$y = F(x) + \text{pad}(x)$$

- ✓ 2. Projection Mapping

$$y = F(x) + W_s x$$

→ where  $W_s$  is a 1x1 convolution used to match dimensions



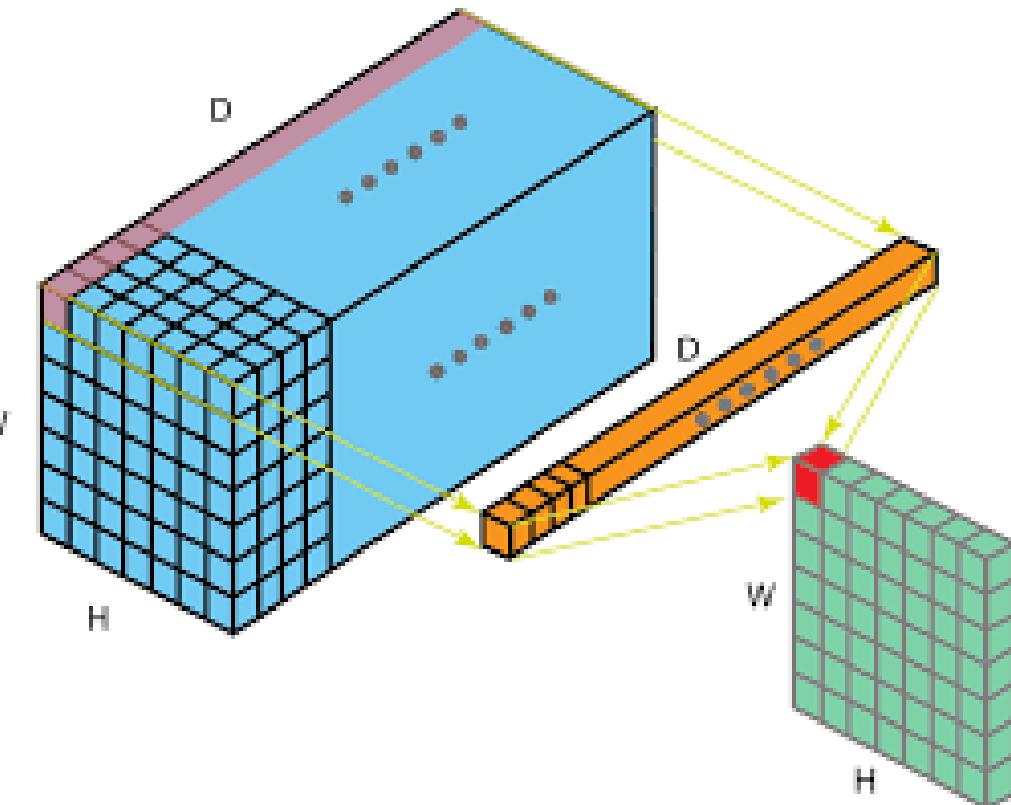
# The Second Key Components of ResNet

- Introduction to Skip (i.e., Shortcut) Connections
  - What is  $1 \times 1$  Convolution?

- A  $1 \times 1$  convolution is a convolutional operation where the filter (kernel) has a spatial dimension of 1 (i.e.,  $1 \times 1$ ), but spans the full depth (channel) of the input.

- Key Purposes of  $1 \times 1$  Convolution

Purpose	Explanation
Channel Adjustment	Reduces or expands the number of channels without changing the spatial size ( $W \times H$ )
Computational Efficiency	Reduces the number of parameters and operations, especially before expensive convolutions (e.g., $3 \times 3$ , $5 \times 5$ )
Adding Nonlinearity	When followed by an activation (e.g., ReLU), it increases model expressiveness
Bottleneck Design Enabler	Used in ResNet's bottleneck blocks ( $1 \times 1 \rightarrow 3 \times 3 \rightarrow 1 \times 1$ ) for compression-expansion



# The Second Key Components of ResNet

## ■ Introduction to Skip (i.e., Shortcut) Connections


### • What is $1 \times 1$ Convolution?

#### ○ Example

✓ Scenario:  $5 \times 5$  Conv on  $28 \times 28 \times 192 \rightarrow 28 \times 28 \times 32$

#### ➤ Direct $5 \times 5$ convolution

- $28 \times 28 \times 32 \times 5 \times 5 \times 192 \approx 120M$  operations



✓ With  $1 \times 1$  Conv Compression


#### ➤ Step 1. Reduce channels to 16 using $1 \times 1$ Conv

- $28 \times 28 \times 16 \times 1 \times 1 \times 192 \approx 2.4M$

#### ➤ Step 2. Apply $5 \times 5$

- $28 \times 28 \times 32 \times 5 \times 5 \times 16 \approx 10M$

→ Total  $\approx 12.4M$  operations ( $\sim 10 \times$  smaller)



# The Third Key Components of ResNet

## ■ Introduction to the Bottleneck Block

### • Residual Block Variants in ResNet Implementation

- Basic Residual Block (used in ResNet-18, 34)

- ✓ **Shortcut Structure:** two  $3 \times 3$  convolutions

- ✓ **Shortcut type**

- **Mostly Identity Mapping** is used when input and output dimensions are the same.

- **BUT! Projection Mapping** (via  $1 \times 1$  convolution) is used when needed
    - The number of channels increases
    - The spatial resolution changes (e.g., stride = 2)


- Bottleneck Block (used in ResNet-50, 101, 152)

- ✓ **Shortcut Structure:** three layers consisting of  $1 \times 1 \rightarrow 3 \times 3 \rightarrow 1 \times 1$

- ✓ **Shortcut type**

- **Mostly Projection Mapping** is used in most blocks, especially because
    - The bottleneck design frequently changes dimensions
    - Downsampling occurs regularly (stride = 2)

- ✓ Reduces and restores dimensions to lower computation while preserving representational power.



# Key Components of ResNet

---


- Experimental Results – Does ResNet Really Work?
  - Purpose
    - To evaluate whether **shortcut connections (residual blocks)** truly improve learning in deep networks.
  - Experimental Setup
    - Dataset: ImageNet
    - Models Compared: (1) Plain-18 vs. Plain-34, (2) ResNet-18 vs. ResNet-34

# Key Components of ResNet

## ■ Experimental Results – Does ResNet Really Work?

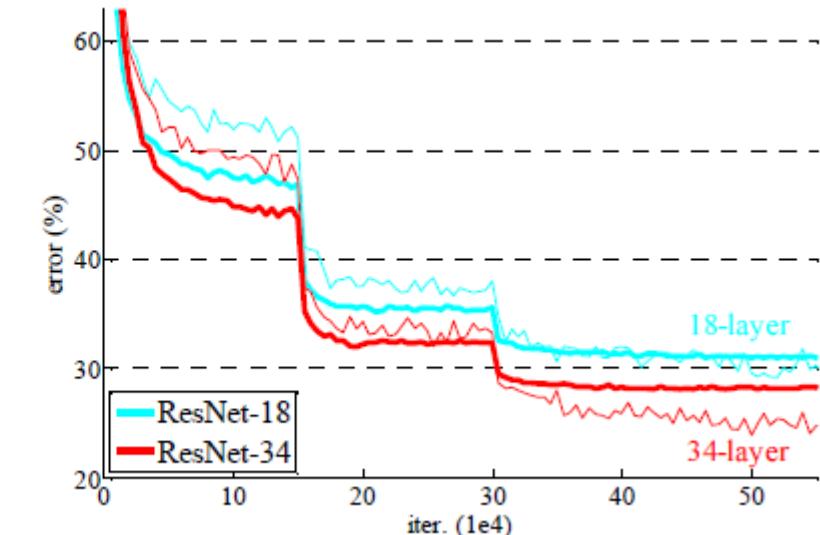
- Observation

- Left: Plain Network



- ✓ As the network **gets deeper (34-layer)**, performance **worsens**
  - ✓ The **Plain-34** model has **higher training and validation error** than Plain-18
  - ✓ This is a classic case of the **degradation problem**

- Right: Residual Network



- ✓ **ResNet-34** outperforms ResNet-18 across all iterations
  - ✓ **Deeper = better performance**, as expected
  - ✓ Residual learning via shortcuts enables effective optimization of deeper network

# Key Components of ResNet

---

## ■ Why Residual Blocks Prevent Vanishing Gradients (i.e., Better Performance)

### • Forward Pass (Residual Formulation)

- Let the activation function be ReLU and the residual mapping be defined as:

$$\checkmark \mathbf{x}_l + \mathbf{1} = \mathbf{f}(\mathbf{y}_l) = \mathbf{f}(\mathbf{x}_l + \mathbf{F}(\mathbf{x}_l, \mathbf{W}_l)) \approx \mathbf{x}_l + \mathbf{F}(\mathbf{x}_l, \mathbf{W}_l)$$

➤ Assuming  $\mathbf{f}$  is identity for simplicity.

➤ Each term in the equation represents

- $\mathbf{x}_l$ : The input to the  $l$ -th layer (or the output from the previous layer)
- $\mathbf{W}_l$ : The learnable weights of the  $l$ -th layer
- $\mathbf{F}(\mathbf{x}_l, \mathbf{W}_l)$ : The residual function applied in the  $l$ -th layer
- $\mathbf{f}$ : The activation function (e.g., ReLU)
- $\mathbf{x}_l + \mathbf{1}$ : The output from the  $l$ -th layer (input to the next layer)

- This leads to a general form across  $L$  layers:

$$\checkmark \mathbf{x}_L = \mathbf{x}_1 + \sum_{i=1}^{L-1} \mathbf{F}(\mathbf{x}_i, \mathbf{W}_i)$$

# Key Components of ResNet

---

- Why Residual Blocks Prevent Vanishing Gradients (i.e., Better Performance)

- Backward Pass (Gradient Computation)

- Let the loss be  $\mathcal{L}$ , then by the chain rule: (where the symbol  $\mathcal{L}$  represents the loss function)

$$\frac{\partial \mathcal{L}}{\partial x_1} = \frac{\partial \mathcal{L}}{\partial x_L} \left( 1 + \frac{\partial}{\partial x_1} \sum_{i=1}^{L-1} F(x_i, W_i) \right)$$

- The gradient splits into two terms

- ✓ (1) Identity Path (i.e., Residual Connection)

$$\frac{\partial \mathcal{L}}{\partial x_L} \cdot 1$$

- This term bypasses all weight layers → gradient flows unimpeded

- ✓ (2) Residual Path

$$\frac{\partial \mathcal{L}}{\partial x_L} \cdot \frac{\partial}{\partial x_1} \sum_{i=1}^{L-1} F(x_i, W_i)$$

- This term propagates through weights → may diminish or explode

# Key Components of ResNet

---

- Why Residual Blocks Prevent Vanishing Gradients (i.e., Better Performance)

- Why This Matters

- Residual connection ensures at least one term in the gradient is preserved (term 1)
    - Unlike plain networks, where gradients might vanish as depth increases, ResNet ensures

$$\frac{\partial L}{\partial x_L} \text{ (i.e., term 1)} \neq 0$$

- ✓ Even if part of the gradient shrinks, the identity shortcut preserves signal, avoiding full gradient collapse
    - Practical Implication

$$\frac{\partial L}{\partial x_1} = \frac{\partial L}{\partial x_L} \left( 1 + \frac{\partial}{\partial x_1} \sum_{i=1}^{L-1} F(x_i, W_i) \right)$$

- ✓ In order to become gradient vanishing in real training (mini-batch SGD), the gradients in the first term become -1. BUT, the gradients vary and often don't exactly become -1.

- ✓ Still, residual design avoids full cancellation

- ✓ Thus, vanishing gradients are prevented and deep networks can be trained effectively

# The Fourth Key Components of ResNet

---

- What is Batch Normalization?

- **Batch Normalization (BN) – Making Deep Networks Easier to Train**

- BN is a technique to normalize the input of each layer so that its distribution remains stable throughout training.

- **Why Do We Need It?**

- Deep networks often suffer from

- ✓ **Gradient Vanishing/Exploding:** Especially with deep networks and activation functions like sigmoid/tanh.

- ✓ **Internal Covariate Shift:** Input distribution to each layer changes during training.

# The Fourth Key Components of ResNet

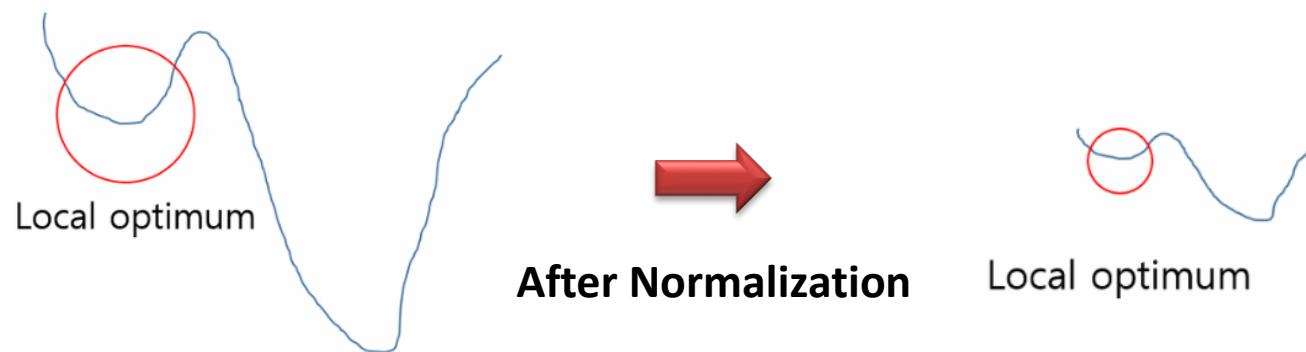
---

## ■ What is Batch Normalization?

- Why Do We Normalize?

- Intuition Behind Normalization

- ✓ Normalization helps the model **learn faster** and **avoid getting stuck** in bad solutions.
    - ✓ During optimization, models can get trapped in **local optima** if the loss surface is highly irregular or poorly scaled.
    - ✓ By normalizing, we **reshape the loss surface**, making it smoother and easier to traverse.



- Without normalization: Optimizer may get stuck in **sharp local minima**
  - With normalization: Loss surface becomes **flatter**, reducing the chance of being stuck in suboptimal regions

*Goal: Help the optimizer reach the **global optimum** by simplifying the search space.*

# The Fourth Key Components of ResNet

---

## ■ Batch Normalization

- Motivation – 1. The Gradient Vanishing / Exploding Problem

- Why Deep Networks Are Hard to Train

- ✓ Training deep neural networks is challenging due to the **Gradient Vanishing** or **Gradient Exploding** problem.
    - ✓ Gradients become **too small** (vanish) or **too large** (explode) as they propagate backward through many layers.
    - ✓ This makes it hard to update parameters effectively, resulting in
      - Poor convergence
      - High error rate
      - Inability to train deep models

- What Causes It?

- ✓ Common activation functions like **sigmoid** or **tanh** squash the output to narrow ranges (e.g., [0, 1]), especially in deep layers.
    - ✓ This results in
      - Large variations in input leading to tiny changes in output
      - Gradients becoming extremely small
    - ✓ A typical solution is to use **ReLU** (Rectified Linear Unit), which preserves gradients better in deep models.

# The Fourth Key Components of ResNet

---

## ■ Batch Normalization

- Motivation – 1. The Gradient Vanishing / Exploding Problem

- Workarounds (but not ideal)

- ✓ Change activation: Use ReLU instead of sigmoid/tanh
    - ✓ Careful weight initialization: Helps maintain gradient scale
    - ✓ Small learning rate: Prevents gradients from exploding

*However, these are patches, not fundamental fixes.*

- Core Idea: Stabilize the Entire Learning Process

- ✓ This leads to the introduction of **Batch Normalization (BN)**

- ✓ BN normalizes inputs to each layer, so that

- The distribution of inputs stays consistent throughout training
    - The network becomes **less sensitive** to initialization
    - Training becomes **faster** and **more stable**

- ✓ BN effectively prevents both gradient vanishing and exploding by keeping activations within a healthy range.

# The Fourth Key Components of ResNet

---

- Batch Normalization

- Motivation – 2. Internal Covariate Shift

- What Is Internal Covariate Shift?

- ✓ During training, the distribution of inputs to each layer **keeps changing** as the parameters of previous layers update.
      - ✓ This shift slows down training and makes it harder for the network to converge.

- **Covariate Shift:** When the input distribution to a model (or a layer) changes due to external factors (e.g., changing dataset).

- **Internal Covariate Shift:** When the input distribution to a layer **changes during training** due to updates in earlier layers.

# The Fourth Key Components of ResNet

---

- Batch Normalization

- Motivation – 2. Internal Covariate Shift

- What Is Internal Covariate Shift?

- ✓ During training, the distribution of inputs to each layer **keeps changing** as the parameters of previous layers update.
      - ✓ This shift slows down training and makes it harder for the network to converge.

- **Covariate Shift:** When the input distribution to a model (or a layer) changes due to external factors (e.g., changing dataset).

- **Internal Covariate Shift:** When the input distribution to a layer **changes during training** due to updates in earlier layers.

# The Fourth Key Components of ResNet



## Batch Normalization

- Motivation – 2. Internal Covariate Shift

- Why It Matters

- ✓ Neural networks are deep compositions of functions.
    - ✓ As each layer is updated during training, the input distribution to the next layer **shifts**, making it difficult for the next layer to adapt.
    - ✓ This leads to **slower convergence** and **unstable training**.

- Empirical Evidence



Batch Normalization reduces internal covariate shift  
→ resulting in **faster training, better stability, and higher accuracy**

- ✓ The graph (right) shows how **BN stabilizes input distributions** across layers.

- ✓ With BatchNorm: input distributions become **more consistent**, helping the network learn faster and more reliably.

# The Fourth Key Components of ResNet

---

## ■ Batch Normalization

### • How BN Works

- Let  $x = \{x_1, x_2, \dots, x_m\}$  be a mini-batch of inputs

#### ✓ 1. Compute mean and variance:

$$\mu_B = \frac{1}{m} \sum_{i=1}^m x_i, \quad \sigma_B^2 = \frac{1}{m} \sum_{i=1}^m (x_i - \mu_B)^2$$

#### ✓ 2. Normalize the inputs:

$$\hat{x}_i = \frac{x_i - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}}$$

#### ✓ 3. Apply scale and shift:

$$y_i = \gamma \hat{x}_i + \beta$$

➤ Where

$\epsilon$ : small constant for numerical stability

$\gamma$ : learnable **scale** parameter

$\beta$ : learnable **shift** parameter

# The Fourth Key Components of ResNet

## ■ Batch Normalization

- Backpropagation in Batch Normalization
  - Learnable Parameters
    - ✓ BN introduces two learnable parameters
      - $\gamma$  (gamma): controls the scale,  $\beta$  (beta): controls the shift

### 1. Compute mean and variance:

$$\mu_B = \frac{1}{m} \sum_{i=1}^m x_i, \quad \sigma_B^2 = \frac{1}{m} \sum_{i=1}^m (x_i - \mu_B)^2$$

### 2. Normalize the inputs:

$$\hat{x}_i = \frac{x_i - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}}$$

### 3. Apply scale and shift:

$$y_i = \gamma \hat{x}_i + \beta$$

### ○ Chain Rule Derivatives

- ✓ Given loss  $L$ , we backpropagate through BN using the following gradients. The normalized input  $\hat{x}_i$  consists of  $\mu_B$  and  $\sigma_B^2$ .

$$\frac{\partial L}{\partial \hat{x}_i} = \frac{\partial L}{\partial y_i} \cdot \gamma$$

➤ Which means that we need to compute the partial derivative of the loss  $L$  with respect to  $\mu_B$  and  $\sigma_B^2$ .

### ○ Why?

- ✓ In Batch Normalization, we normalize the input using the batch mean  $\mu_B$  and variance  $\sigma_B^2$ , then scale and shift it using learnable parameters  $\gamma$  and  $\beta$ .
- ✓ To update the model through backpropagation, we must compute the gradient of the loss with respect to these internal variables, including the variance.

# The Fourth Key Components of ResNet

## ■ Batch Normalization

- Backpropagation in Batch Normalization

- Chain Rule Derivatives

- ✓ Step 1. Compute the partial derivative of the loss  $L$  with respect to the variance  $\sigma_B^2$  of a mini-batch.

$$\begin{aligned}\frac{\partial L}{\partial \sigma_B^2} &= \sum_{i=1}^m \frac{\partial L}{\partial \hat{x}_i} \cdot \frac{\partial \hat{x}_i}{\partial \sigma_B^2} \\ &= \sum_{i=1}^m \frac{\partial L}{\partial \hat{x}_i} \cdot (x_i - \mu_B) \cdot \frac{\partial}{\partial \sigma_B^2} (\sigma_B^2 + \epsilon)^{-\frac{1}{2}} \\ &= \sum_{i=1}^m \frac{\partial L}{\partial \hat{x}_i} \cdot (x_i - \mu_B) \cdot \left(-\frac{1}{2}\right) (\sigma_B^2 + \epsilon)^{-\frac{3}{2}}\end{aligned}$$

**1. Compute mean and variance:**

$$\mu_B = \frac{1}{m} \sum_{i=1}^m x_i, \quad \sigma_B^2 = \frac{1}{m} \sum_{i=1}^m (x_i - \mu_B)^2$$

**2. Normalize the inputs:**

$$\hat{x}_i = \frac{x_i - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}}$$

**3. Apply scale and shift:**

$$y_i = \gamma \hat{x}_i + \beta$$

# The Fourth Key Components of ResNet

## Batch Normalization

- Backpropagation in Batch Normalization
  - Chain Rule Derivatives
    - ✓ Step 2. Compute the partial derivative of the loss  $L$  with respect to the variance  $\mu_B$  of a mini-batch.

$$\frac{\partial L}{\partial \mu_B} = \sum_{i=1}^m \frac{\partial L}{\partial \hat{x}_i} \cdot \frac{\partial \hat{x}_i}{\partial \mu_B} + \frac{\partial L}{\partial \sigma_B^2} \cdot \frac{\partial \sigma_B^2}{\partial \mu_B} = \sum_{i=1}^m \frac{\partial L}{\partial \hat{x}_i} \cdot (\sigma_B^2 + \epsilon)^{-\frac{1}{2}} \cdot \frac{\partial}{\partial \mu_B} (x_i - \mu_B) + \frac{\partial L}{\partial \sigma_B^2} \cdot \frac{\partial \sigma_B^2}{\partial \mu_B}$$

$$= \sum_{i=1}^m \frac{\partial L}{\partial \hat{x}_i} \cdot (\sigma_B^2 + \epsilon)^{-\frac{1}{2}} \cdot -1 + \frac{\partial L}{\partial \sigma_B^2} \cdot \frac{\partial \sigma_B^2}{\partial \mu_B}$$

$$= \left( \sum_{i=1}^m \frac{\partial L}{\partial \hat{x}_i} \cdot (\sigma_B^2 + \epsilon)^{-\frac{1}{2}} \cdot -1 \right) + \frac{\partial L}{\partial \sigma_B^2} \cdot \frac{\partial}{\partial \mu_B} \left( \frac{1}{m} \sum_{i=1}^m (x_i - \mu_B)^2 \right)$$

$$= \left( \sum_{i=1}^m \frac{\partial L}{\partial \hat{x}_i} \cdot (\sigma_B^2 + \epsilon)^{-\frac{1}{2}} \cdot -1 \right) + \frac{\partial L}{\partial \sigma_B^2} \cdot \left( \frac{-2}{m} \sum_{i=1}^m (x_i - \mu_B) \right)$$

### 1. Compute mean and variance:

$$\mu_B = \frac{1}{m} \sum_{i=1}^m x_i, \quad \sigma_B^2 = \frac{1}{m} \sum_{i=1}^m (x_i - \mu_B)^2$$

### 2. Normalize the inputs:

$$\hat{x}_i = \frac{x_i - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}}$$

### 3. Apply scale and shift:

$$y_i = \gamma \hat{x}_i + \beta$$

# The Fourth Key Components of ResNet

## ■ Batch Normalization

- Backpropagation in Batch Normalization

- Chain Rule Derivatives

➤ Step 3. Compute the partial derivative of the loss  $L$  with respect to the input  $x_i$  of a mini-batch.

- (1) Direct path through  $\hat{x}_i$

### 1. Compute mean and variance:

$$\mu_B = \frac{1}{m} \sum_{i=1}^m x_i, \quad \sigma_B^2 = \frac{1}{m} \sum_{i=1}^m (x_i - \mu_B)^2$$

### 2. Normalize the inputs:

$$\hat{x}_i = \frac{x_i - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}}$$

### 3. Apply scale and shift:

$$y_i = \gamma \hat{x}_i + \beta$$

$$\frac{\partial L}{\partial x_i} = \frac{\partial L}{\partial \hat{x}_i} \cdot \frac{\partial \hat{x}_i}{\partial x_i} = \frac{\partial L}{\partial \hat{x}_i} \cdot \frac{\partial}{\partial x_i} \left( \frac{x_i - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}} \right) = \frac{\partial L}{\partial \hat{x}_i} \cdot (\sigma_B^2 + \epsilon)^{-\frac{1}{2}}$$

- (2) Direct path through  $\sigma_B^2$

$$\frac{\partial L}{\partial x_i} = \frac{\partial L}{\partial \sigma_B^2} \cdot \frac{\partial \sigma_B^2}{\partial x_i} = \frac{\partial L}{\partial \sigma_B^2} \cdot \frac{\partial}{\partial x_i} \left( \frac{1}{m} \sum_{i=1}^m (x_i - \mu_B)^2 \right) = \frac{\partial L}{\partial \sigma_B^2} \cdot \frac{2(x_i - \mu_B)}{m}$$

# The Fourth Key Components of ResNet

## ■ Batch Normalization

- Backpropagation in Batch Normalization

- Chain Rule Derivatives

➤ Step 3. Compute the partial derivative of the loss  $L$  with respect to the input  $x_i$  of a mini-batch.

- (3) Direct path through  $\mu_B$

### 1. Compute mean and variance:

$$\mu_B = \frac{1}{m} \sum_{i=1}^m x_i, \quad \sigma_B^2 = \frac{1}{m} \sum_{i=1}^m (x_i - \mu_B)^2$$

### 2. Normalize the inputs:

$$\hat{x}_i = \frac{x_i - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}}$$

### 3. Apply scale and shift:

$$y_i = \gamma \hat{x}_i + \beta$$

$$\frac{\partial L}{\partial x_i} = \frac{\partial L}{\partial \mu_B} \cdot \frac{\partial \mu_B}{\partial x_i} = \frac{\partial L}{\partial \mu_B} \cdot \frac{\partial}{\partial x_i} \left( \frac{1}{m} \sum_{i=1}^m x_i \right) = \frac{\partial L}{\partial \mu_B} \cdot \frac{1}{m}$$

- (4) Final Expression Combining All Three Terms

$$\frac{\partial L}{\partial x_i} = \frac{\partial L}{\partial \hat{x}_i} \cdot (\sigma_B^2 + \epsilon)^{-\frac{1}{2}} + \frac{\partial L}{\partial \sigma_B^2} \cdot \frac{2(x_i - \mu_B)}{m} + \frac{\partial L}{\partial \mu_B} \cdot \frac{1}{m}$$

# The Fourth Key Components of ResNet

---

## ■ Batch Normalization

- Backpropagation in Batch Normalization

- Chain Rule Derivatives

➤ Step 4. Complete the final expression combining the results from Step1 and Step2.

- Result from Step 1

$$\frac{\partial L}{\partial \sigma_B^2} = \sum_{i=1}^m \frac{\partial L}{\partial \hat{x}_i} \cdot (x_i - \mu_B) \cdot \left(-\frac{1}{2}\right) (\sigma_B^2 + \epsilon)^{-\frac{3}{2}}$$

- Result from Step 2

$$\frac{\partial L}{\partial \mu_B} = \sum_{i=1}^m \frac{\partial L}{\partial \hat{x}_i} \cdot (\sigma_B^2 + \epsilon)^{-\frac{1}{2}} + \frac{\partial L}{\partial \sigma_B^2} \cdot \left(\frac{2}{m} \sum_{i=1}^m (x_i - \mu_B)\right)$$

- Final Expression

$$\frac{\partial L}{\partial x_i} = \frac{\partial L}{\partial \hat{x}_i} \cdot (\sigma_B^2 + \epsilon)^{-\frac{1}{2}} + \frac{\partial L}{\partial \sigma_B^2} \cdot \frac{2(x_i - \mu_B)}{m} + \frac{\partial L}{\partial \mu_B} \cdot \frac{1}{m}$$

# The Fourth Key Components of ResNet

## ■ Batch Normalization

- Backpropagation in Batch Normalization

- Chain Rule Derivatives

➤ Step 4. Complete the final expression combining the results from Step1 and Step2.

- Final Expression

$$\begin{aligned}\frac{\partial L}{\partial x_i} &= \frac{\partial L}{\partial \hat{x}_i} \cdot (\sigma_B^2 + \epsilon)^{-\frac{1}{2}} + \left( \sum_{i=1}^m \frac{\partial L}{\partial \hat{x}_i} \cdot (x_i - \mu_B) \cdot \left( -\frac{1}{2} \right) (\sigma_B^2 + \epsilon)^{-\frac{3}{2}} \right) \cdot \frac{2(x_i - \mu_B)}{m} + \sum_{i=1}^m \frac{\partial L}{\partial \hat{x}_i} \cdot (\sigma_B^2 + \epsilon)^{-\frac{1}{2}} \\ &\quad + \left( \sum_{i=1}^m \frac{\partial L}{\partial \hat{x}_i} \cdot (x_i - \mu_B) \cdot \left( -\frac{1}{2} \right) (\sigma_B^2 + \epsilon)^{-\frac{3}{2}} \right) \cdot \left( \frac{2}{m} \sum_{i=1}^m (x_i - \mu_B) \right) \cdot \frac{1}{m}\end{aligned}$$

The final expression can be written entirely in terms of  $\frac{\partial L}{\partial \hat{x}_i}$ . It tells us how much the normalized input  $\hat{x}_i$  contributes to the final loss.

$$y_i = \gamma \hat{x}_i + \beta \Rightarrow \frac{\partial L}{\partial \hat{x}_i} = \frac{\partial L}{\partial y_i} \cdot \frac{\partial y_i}{\partial \hat{x}_i} = \frac{\partial L}{\partial y_i} \cdot \gamma$$

### 3. Apply scale and shift:

$$y_i = \gamma \hat{x}_i + \beta$$

# The Fourth Key Components of ResNet

---

## ■ Batch Normalization

- Backpropagation in Batch Normalization

- Chain Rule Derivatives

✓ The final expression can be written entirely in terms of  $\frac{\partial L}{\partial \hat{x}_i}$ . It tells us how much the normalized input  $\hat{x}_i$  contributes to the final loss.

$$y_i = \gamma \hat{x}_i + \beta \Rightarrow \frac{\partial L}{\partial \hat{x}_i} = \frac{\partial L}{\partial y_i} \cdot \frac{\partial y_i}{\partial \hat{x}_i} = \frac{\partial L}{\partial y_i} \cdot \gamma$$

- Summary of the Full Process

- Forward Pass

$$\hat{x}_i = \frac{x_i - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}} \quad y_i = \gamma \hat{x}_i + \beta$$

- Backward Pass Goal

Compute  $\frac{\partial L}{\partial y_i}$

We don't "observe"  $\frac{\partial L}{\partial y_i}$  directly. BUT! we can compute it from  $\frac{\partial L}{\partial y_i}$ , which is what backprop actually gives us by using the final expression with  $\frac{\partial L}{\partial \hat{x}_i}$ .

# The Fourth Key Components of ResNet

## ■ Batch Normalization

### • How BN is used in ResNet


- In the original ResNet architecture (shown on the left), **Batch Normalization** is applied **after each convolutional layer**, before the ReLU activation.

- The sequence inside a Residual Block is

- ✓ 1. Conv (i.e., weight in the figure) → BN → ReLU
- ✓ 2. Conv → BN
- ✓ 3. Addition
- ✓ 4. ReLU

- Pre-activation ResNet (Proposed Improvement)

- ✓ In deeper networks (like ResNet-164), researchers observed that training becomes harder.
- ✓ They proposed a **pre-activation variant**, where **BN and ReLU come before convolution**.
- ✓ This leads to smoother optimization and better gradient flow.




# The Fourth Key Components of ResNet

## Batch Normalization

- Performance Comparison (ResNet-164)
  - Yellow: Original post-activation ResNet
  - Blue: Proposed pre-activation ResNet
  - Training Loss
    - ✓ Blue curve (pre-activation) shows worse training loss (highlighted in red).
  - Test Error
    - ✓ Blue curve generalizes better with lower test error (highlighted in black).

Indicates less overfitting.  
Leads to improved generalization performance.



# Summary – Key Ideas Behind ResNet

---

## ■ Why Do We Need ResNet?

- Deeper networks offer stronger representational power.
- But **deeper ≠ better**, due to training difficulties.
- Challenges like **Vanishing/Exploding Gradients** and the **Degradation Problem** hinder performance as depth increases.

## ■ Key Components of ResNet

Component	Role in the Network
Residual Block	Learns a residual function: $F(x)$ , and outputs $F(x) + x$
Skip Connection	Directly passes input through identity or projection mapping
Bottleneck Block	Reduces → processes → restores dimensions via $1 \times 1$ / $3 \times 3$ / $1 \times 1$ convs
BatchNorm + ReLU	Applied after each convolution layer

# Summary – Key Ideas Behind ResNet

---

## ■ Role of Batch Normalization

- Mitigates gradient vanishing/exploding
- Addresses **Internal Covariate Shift** → keeps input distribution stable across layers
- Used in the sequence: Conv → BN → ReLU  
(Later, **Pre-activation**: BN → ReLU → Conv is proposed and found to perform better)

## ■ Takeaways

- **Residual learning** enables the training of very deep networks
- ResNet architecture is not just about going deeper, but about training **effectively and reliably**
- Still widely used as a **backbone** in modern deep learning models