

# **ComputerVision**

## **Week7**

2025-2

Mobile Systems Engineering  
Dankook University

# From Classification to Segmentation

- Understanding different vision tasks

- 1. Classification – Translation Invariance Task

- “*What is in the image?*”
    - Single label for the entire image.

- 2. Classification + Localization – Translation Variance Task

- “*What is it, and where is it?*”
    - Predicts class label and bounding box for **one object**.

- 3. Object Detection – Translation Variance Task

- “*What objects are there, and where are they?*”
    - Detects **multiple objects**, each with a bounding box and label.

- 4. Instance Segmentation – Translation Variance Task

- “*Which pixels belong to each object?*”
    - Pixel-level masks that separate even objects of the same class.

Classification      Classification + Localization



Single object

Object Detection

Instance Segmentation




Multiple objects

# From Classification to Segmentation

- Segmentation Types: Semantic vs. Instance
  - Segmentation Types in Computer Vision

- 1. Semantic Segmentation




Segmented

**Semantic Labels**

- 1: Person
- 2: Purse
- 3: Plants/Grass
- 4: Sidewalk
- 5: Building/Structures


3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	5	5	5	5	5	
3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	5	5	5	5	5
3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	5	5	5	5	5
3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	5	5	5	5	5
3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	5	5	5	5	5
3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	3	5	5	5	5	5
5	5	3	3	3	3	3	3	3	3	3	3	1	1	3	3	3	3	3	3	3	3	5	5	5	5	5
4	4	3	4	1	1	1	1	1	1	1	1	1	1	1	1	1	1	4	4	4	4	5	5	5	5	5
4	4	3	4	1	1	1	1	1	1	1	1	1	1	1	1	1	1	4	4	4	4	4	5	5	5	5
4	4	4	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	4	4	4	4	4	4	4	4	4
3	3	3	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	4	4	4	4	4	4	4	4	4
3	3	3	1	2	2	1	1	1	1	1	1	1	1	1	1	1	1	4	4	4	4	4	4	4	4	4
3	3	3	1	2	2	1	1	1	1	1	1	1	1	1	1	1	1	4	4	4	4	4	4	4	4	4

- ✓ Assigns a **class label** to each **pixel** in the image.
- ✓ All objects of the same class share the **same label**.
  - Example: All cars in the image → labeled as “Car” (same color in mask).
- ✓ Limitation: Cannot distinguish **different instances** of the same class.



# From Classification to Segmentation

- Segmentation Types: Semantic vs. Instance
  - Segmentation Types in Computer Vision
    - 2. Instance Segmentation



- ✓ Assigns **class + instance ID** to each pixel.
- ✓ Objects with the same class but different instances have **separate masks**.
  - Example: Two cars → both “Car” class, but with different colors/masks (i.e., different id – car1 and car2).
- ✓ Combines object detection’s **localization** with segmentation’s **pixel-level detail**.


# From Classification to Segmentation

---

- Segmentation Types: Semantic vs. Instance

- Segmentation Types in Computer Vision

- Key Differences



Aspect	Semantic Segmentation	Instance Segmentation
<b>Output Label</b>	Class ID per pixel	Class ID + Instance ID per pixel
<b>Same-class objects</b>	Same label	Different labels (instance-aware)
<b>Example Use Cases</b>	Road scene parsing, medical imaging	Autonomous driving, object counting

# From Classification to Segmentation

## ■ Why Standard CNN Architectures for Classification Fail at Segmentation

### • 1. How CNNs Learn Features

- Early Layers – Low-Level Features

- ✓ Learn basic visual patterns
  - Edges, corners, simple colors.

- ✓ Capture fine-grained spatial information.

- ✓ Examples


- Layer 1 detects straight lines, edges at different orientations.

- Middle Layers – Mid-Level Features

- ✓ Combine low-level patterns into **textures** and **object parts**.

- ✓ Example

- Layer 3 detects a dog's ear, a car wheel, or repeated surface patterns.



# From Classification to Segmentation

## ■ Why Standard CNN Architectures for Classification Fail at Segmentation

### • 1. How CNNs Learn Features


- Deep Layers – High-Level Features

- ✓ Detect entire objects or complex shapes.

- ✓ Focus on semantic meaning rather than precise location.


- ✓ Example

- Layer 5 activates for a whole dog or bicycle.



# From Classification to Segmentation

- Why Standard CNN Architectures for Classification Fail at Segmentation
  - 2. Important Observations for Segmentation



- **Standard CNN is “Translation Invariance”!**


- ✓ 1. As we move deeper, **spatial resolution decreases** due to **pooling**, **stride**, and **downsampling**.

*➤ This leads to loss of spatial information → harder to know exact pixel positions.*

- ✓ 2. Fully connected layers discard almost all spatial info.

# From Classification to Segmentation

- Why Standard CNN Architectures for Classification Fail at Segmentation
  - 3. For segmentation, we need both for "Translation Variance"!




- 1. Low-level features → for precise boundaries (locality).
- 2. High-level features → for semantic understanding.

# From Classification to Segmentation

## ■ How CNN Achieves Translation Invariance


### • 1. Translation Equivariance vs. Translation Invariance

- **Equivariance:** Output changes position in the same way as input.
  - ✓ Example: If a cat's face moves in the input, the feature map also moves.
- **Invariance:** Output remains the same even if the input shifts.
  - ✓ Example: A cat is still labeled "cat" no matter where it appears.



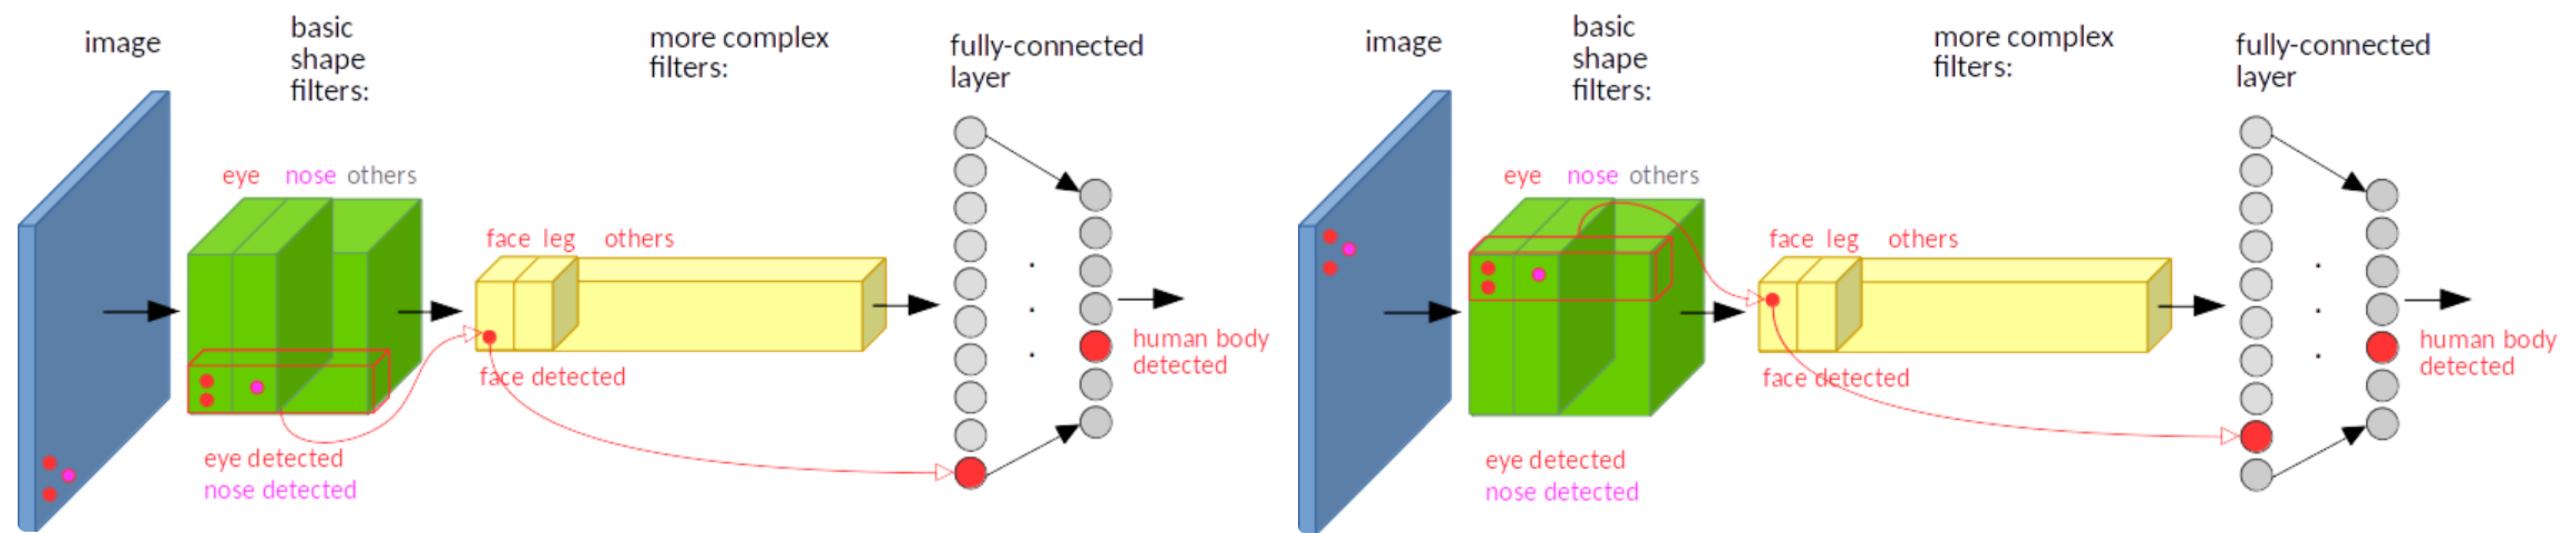
### • 2. Why CNN Has Translation Invariance

- Convolution operation → **translation equivariant** by nature.
- **Weight sharing:** Same filters detect patterns anywhere in the image.
- **Pooling:** Reduces sensitivity to small shifts (small-scale invariance).
- **Fully connected layers + Softmax:** Final classification ignores position.



# From Classification to Segmentation


- How CNN Achieves Translation Invariance
  - 3. Example Flow in CNN





- Early convolution layers: Detect low-level features (eyes, nose) at exact positions → still equivariant.
- Deeper convolution layers: Combine features into higher-level patterns (face, leg).
- FC layer & Softmax: Output label probability is **position-independent** → invariant.

# From Classification to Segmentation


- How CNN Achieves Translation Invariance
  - 4. Why This Matters




Classification Task



Face



Segmentation Task



Not Face

- For **classification**: Translation invariance is beneficial.
- For **segmentation**: It can cause **loss of locality**, harming pixel-level predictions.

# From Classification to Segmentation

---

## ■ Detection Needs Translation Variance

- Why Detection Needs Translation Variance

- Classification: Output label does not depend on location → *translation invariant*.
- Detection: Must **localize** the object → output changes when the object moves.
- Translation variance means the network output must change according to the object's position.

- Key Requirements for Detection

- Preserve Spatial Information


- ✓ **Spatial information** = exact positional coordinates of features in the image.
    - ✓ Allows model to determine *where* the object is located.
    - ✓ Loss of spatial info → bounding boxes drift or fail completely.

- Preserve Locality

- ✓ **Locality** = maintaining the relationship between features **within a region** of the image.
    - ✓ Important for capturing an object's **internal structure** (e.g., face = eyes + nose).


# From Classification to Segmentation

- Detection Needs Translation Variance
  - Spatial Information vs. Locality – What's the Difference?



Detected Object

Spatial Information-based Facial Identification



Detected Object

Locality-based Facial Identification

Term	Meaning	Example in Detection Task
Spatial Information	Absolute and relative coordinates of features in the image	"The face is at (x=150, y=200)"
Locality	Preservation of spatial relationships within a region	"The ears are above the eyes, the nose is below the eyes"

# From Classification to Segmentation

---


## ■ How to Preserve Spatial Information and Locality

### • Problem Recap

- 1. Pooling & stride in CNN → **downsampling** → loss of spatial resolution.
- 2. **Spatial information:** Absolute position of features in the image.
- 3. **Locality:** Relative arrangement of features within a region.
- 4. Loss of either → blurry boundaries & inaccurate localization.

# From Classification to Segmentation

- How to Preserve Spatial Information and Locality
  - Solution 1: Encoder–Decoder Structure



- Encoder

- Encoder
  - ✓ Uses **convolutions + pooling** to extract high-level semantic features.
  - ✓ Gradually reduces spatial dimensions ( $W \times H \rightarrow w \times h$ ).

- Decoder

- Decoder
  - ✓ Uses **upsampling** to restore feature maps to original size ( $w \times h \rightarrow W \times H$ ).
  - ✓ Generates segmentation map where each pixel has a class label.


- Problem in Simple Encoder–Decoder

- Problem in Simple Encoder–Decoder
  - ✓ High-level features from encoder lack **low-level details** (edges, fine structures).
  - ✓ Upsampling alone cannot reconstruct precise boundaries.

# From Classification to Segmentation

---

- How to Preserve Spatial Information and Locality
  - Solution 2: Skip Connections



- Directly connect encoder layers to corresponding decoder layers.
- Merge **low-level spatial details** from encoder with **high-level semantics** from decoder.
- Restores sharp object boundaries and preserves locality.
- Benefits
  - ✓ Retain **low-level details** (edges, textures).
  - ✓ Preserve **spatial resolution** across network.
  - ✓ Combine **locality + semantic context** for accurate segmentation.

# Introduction to Fully Convolutional Networks (FCNs)

## ▪ Fully Convolutional Networks for Semantic Segmentation


- Goal

- Perform **semantic segmentation**: assign a **class label to every pixel** in an image.
- Achieve **end-to-end learning** from input image to dense output.
- Balance **semantic meaning** (*what is present*) with **spatial precision** (*where it is located*).


Fully-connected layers

- Key Idea

- Transform existing high-performing classification CNNs (e.g., AlexNet, VGG, GoogLeNet) into **fully convolutional networks**.
- Replace **fully connected layers** with **convolutional layers** to preserve spatial information.
- Adapt these networks for **dense prediction** tasks.



VGG



# Introduction to Fully Convolutional Networks (FCNs)

---

## ▪ Fully Convolutional Networks for Semantic Segmentation

- Key Insight

- Fully Convolutional Design
  - ✓ 1. Accepts **arbitrary-sized inputs**.
  - ✓ 2. Produces outputs of **corresponding spatial size**.
  - ✓ 3. Enables **efficient whole-image training & inference**.
- Avoids complex pre/post-processing (superpixels, proposals, CRFs).
- Fine-tune from supervised pre-training for strong initial feature representations.

- Core Steps

- Convolutionalization
  - ✓ Replace fully connected layers with equivalent convolution layers.
  - ✓ Maintain spatial correspondence and allow variable input sizes.
- Upsampling (Deconvolution)
  - ✓ Convert coarse feature maps into dense, high-resolution predictions.
- Skip Architecture
  - ✓ Fuse **deep, coarse, semantic** features with **shallow, fine, appearance** features for accurate boundaries.

# Convolutionalization: Motivation

## ■ Why Replace Fully Connected Layers for Segmentation?

### • Problems with Fully Connected Layers

#### ○ 1. Loss of Spatial Information


✓ FC layers flatten feature maps, discarding pixel location data.

✓ Critical spatial correspondences (what is where) are lost.

#### ○ 2. Fixed Input Size Constraint

✓ FC layers require a fixed vector size  
→ fixed input image dimensions.

✓ Limits flexibility for variable-sized images.




# Convolutionalization: Motivation

- Why Replace Fully Connected Layers for Segmentation?

- Problems with Fully Connected Layers

- 3. Coarse Outputs for Dense Prediction

- ✓ Output is a single class label for the whole image, not per-pixel.



- Summary – Why This is a Problem for Segmentation


- ✓ 1. Semantic segmentation needs **dense, pixel-to-pixel predictions**.
      - ✓ 2. Requires **both semantic meaning** (what the object is) and **location information** (where it is).
      - ✓ 3. FC layers destroy the "where" component.

# Convolutionalization: Motivation

- Why Replace Fully Connected Layers for Segmentation?

- The FC-to-Conv Transformation (*Convolutionalization*)

- Replace FC layers with equivalent convolutional layers.
      - ✓ Example: VGG16's first FC layer  
→ 7×7 Conv layer with the same number of parameters.
      - ✓ Final FC layer  
→ **1×1 Conv layer** with channels = number of classes.
    - Benefits
      - ✓ **(1)** Accepts arbitrary input sizes
      - ✓ **(2)** Produces spatial output maps that preserve location cues
      - ✓ **(3)** Enables end-to-end pixel-level learning.



# Convolutionalization: Motivation

## ■ Why Replace Fully Connected Layers for Segmentation?

### • The FC-to-Conv Transformation (*Convolutionalization*)


- Example: Replacing FC Layers with Convolutional Layers

#### ✓ Problem with FC Layers

- Fixed input dimension  
(e.g., VGG-16 requires a vector of length **4096**).
- Spatial coordinates are lost when flattening.

#### ✓ Key Idea

- Treat an FC layer as a **convolution** with a kernel covering the entire input feature map.
- Replace
  - First FC layer  
→ **7x7 Conv** with 4096 filters.
  - Final FC layer  
→ **1x1 Conv** with channels = number of classes.



# Convolutionalization: Motivation

## ■ Why Replace Fully Connected Layers for Segmentation?

### • The FC-to-Conv Transformation (*Convolutionalization*)

- Example: Replacing FC Layers with Convolutional Layers

#### ✓ VGG-16 Example

##### ➤ Before



- Input feature map:  $7 \times 7 \times 512$ .
- Flatten to  $1 \times 1 \times (512 \times 7 \times 7) = 1 \times 1 \times 25088$ .
- Fully Connected to 4096 units.

##### ➤ After

- Apply **7x7 Conv** (512 input channels, 4096 filters) → Output  $1 \times 1 \times 4096$ .
- Spatial structure is preserved for larger inputs.

#### ✓ Advantages

- Works with **any input size** (no FC size constraint).
- Produces **spatial output maps**  
→ essential for segmentation.
- Computation is **faster** for large images  
→  $\approx 5 \times$  speedup over patch-by-patch.




# Coarse Outputs and the Need for Upsampling

## ■ Why Upsampling is Essential for Segmentation

### • 1. Downsampling in CNNs

- Pooling layers and strided convolutions progressively reduce spatial resolution.



- Example: VGG-16


✓ Input:  $224 \times 224$

✓ After 5 pooling layers (stride 2 each) → Output feature map:  $7 \times 7$

✓ This is a **stride 32** reduction.

# Coarse Outputs and the Need for Upsampling

- Why Upsampling is Essential for Segmentation
  - 2. The Problem for Segmentation



- Fully convolutional conversion preserves spatial arrangement, *but resolution is still coarse*.
- For pixel-wise prediction, we need **near-original resolution** output.
- Without upsampling, fine details and object boundaries are lost.

# Upsampling with Learnable Deconvolution

## From Coarse to Dense: Deconvolution Layers

### 1. Upsampling Methods

- Fixed

- ✓ Bilinear Interpolation

- Simple, non-learnable resizing.

- ✓ Unpooling

- Reverses pooling using stored indices.

- Learnable

- ✓ Deconvolution

- (=Transposed Convolution, Backwards Convolution)


- Parameters learned during training.

### 2. Backwards Convolution Concept

- Standard convolution (stride > 1) **reduces** spatial size (downsampling).

- Reversing this process **increases** spatial size (upsampling).

- Uses **learnable filters** so that the upsampling is **task-optimized**.



# Upsampling with Learnable Deconvolution



## From Coarse to Dense: Deconvolution Layers

- How FCN Uses It

- FCNs employ deconvolution layers for upsampling coarse feature maps to the input size.*

- Initial weights are set to bilinear interpolation values.
  - During training, these weights are updated to improve segmentation accuracy.
  - Allows **end-to-end training** with pixel-wise loss.

**Deconvolution  
(i.e., upsampling)**



# The Problem with Coarse Outputs

## ■ Why Coarse Outputs are Not Enough

- Deconvolution alone is insufficient

- Even with trainable deconvolution, a coarse feature map (e.g., stride 32) lacks detailed spatial information.

- Loss of detail

- Large stride means much information is lost during down sampling, leading to blurry and imprecise boundaries.



- Example

- FCN-32s often produces smooth, rounded edges rather than accurate object contours.

- Reason

- The receptive field covers a large area, so fine-grained boundaries from the original image are not preserved.

Deconvolution  
(i.e., upsampling)



Prediction



Ground truth


Loss of detail

# Skip Architecture

## Skip Architecture for Better Segmentation

- Core idea – Fuse

Layer 1   Layer 2   Layer 3   Layer 4   Layer 5




Shallow

Deep

Fine  
Location  
Local  
Detail

Coarse  
Semantic  
Global  
Abstract


- Deep & coarse layers → rich semantic (meaning) information.
- Shallow & fine layers → detailed appearance and boundary information.




# Skip Architecture

## Skip Architecture for Better Segmentation

### Architecture variants



**FCN-8s: Final + pool4 + pool3  
(stride 8) — most detailed boundaries.**




**FCN-16s: Final + pool4 layer  
(stride 16) — better boundaries.**

# FCN Results and Key Takeaways

## ■ FCN Performance Highlights

- **Performance Trend**




- **FCN-32s → FCN-16s → FCN-8s:** Steady improvement in segmentation accuracy.
- Finer strides capture more precise object boundaries.

- **Quantitative Impact**

- Significant mean IU improvement over prior state-of-the-art.
- ~20% relative gain on benchmarks such as **PASCAL VOC 2011/2012**, **NYUDv2**, and **SIFT Flow**.

- **Legacy**

- Provided the foundation for later architectures (U-Net, SegNet, DeepLab, etc.).



# Introduction to U-Net

---

- **U-Net: Convolutional Networks for Biomedical Image Segmentation**

- Published at **MICCAI 2015** by Olaf Ronneberger et al.
- Designed for **biomedical image segmentation** with **limited training data**.
- Key innovation: **U-shaped architecture** combining
  - **Contracting Path** → capture context
  - **Expansive Path** → enable precise localization
- Outperformed previous state-of-the-art on ISBI 2015 challenges.

# Motivation – U-Net

---

## ■ Why U-Net?

- **1. Problem in Biomedical Segmentation**


- Small, annotated datasets
- Need for **pixel-level** classification (semantic segmentation)
- Critical to preserve both

✓ **Context**


➤ Global semantic understanding (what is in the image)

✓ **Localization**

➤ Precise spatial boundaries (where the object is)



Biomedical Image



Ground-truth (Label)

- **2. Limitations of Sliding-Window CNNs (Patch-Based Methods)**

- **Inefficient computation:** neighboring patches overlap heavily → redundant convolution operations.
- **Context–Localization trade-off**
  - ✓ Large patch → captures context but loses fine detail.
  - ✓ Small patch → keeps detail but misses context.
- Cannot leverage **shared computation** across overlapping regions.

# Motivation – U-Net

## ■ Why U-Net?

### • 3. U-Net's Approach

- (1) Fully Convolutional


- ✓ Process the whole image at once.

- (2) Encoder–Decoder Structure

- ✓ Encoder (contracting path) captures context.
  - ✓ Decoder (expansive path) restores spatial resolution.

- (3) Skip Connections

- ✓ Combine fine detail from early layers with deep semantic info.
  - ✓ Designed to work **efficiently** with **small biomedical datasets**.



# Architecture Overview

## The U-Shaped Architecture

### 1. Contracting Path (Encoder)

- Purpose

- ✓ Capture context by progressively downsampling the input image while increasing the number of feature channels.

- Structure per Downsampling Step

- ✓ Step 1. 3x3 Convolution (stride=1, valid padding) → ReLU


- (optional) BatchNorm

- No padding → feature map shrinks by 2 pixels per conv in each spatial dimension.

- ✓ Step 2. 3x3 Convolution (stride=1, valid padding) → ReLU

- (optional) BatchNorm

- ✓ Step 3. 2x2 Max Pooling (stride=2) for downsampling.




# Architecture Overview

## ■ The U-Shaped Architecture

### • 1. Contracting Path (Encoder)

- After **each** downsampling step, the number of channels **doubles**
- ✓ E.g.,  $1 \rightarrow 64 \rightarrow 128 \rightarrow 256 \rightarrow 512 \rightarrow 1024$  (bottleneck).
- Pooling **increases the receptive field**, allowing the network to capture more context from the input image.

→ conv 3x3, ReLU  
→ copy and crop  
↓ max pool 2x2



# Architecture Overview

---

## ■ The U-Shaped Architecture

### • 2. Bottleneck

- Purpose

- ✓ Bridge between encoder and decoder, containing the most abstract representation of the input.

- Structure



- ✓ 3x3 Convolution (stride=1, valid padding) → ReLU → (optional) BatchNorm

- ✓ 3x3 Convolution (stride=1, valid padding) → ReLU → (optional) BatchNorm

- Dropout Layer (optional in modern versions; not in original)

- ✓ Helps prevent overfitting and improves robustness to noise.

- This layer has the highest channel count and smallest spatial resolution.

# Architecture Overview

## ■ The U-Shaped Architecture

### • 3. Expansive Path (Decoder)

- Purpose

- ✓ Recover spatial resolution and produce dense

- Pixel-level predictions by combining **contextual features** from the encoder with **localization features** from early layers.

- Structure per Upsampling Step


- ✓ Step 1. 2x2 Transposed Convolution (Deconvolution) (stride=2)

- Upsamples the feature map by a factor of 2.
  - Halves the number of channels.

- ✓ Step 2. Crop & Concatenate


- Take the corresponding feature map from the encoder (same depth level).

- Because of **no padding** in 3x3 convolutions, encoder feature maps are slightly smaller → **crop** to match dimensions before concatenation.



# Architecture Overview

- The U-Shaped Architecture
  - 3. Expansive Path (Decoder)



- Purpose

✓ **Recover spatial resolution** and produce dense

➤ Pixel-level predictions by combining **contextual features** from the encoder with **localization features** from early layers.

# Architecture Overview

## ■ The U-Shaped Architecture

### • 3. Expansive Path (Decoder)

#### ○ Structure per Upsampling Step

##### ✓ Step 1. 2×2 Transposed Convolution (Deconvolution) (stride=2)


- Upsamples the feature map by a factor of 2.
- Halves the number of channels.

##### ✓ Step 2. Crop & Concatenate

- Take the corresponding feature map from the encoder (same depth level).
- Because of **no padding** in 3×3 convolutions, encoder feature maps are slightly smaller → **crop** to match dimensions before concatenation.

##### ✓ Step 3. 3×3 Convolution (stride=1, valid padding) → ReLU → (optional) BatchNorm

##### ✓ Step 4. 3×3 Convolution (stride=1, valid padding) → ReLU → (optional) BatchNorm




# Architecture Overview

## ■ The U-Shaped Architecture

### • 4. Output Layer

- **1x1 Convolution** with **C filters**, where  $C = \text{number of segmentation classes}$ .
- Produces a  $(H \times W \times C)$  score map, where each pixel has a vector of class probabilities.



### • 5. Skip Connections

#### ○ Purpose

✓ Preserving Spatial Detail with Skip Connections

#### ✓ Without skips


➤ decoder relies solely on deep, coarse features → blurry outputs.

#### ✓ With skips – combine

➤ Low-level appearance features (from early layers)

➤ High-level semantic features (from deep layers)

- Requires **crop & copy** due to valid convolutions reducing feature map size.




# Training Strategies in U-Net

## ■ Key Training Strategies for High-Quality Segmentation

- **Strategy 1. Overlap-Tile Strategy**

- Purpose: Handle large images and ensure predictions are accurate at boundaries.



- How it works

- ✓ Split large input images into smaller tiles with overlapping regions.
    - ✓ Each tile is fed into the network, producing a smaller output region due to valid padding in convolutions.
    - ✓ Overlap ensures full coverage and avoids loss of context at the borders.

- Key Benefit

- ✓ Maintains context for edge regions and prevents border artifacts.

# Training Strategies in U-Net

## ■ Key Training Strategies for High-Quality Segmentation

- **Strategy 2. Mirroring Extrapolation**


- **Purpose:** Provide meaningful context for border pixels without introducing artificial padding values.

- **How it works**


- ✓ Extend the input image by mirroring its borders.
    - ✓ This mirrored extension surrounds the original image before tiling.

- **Why it's useful in biomedical images**

- ✓ Many biological structures (e.g., cells) have symmetrical patterns.
    - ✓ Mirroring preserves realistic textures at the borders, helping the network learn better boundary representations.



Origin Image



After Mirroring

# Training Strategies in U-Net


## ■ Key Training Strategies for High-Quality Segmentation

### • Strategy 3. Weighted Loss for Boundary Separation

- Purpose: Improve segmentation in regions where objects touch or have thin boundaries.

- How it works

Segmented Image



Visualized Weight Map

- ✓ Compute a **weight map** from ground truth
  - Higher weights near boundaries between touching objects.
  - Lower weights elsewhere.
- ✓ Use this weight map in the pixel-wise cross-entropy loss to emphasize difficult boundary pixels.

- Effect

- ✓ The network learns to focus on separating closely packed or touching objects.

# Summary & Key Takeaways

---

## ■ FCN & U-Net: Foundations of Modern Segmentation

### • 1. FCN (Fully Convolutional Network)

- **Key Idea:** Replace fully connected layers with convolution layers → accept arbitrary input size.
- **Upsampling:** Learnable deconvolution layers (backwards convolution) to restore resolution.
- **Skip Architecture:** Combines deep semantic info with shallow spatial info for sharper boundaries (FCN-32s → FCN-16s → FCN-8s).
- **Impact:** First end-to-end trainable CNN for dense prediction; influenced most later segmentation models.

### • 2. U-Net

- **Architecture:** Symmetric **U-shape** with **Contracting Path** (context) and **Expansive Path** (localization).
- **Skip Connections:** Concatenate encoder features to decoder features at matching resolution for precise segmentation.
- **Specialized Training Strategies:**
  - ✓ **Overlap-Tile** for large images
  - ✓ **Mirroring Extrapolation** for border pixels
  - ✓ **Weighted Loss** to separate touching objects
- **Impact:** State-of-the-art in biomedical segmentation, strong performance with limited data.

# Summary & Key Takeaways

---

- FCN & U-Net: Foundations of Modern Segmentation

- 3. Key Insights

- **Preserving spatial information** is crucial for pixel-wise tasks → skip connections are a common solution.
    - **Upsampling quality** directly affects boundary accuracy.
    - **Transfer learning** and **augmentation** help overcome limited labeled data.