

Computer Vision

Week 12-13

2025-2

Mobile Systems Engineering
Dankook University

Motivation: Why Study Vision Transformers?

■ From Inductive Bias in CNNs to Data-Driven Transformers

• Why Transformers in Vision?

- CNNs have been the backbone of computer vision for decades.

- Their strength comes from strong *inductive biases*

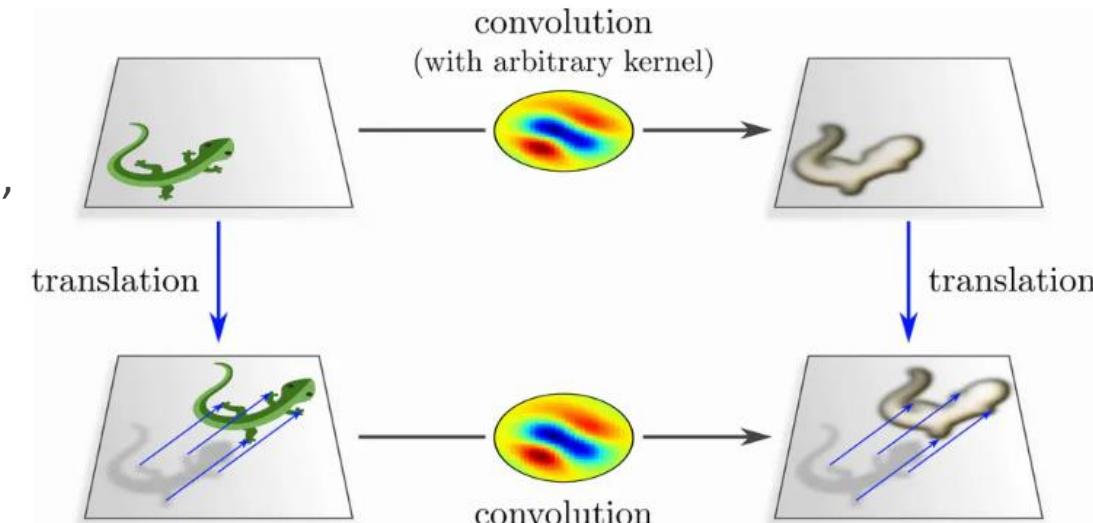
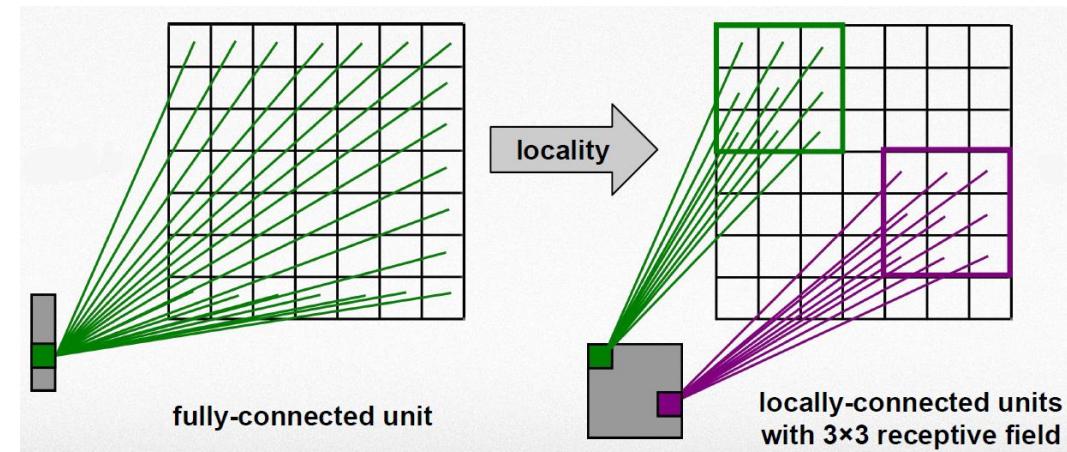
- ✓ **Locality**

- Convolutional filters focus on local neighborhoods.

- ✓ **Translation Equivariance**

- Features shift predictably when the input image shifts.

- These biases allow CNNs to learn effectively from **limited data**, but they can also **restrict scalability** when moving to larger datasets or more complex patterns.



Motivation: Why Study Vision Transformers?

- From Inductive Bias in CNNs to Data-Driven Transformers

- Recap – What is Translation Equivariance?

- Why CNNs Can Detect the Same Pattern Anywhere in the Image

- ✓ Meaning

- Translation equivariance means that if the input shifts in space, the output feature map also shifts in the same direction.

- Example

- If a cat in an image moves **10 pixels to the right**, the convolution feature map of the cat will also appear **10 pixels to the right**.

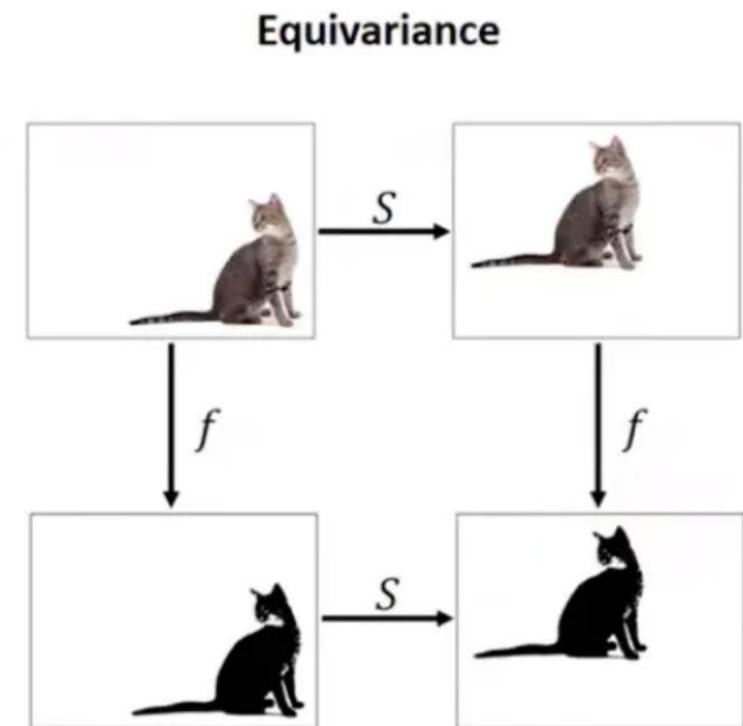
- ✓ Role of Kernels (Filters)

- The **kernel weights do NOT change** when the input moves.

- During training, kernels learn to detect specific local patterns (e.g., cat's ear, eye, whisker).

- The **same kernel** is applied across the entire image (sliding operation).

- As a result, the same feature can be detected **regardless of its position** in the image.



Motivation: Why Study Vision Transformers?

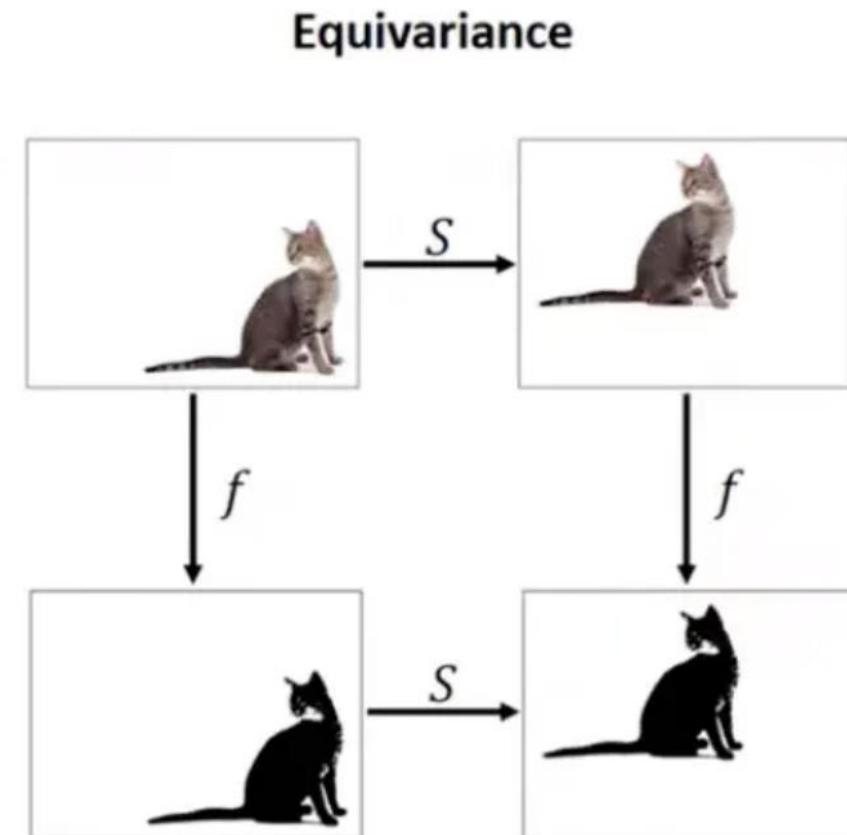
- From Inductive Bias in CNNs to Data-Driven Transformers

- Recap – What is Translation Equivariance?

- Why CNNs Can Detect the Same Pattern Anywhere in the Image

- ✓ Key Takeaway

- Translation equivariance = input shift \rightarrow output shift.
 - Kernel weights remain fixed after training.
 - Weight sharing across the whole image ensures that features are recognized anywhere, independent of position.



Motivation: Why Study Vision Transformers?

■ From Inductive Bias in CNNs to Data-Driven Transformers

• What is Inductive Bias?

○ *Inductive Bias = the structural assumptions that a model inherently carries.*

✓ It reflects how a model internally views the relationship between **input data** and **output predictions**.

○ Example

Component	Entities	Relations	Rel. inductive bias	Invariance
Fully connected	Units	All-to-all	Weak	-
Convolutional	Grid elements	Local	Locality	Spatial translation
Recurrent	Timesteps	Sequential	Sequentiality	Time translation
Graph network	Nodes	Edges	Arbitrary	Node, edge permutations

✓ **CNNs** assume that local pixel groups contain meaningful features (locality).

✓ **RNNs** assume sequential dependencies between timesteps.

✓ **Fully connected layers** assume all inputs may interact equally.

○ These built-in assumptions **guide learning**, reducing data requirements but also limiting flexibility.

Motivation: Why Study Vision Transformers?

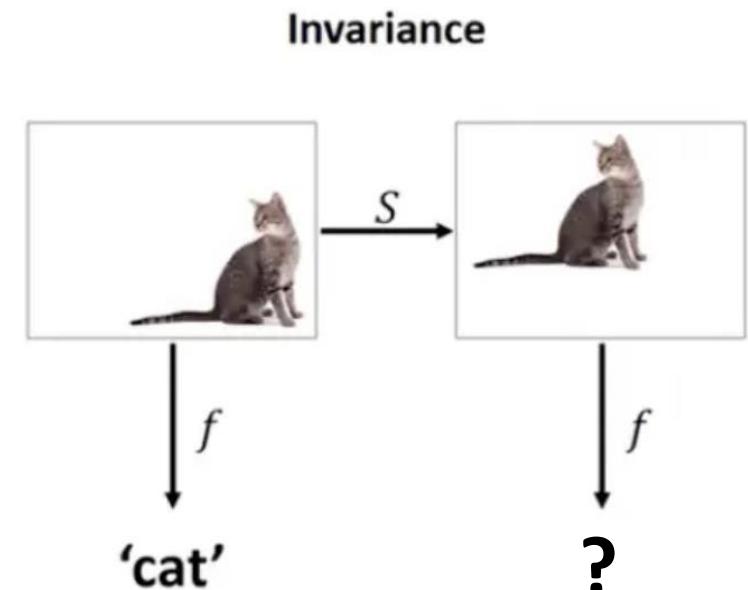
- From Inductive Bias in CNNs to Data-Driven Transformers

- **Transformers and Weak Inductive Bias**

- **Transformers lack strong inductive biases** like locality or translation equivariance.
 - Instead, they rely on **global self-attention** to learn all relationships directly from data.
 - This makes them
 - ✓ **Data-hungry**: need very large datasets for effective training.
 - ✓ **Flexible and scalable**: can capture global dependencies more efficiently once enough data is available.

- **Example**

- If a cat shifts position in the image (see figure), CNNs still classify it as “*cat*” due to weight sharing and invariance.
 - A Transformer, however, must **learn from data** that “cat on the left” and “cat on the right” are the same class.
 - Without large training data, this generalization may fail.



Vision Transformers – Overall Architecture

■ What is a Vision Transformer (ViT)?

- A Vision Transformer (ViT) is a model that applies the **Transformer architecture to images** by

- **Step 1.**
Splitting the image into fixed-size patches.
 - **Step 2.**
Flattening and linearly projecting
each patch into an embedding vector.
 - **Step 3.**
Adding **positional encodings** to
preserve spatial information.
 - **Step 4.** Adding a **learnable class token**
for classification.
 - **Step 5.** Passing the sequence of embeddings
through **Transformer encoder layers**.



Vision Transformers – Overall Architecture

■ What is a Vision Transformer (ViT)?

- One-Sentence Summary

- ViT treats an image as a **sequence of patch tokens**, just like words in NLP, and applies a Transformer encoder for classification or transfer learning.

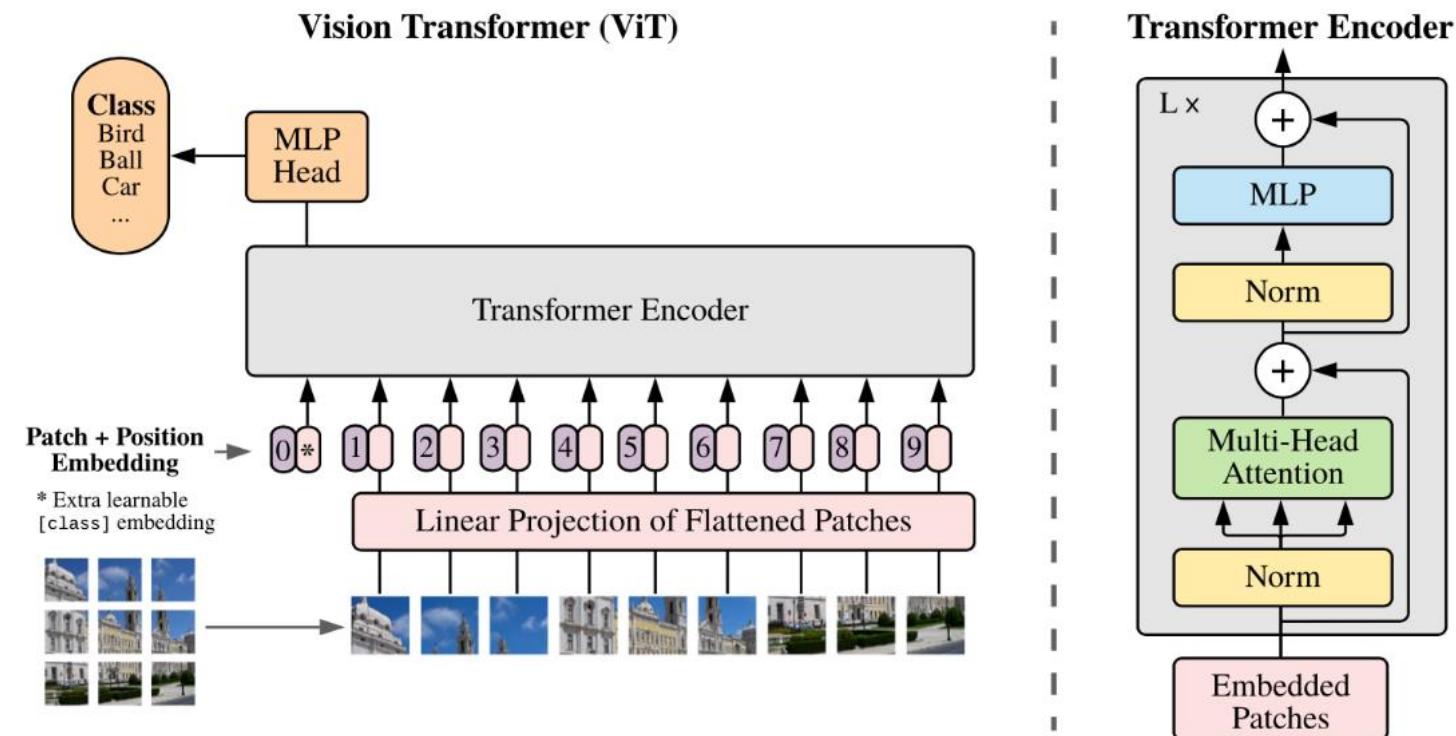
■ Architecture Overview

- Input = patch embeddings + [CLS] + positional embeddings

- Each of L blocks

- LayerNorm
- Multi-Head Self-Attention (MSA) + residual
- LayerNorm
- MLP (GELU) + residual

- Final [CLS] embedding → image representation → MLP/linear head



Data Preparation for Vision Transformer Input

■ How Images Become Tokens in Vision Transformer

- From pixels to patch embeddings

- Steps

- ✓ Step 1. Image Patch Creation**

- Split the image into fixed-size non-overlapping patches.
 - Example: For a 224×224 image and 16×16 patch size $\rightarrow 14 \times 14 = 196$ patches.

- ✓ Step 2. Patch Embedding**

- Flatten each patch ($P \times P \times C$) into a vector of length $P^2 \cdot C$.
 - Apply a **linear projection** to map it into **D-dimensional embedding space**.

- ✓ Step 3. Add Class Token**

- Add a special learnable vector [CLS] to serve as the **image representation** for classification.

- ✓ Step 4. Positional Embedding**

- Since Transformers lack spatial structure, add positional encodings so the model knows patch order.

Data Preparation for Vision Transformer Input

■ How Images Become Tokens in Vision Transformer

- From pixels to patch embeddings

- Steps

- ✓ Step 1. Image Patch Creation**

- Split the image into fixed-size non-overlapping patches.
 - Example: For a 224×224 image and 16×16 patch size $\rightarrow 14 \times 14 = 196$ patches.

- ✓ Step 2. Patch Embedding**

- Flatten each patch ($P \times P \times C$) into a vector of length $P^2 \cdot C$.
 - Apply a **linear projection** to map it into **D-dimensional embedding space**.

- ✓ Step 3. Class Token**

- Add a special learnable vector [CLS] to serve as the **image representation** for classification.

- ✓ Step 4. Positional Embedding**

- Since Transformers lack spatial structure, add positional encodings so the model knows patch order.

Data Preparation for Vision Transformer Input

■ Step 1. Image Patch Creation

- From Image to Patches

- Input image: $x \in \mathbb{R}^{C \times H \times W}$

- ✓ C : number of channels (e.g., 3 for RGB)

- ✓ H : image height

- ✓ W : image width

- Divide the image into non-overlapping patches of size $P \times P$.

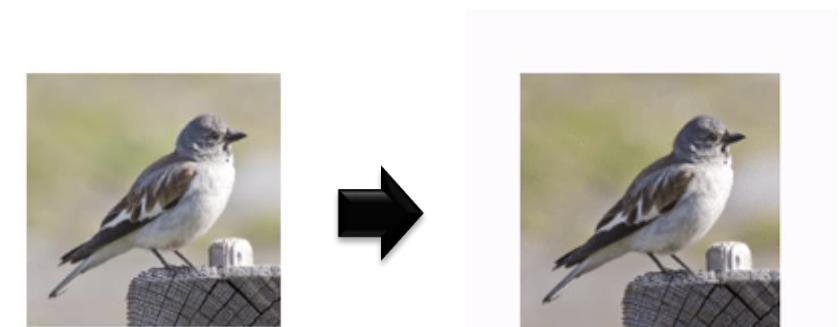
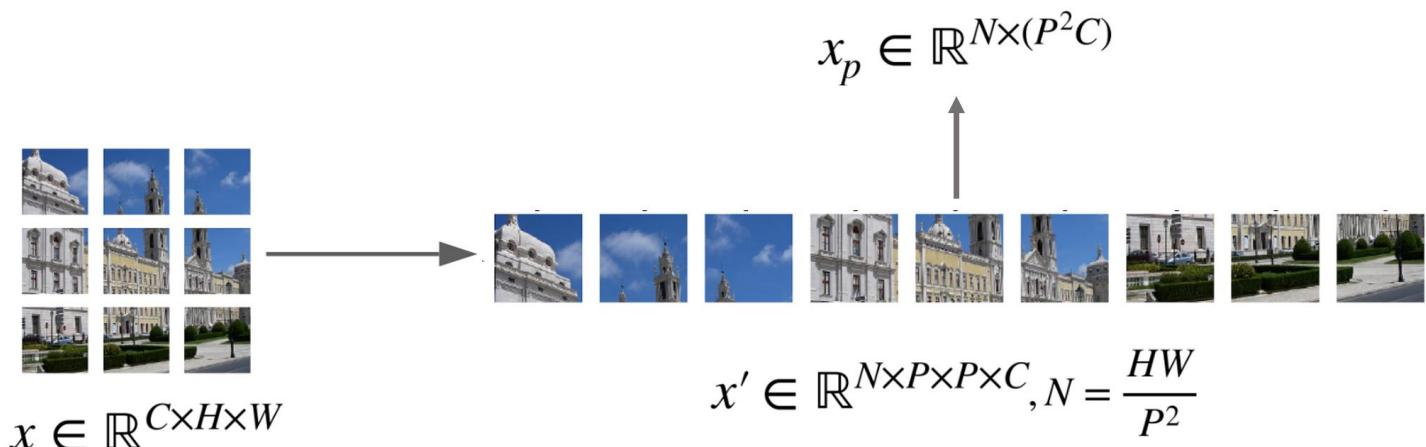
- Total number of patches

$$N = \frac{H \cdot W}{P^2}$$

- After splitting, each patch is represented as

$$x' \in \mathbb{R}^{N \times P \times P \times C}$$

- ✓ where x' denotes the collection of all image patches before flattening.



Data Preparation for Vision Transformer Input

■ Step 2. Patch Embedding

• Step 2-1. Flattening

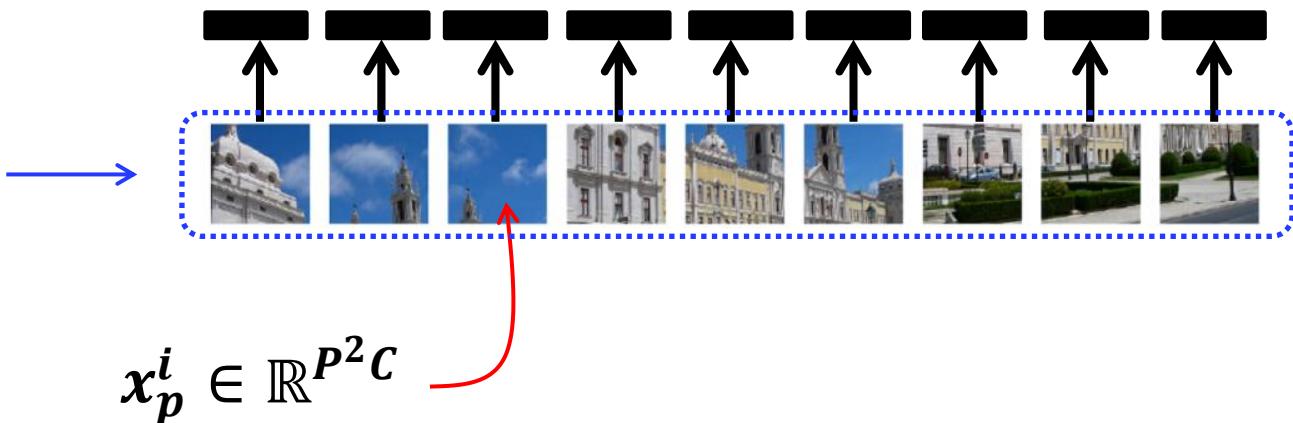
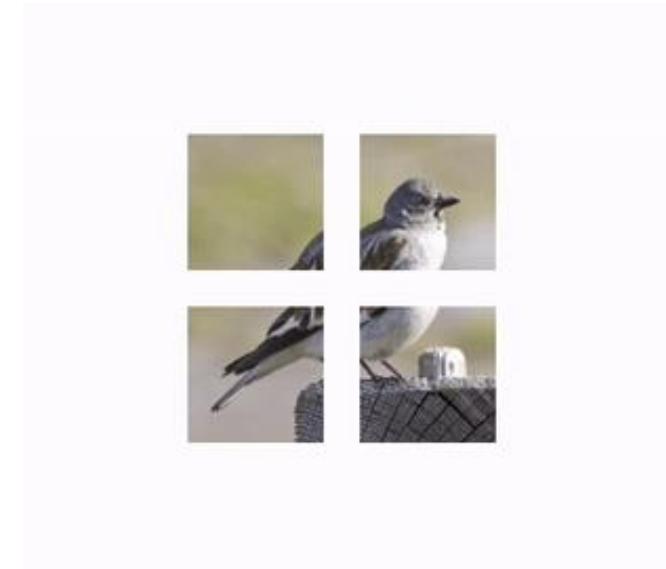
- Each patch x_p^i (the i -th patch) is **flattened** into a vector of length $P^2 \cdot C$

$$x_p^i \in \mathbb{R}^{P^2 C}$$

- Collecting all patches

$$x_p \in \mathbb{R}^{N \times (P^2 C)}$$

$$x_p \in \mathbb{R}^{N \times (P^2 C)}; \text{ where } N = 9$$



Data Preparation for Vision Transformer Input

■ Step 2. Patch Embedding

• Step 2-1. Flattening

- Why Flatten? – CNN vs. ViT

- ✓ CNN (Convolutional Neural Networks)**

- Keep 2D grid structure of images.

- Convolution kernel slides locally across height × width.

- Built-in inductive biases

- Locality (neighboring pixels are strongly related)

- Translation equivariance (shifted object still gives similar features)

Key Difference

- CNN: “I already know neighbors matter, I’ll focus there.”

- ViT: “I don’t assume anything, *I’ll compare all patches and figure out which relations are important.*”

- Flattening is what makes ViT **closer to NLP tokenization** than to CNN feature maps.

✓ Vision Transformer (ViT)

- Flatten patches → treat them as independent tokens.

- No spatial grid is preserved explicitly.

- Purpose: remove hard-coded locality → force the model to **learn relationships through self-attention.**

- Each patch embedding is just a vector — any spatial relation must be discovered by comparing embeddings with others.

Data Preparation for Vision Transformer Input

■ Step 2. Patch Embedding

• Step 2-2. Linear Projection

- Apply a **linear projection** using matrix $E \in \mathbb{R}^{(P^2C) \times D}$

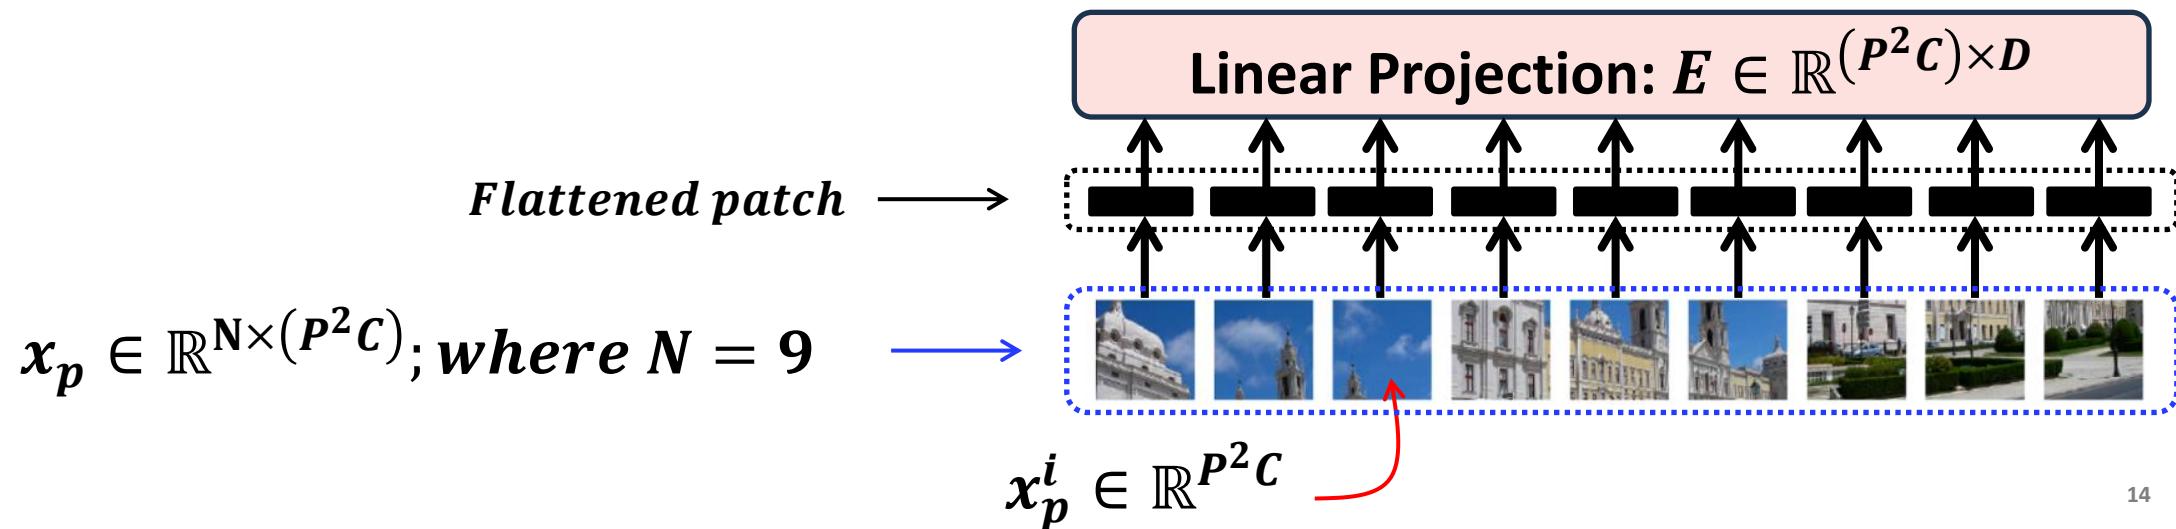
$$x_p \in \mathbb{R}^{N \times (P^2C)} \rightarrow E \in \mathbb{R}^{(P^2C) \times D} \rightarrow [x_p^1 E; x_p^2 E; \dots; x_p^N E] \in \mathbb{R}^{N \times D}$$

✓ Where ...

➤ D : embedding dimension (hidden size for Transformer input)

➤ E : learnable projection matrix

$$[x_p^1 E; x_p^2 E; \dots; x_p^N E] \in \mathbb{R}^{N \times D}$$



Data Preparation for Vision Transformer Input

- Step 2. Patch Embedding

- Step 2-2. Linear Projection

- What is Linear Projection?

- ✓ Definition

- A **Linear Projection** is a simple operation that maps an input vector into another space using a **weight matrix**.

- ✓ Mathematically

$$\mathbf{y} = \mathbf{x}\mathbf{W} + \mathbf{b}$$

- where

- $\mathbf{x} \in \mathbb{R}^{P^2C}$: flattened patch
- $\mathbf{W} \in \mathbb{R}^{(P^2C) \times D}$: projection (weight matrix)
- $\mathbf{b} \in \mathbb{R}^D$: bias term
- $\mathbf{y} \in \mathbb{R}^D$: patch embedding

Data Preparation for Vision Transformer Input

- Step 2. Patch Embedding

- Step 2-2. Linear Projection

- What is Linear Projection?

- ✓ Key Insight

- This operation is exactly the same as a Linear Layer (Fully Connected Layer) in neural networks.

- ✓ Why do we call it “Projection”?

- 1. Because we are *projecting* a high-dimensional patch vector ($P^2 C$) into the embedding dimension D .
 2. “Projection” highlights the idea of dimensionality change (compression or expansion).

- ✓ Connection to Neural Networks

- Linear Projection = Linear Layer = Fully Connected (FC) Layer = Dense Layer (Keras)

- No activation function → purely linear transformation

- In ViT

- Each patch x_p^i is passed through the same Linear Layer → produces a D -dimensional embedding
 - All patches share the same projection weights W .

Data Preparation for Vision Transformer Input

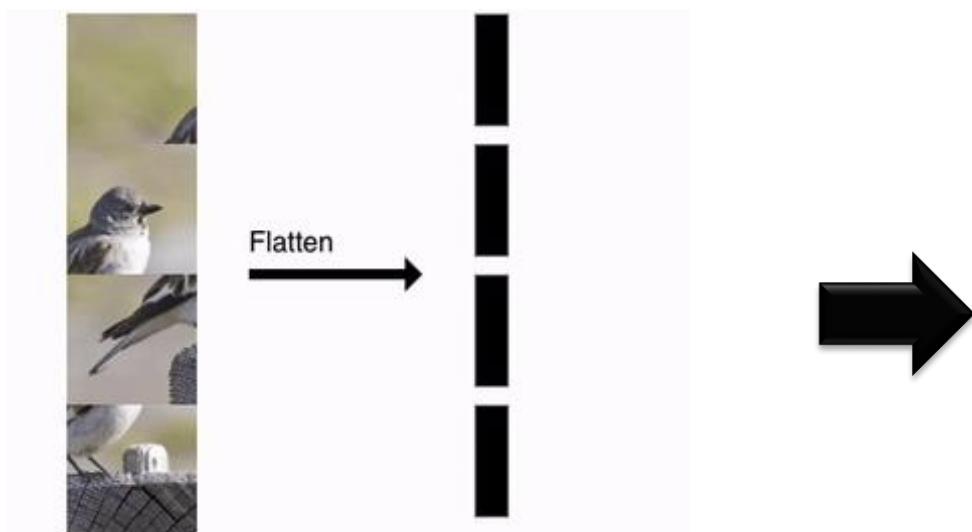
■ Step 2. Patch Embedding

- **Step 2-1. Flattening**

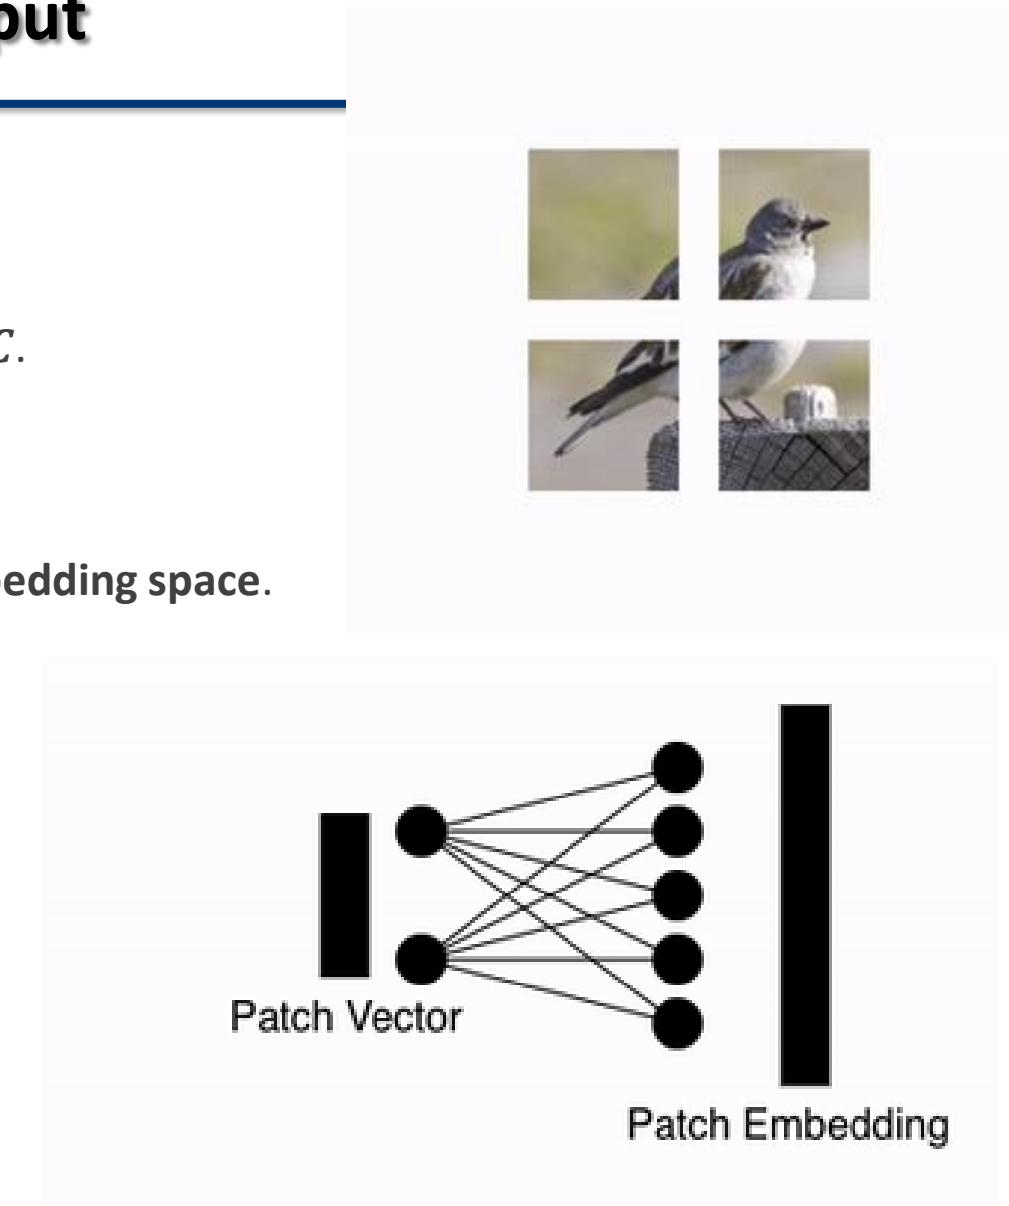
- Flatten each patch ($P \times P \times C$) into a vector of length $P^2 \cdot C$.

- **Step 2-2. Linear Projection**

- Apply a **linear projection** to map it into D -dimensional embedding space.



Creating a patch embedding applying a linear projection



Embedding all patches by applying
a linear projection to all flattened patches

Data Preparation for Vision Transformer Input

■ Step 3. Add Class Token

• Class Token

○ What is a Class Token?

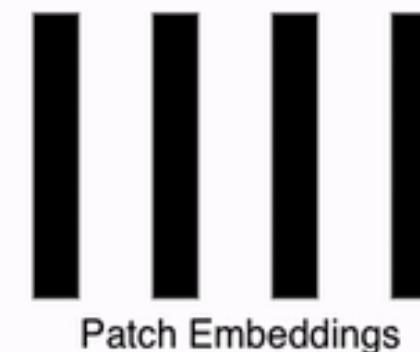
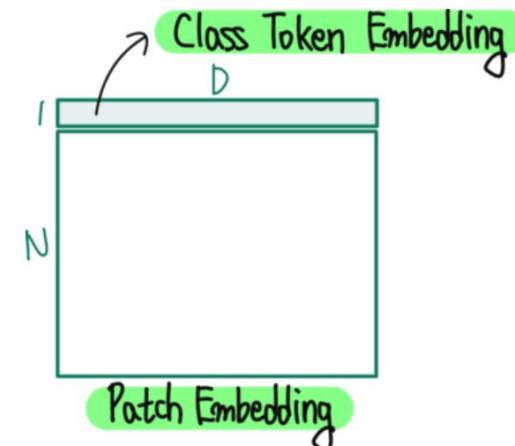
✓ A learnable vector $x_{cls} \in \mathbb{R}^D$ that is prepended to the sequence of patch embeddings.

✓ After appending, the sequence size changes from $\mathbb{R}^{N \times D} \rightarrow \mathbb{R}^{(N+1) \times D}$.

$$[x_{cls}; x_p^1 E; x_p^2 E; \dots; x_p^N E;] \in \mathbb{R}^{(N+1) \times D}$$

✓ This token does NOT correspond to any image patch.

➤ It is an **extra parameter** introduced for classification.



Data Preparation for Vision Transformer Input

■ Step 3. Add Class Token

- Motivation: Why Add a Class Token?

- Inspired by NLP (BERT)

- ✓ In BERT, a [CLS] token is added to summarize a sequence for classification tasks.

- ✓ ViT adopts the same idea: let the model ***gather global information*** into one special token.

- Avoid Pooling

- ✓ CNNs often use **global average pooling** before classification.

- ✓ Instead, ViT uses the [CLS] token to ***aggregate information*** without predefined pooling rules.

- How Does It Work?

- During Multi-Head Self-Attention,

- ✓ The [CLS] token attends to all patch tokens, **collecting information from the entire image**.

- ***Self-Attention allows all tokens to interact from the first layer.***

- ✓ At the same time, all patches can also attend to the [CLS] token.

- This bidirectional attention allows [CLS] to act as a **global information hub**.

Data Preparation for Vision Transformer Input

■ Step 3. Add Class Token

- What Information Does It Contain?

- After passing through all Transformer encoder layers
 - ✓ The [CLS] token embedding z_{cls} becomes a compact representation of the entire image.
 - ✓ It has aggregated class-relevant patterns from all patches.

- Final Use in Classification

- Only the final [CLS] token embedding is passed to the classification head (MLP/Linear Layer).
 - ✓ This single vector is sufficient for classification tasks

$$y = \text{softmax}(W \cdot z_{cls})$$

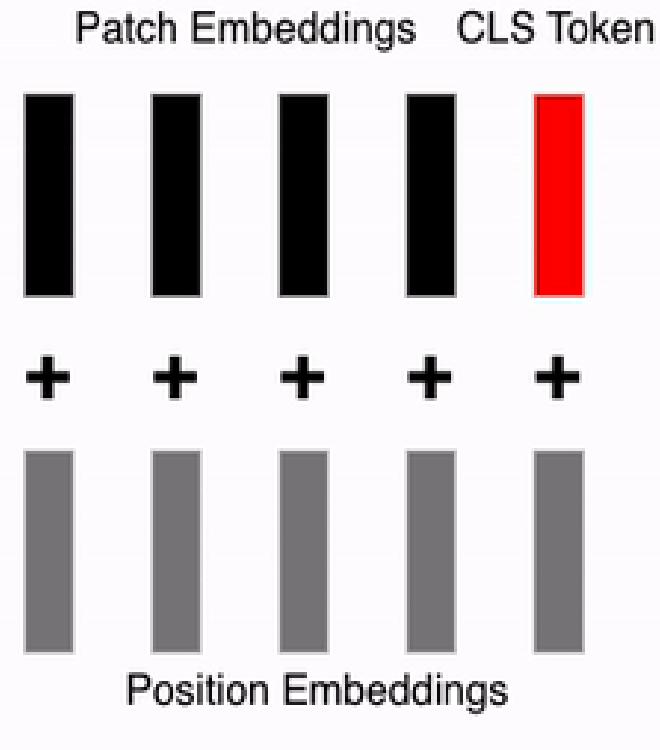
- ✓ In practice, this performs as well (or better) than pooling strategies.
 - “By the end of training, it contains enough global information to decide the class label.”

Data Preparation for Vision Transformer Input

■ Step 4. Positional Embedding

- **Positional Embedding**

- A vector that encodes the **position (index or location)** of each token (or patch) in the input sequence, which is added to the token embedding so that the model knows **where** each token comes from.



- **Why Do We Need Positional Embeddings?**

- Transformer's **Self-Attention** is permutation-invariant
 - ✓ It only looks at the relationships between tokens.
 - ✓ Without extra information, it cannot distinguish whether a patch comes from the **top-left corner** or the **bottom-right corner**.
- But in vision, **spatial order matters** (e.g., eyes above the nose, not below).
- Positional embeddings inject this spatial information into the patch tokens.

Data Preparation for Vision Transformer Input

■ Step 4. Positional Embedding

- How It Works

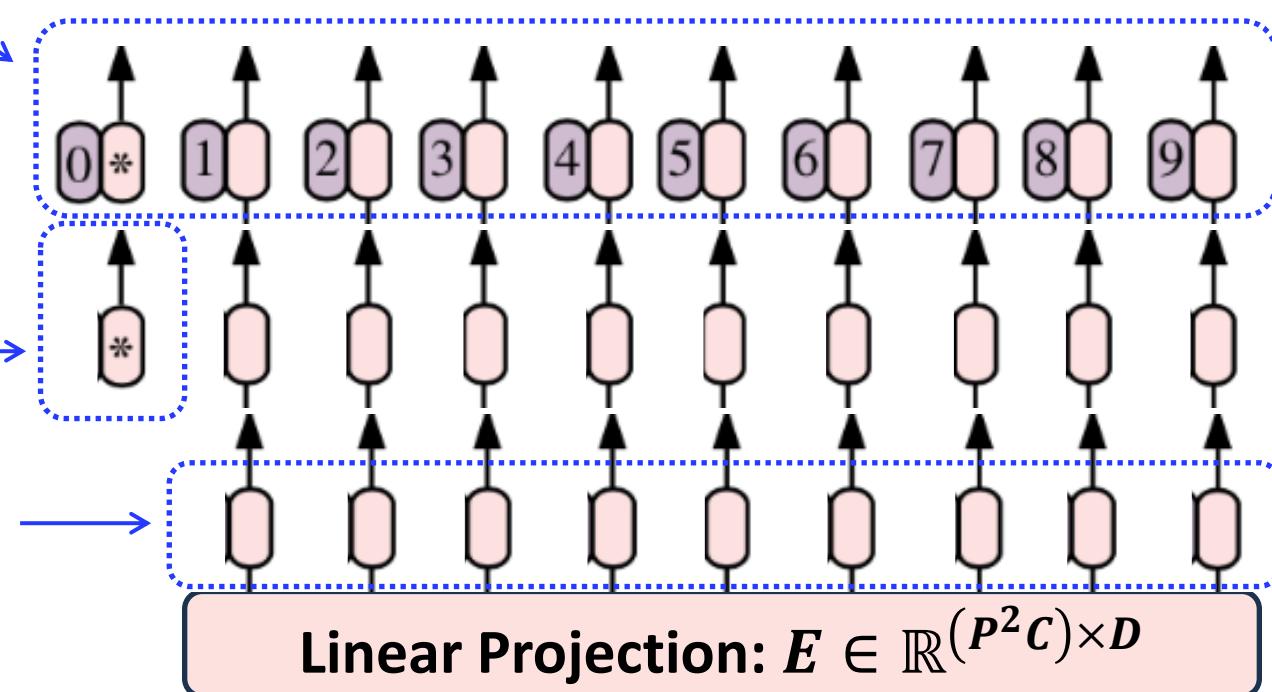
1. After adding the [CLS] token, we have $(N + 1) \times D$ embeddings.
2. A learnable positional embedding matrix $E_{pos} \in \mathbb{R}^{(N+1) \times D}$ is added element-wise

$$z_0 = [x_{cls}; ; x_p^1; E; x_p^2E; \dots; x_p^N E] + E_{pos} \in \mathbb{R}^{(N+1) \times D}$$

3. Each position (patch index) has its own embedding vector, encoding the **absolute location** in the image grid.

$$[x_{cls}; x_p^1E; x_p^2E; \dots; x_p^N E;] \in \mathbb{R}^{(N+1) \times D}$$

$$[x_p^1E; x_p^2E; \dots; x_p^N E;] \in \mathbb{R}^{N \times D}$$



Data Preparation for Vision Transformer Input

■ Step 4. Positional Embedding

- **Variants – Positional Embedding**

- **Fixed Sinusoidal Embeddings (NLP Transformer)**
 - ✓ Use sine and cosine functions with different frequencies.
 - ✓ Deterministic and not learned.
- **Learnable Embeddings (ViT default)**
 - ✓ Each position ($0 = [\text{CLS}]$, $1 \dots N = \text{patches}$) has its own **trainable vector**.
 - ✓ These vectors are initialized randomly and **optimized during training**.

- **Clarification on Learnable Embeddings**

- They do **not** encode positions by “simply numbering patches sequentially.”
- Instead, each position i has a dedicated **learnable vector** $e_i \in \mathbb{R}^D$.
- During training, the model **learns the best way to use position info** for image tasks.

Data Preparation for Vision Transformer Input

■ Step 4. Positional Embedding

- Why Learnable?

1. Images are 2D and far more complex than word order in NLP.
2. A fixed formula (like sinusoids) might be too rigid to capture subtle spatial patterns.
3. Learnable embeddings allow the model to adaptively learn spatial prior knowledge from the dataset.
→ Prior knowledge = Inductive bias
4. Empirically shown to perform better than fixed embeddings for vision.

Data Preparation for Vision Transformer Input

▪ Overall Processes – Data Preparation

Step 4. Adding Positional Embedding

Step 3. Adding Class Token

$$[x_{cls}; x_p^1E; x_p^2E; \dots; x_p^N E] \in \mathbb{R}^{(N+1) \times D}$$

* Learnable Class Token

Step 2-2. Patch Embedding: Linear Projection

$$[x_p^1E; x_p^2E; \dots; x_p^N E] \in \mathbb{R}^{N \times D}$$

Step 2-1. Patch Embedding: Flattening

Flattened patch

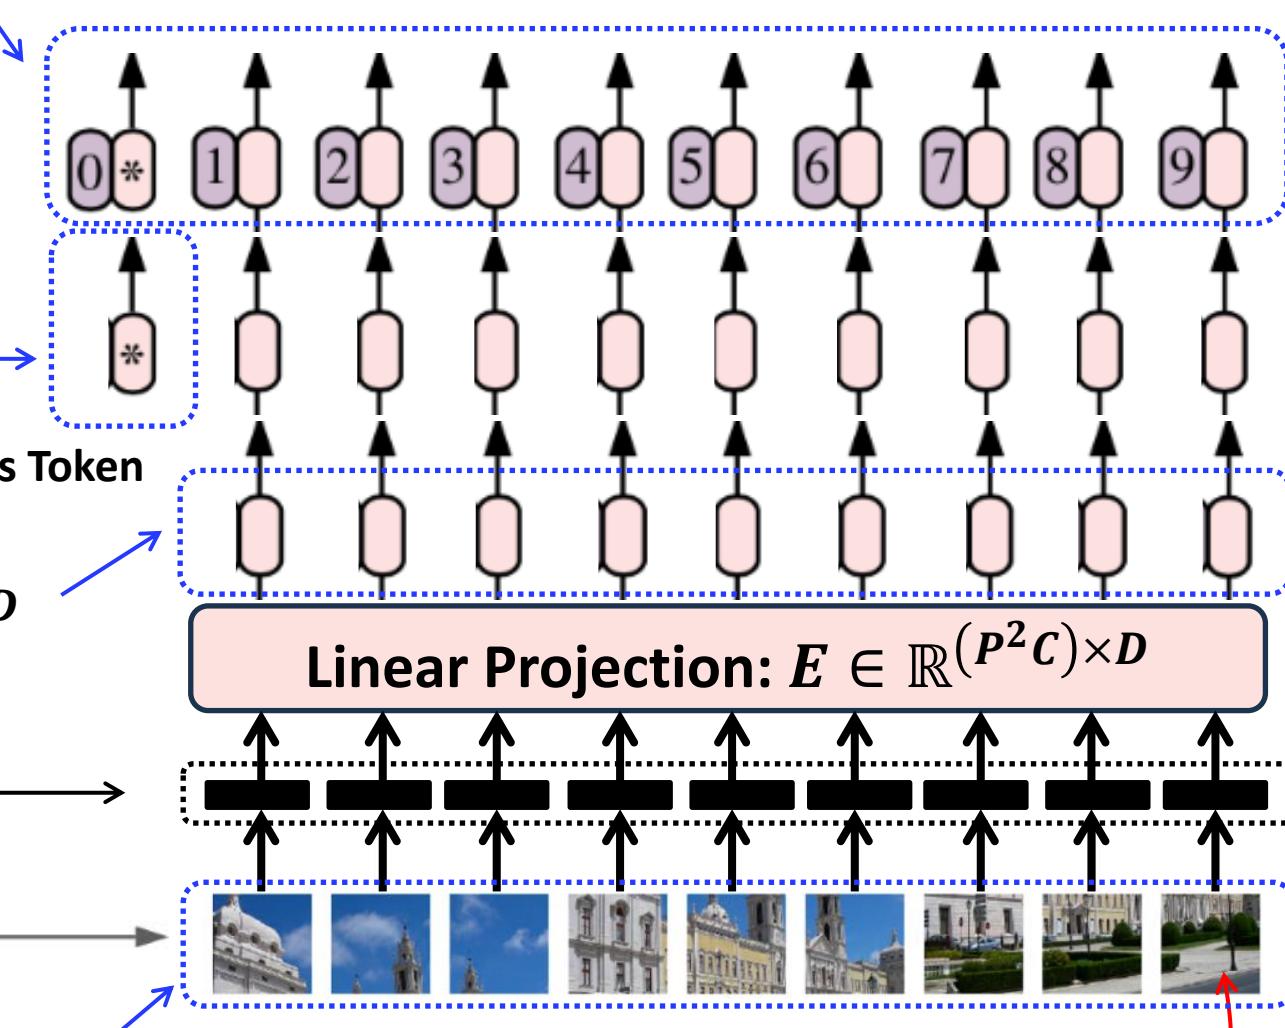
Step 1. Image Patch Creation



$$x_p \in \mathbb{R}^{N \times (P^2C)}; \text{ where } N = 9$$

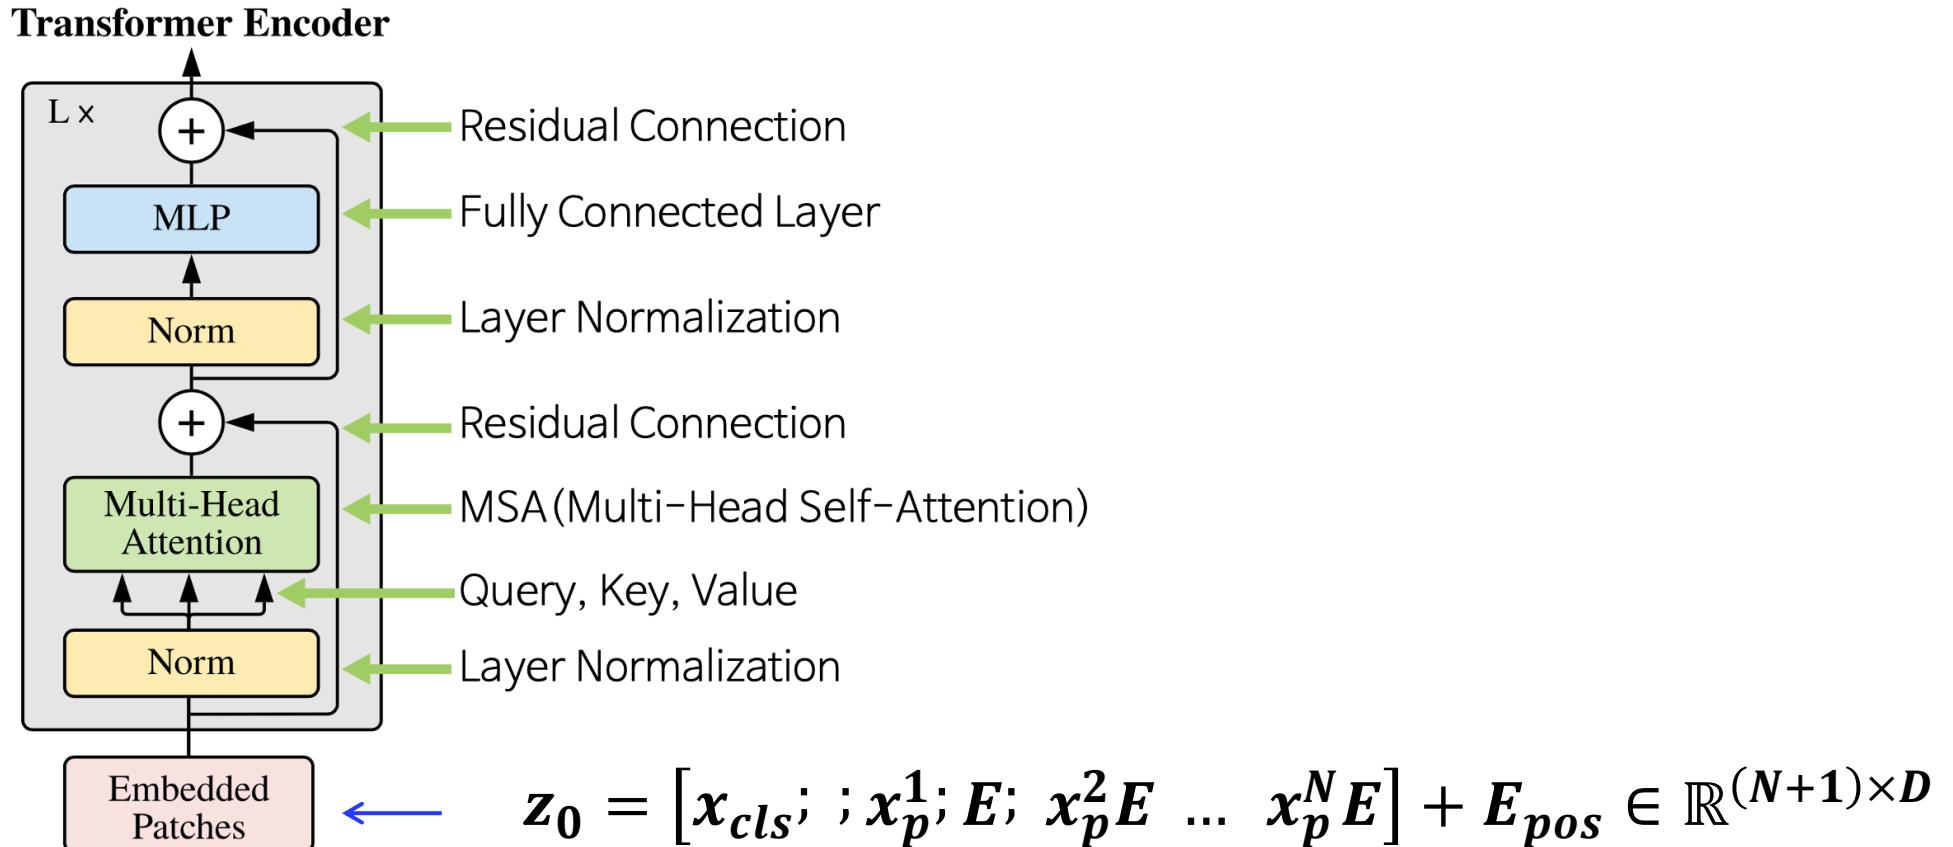
$$x_p^i \in \mathbb{R}^{P^2C}$$

$$z_0 = [x_{cls}; ; x_p^1E; x_p^2E \dots x_p^N E] + E_{pos} \in \mathbb{R}^{(N+1) \times D}$$



Vision Transformer Encoder Overview

- Transformer Encoder in Vision Transformer (ViT)
 - The Transformer Encoder in ViT is built by stacking **L identical blocks**.



- Each block refines the patch embeddings step by step with **attention and feed-forward processing**, while **residual connections** and **layer normalization** ensure stability.

Vision Transformer Encoder Overview

■ Transformer Encoder in Vision Transformer (ViT)

• Flow of Operations (Inside One Encoder Block)

- Input: Embedded Patches

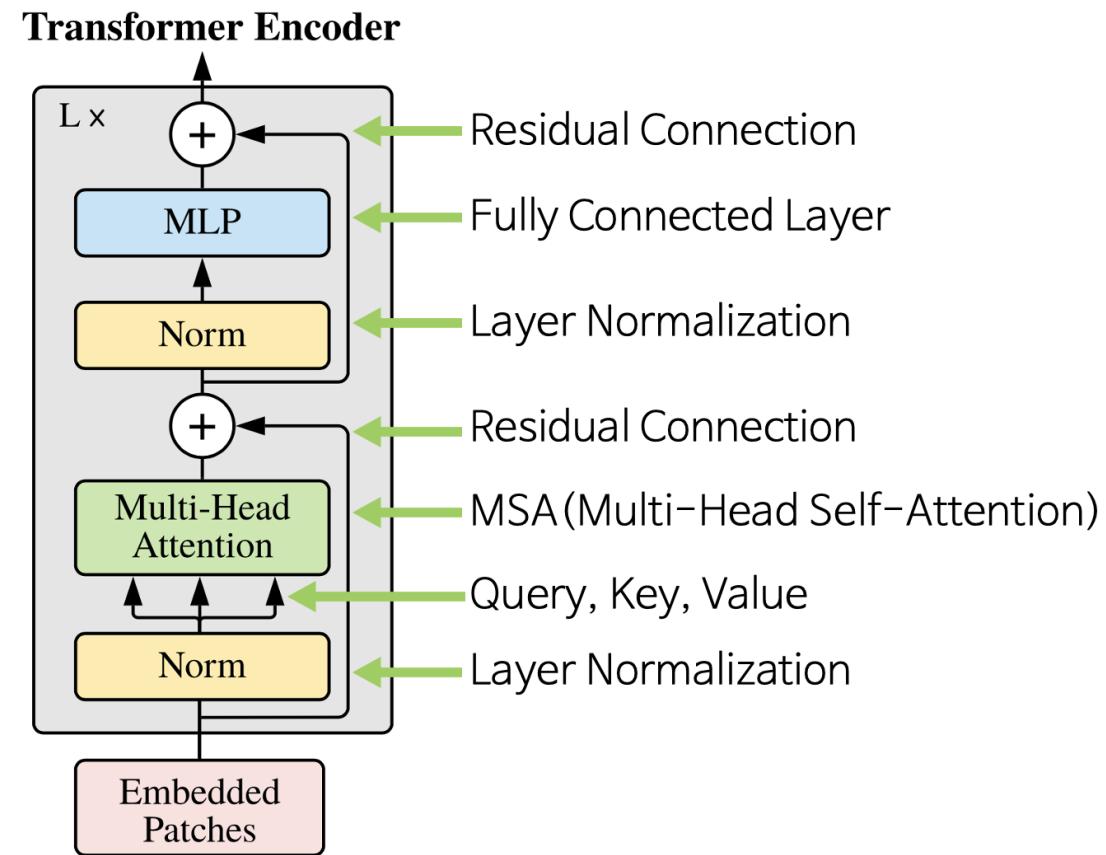
- ✓ These are the prepared tokens:
Patch Embeddings + Class Token + Positional Embeddings.

- Step 1. Layer Normalization (Norm)

- ✓ Normalizes inputs across the feature dimension.
 - ✓ Prepares stable input for Attention.

- Step 2. Q, K, V Projection → Multi-Head Self-Attention (MSA)

- ✓ Each token embedding is projected into **Query (Q)**, **Key (K)**, and **Value (V)**.
 - ✓ Self-Attention lets each token attend to **all others**, including the [CLS] token.
 - ✓ Multiple heads allow the model to capture **different types of relations** in parallel.



Vision Transformer Encoder Overview

■ Transformer Encoder in Vision Transformer (ViT)

• Flow of Operations (Inside One Encoder Block)

- Step 3. Residual Connection (+ Skip Connection)

- ✓ Adds the original input back to the attention output.
 - ✓ Ensures stable gradient flow and prevents loss of original information.

- Step 4. Second Layer Normalization (Norm again)

- ✓ Normalizes before the Feed-Forward stage.

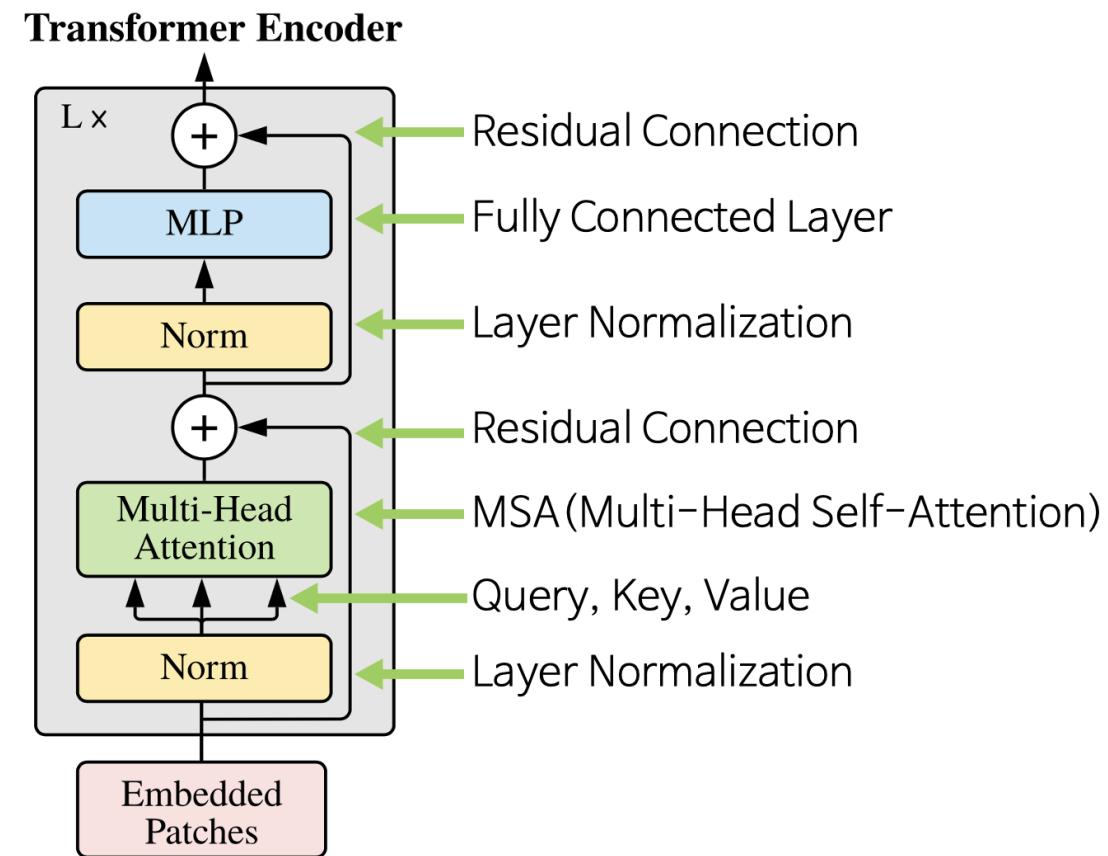
- Step 5. Feed-Forward Network (MLP Block)

- ✓ A two-layer fully connected network with GELU activation.

- ✓ Expands representation capacity by applying a non-linear transformation.

- Step 6. Residual Connection

- ✓ Again, adds the input of the MLP back to its output for stability.



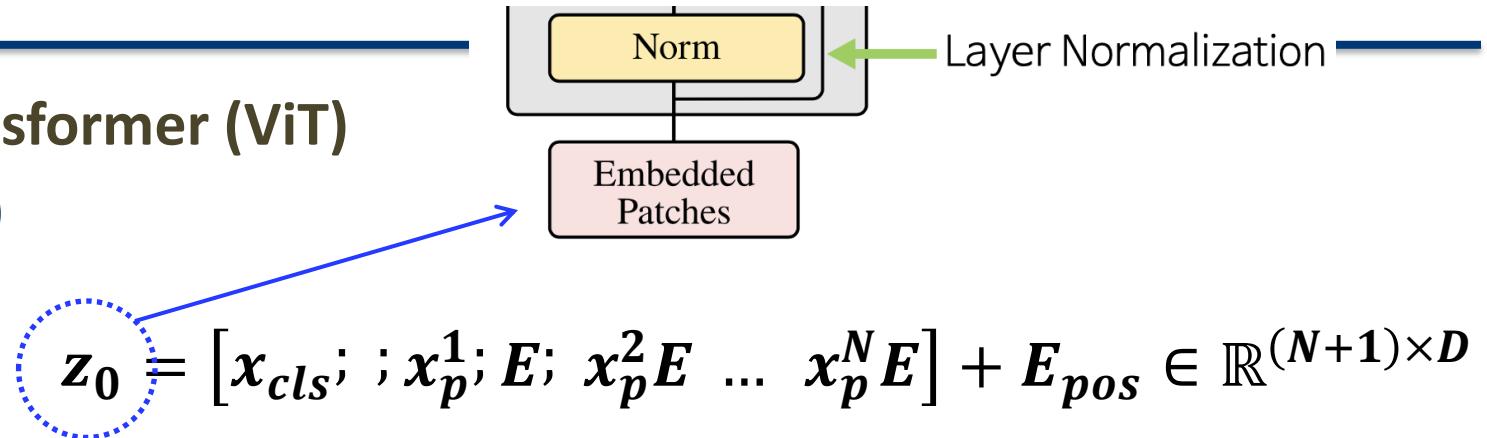
Vision Transformer Encoder Overview

■ Transformer Encoder in Vision Transformer (ViT)

• Step 1: Layer Normalization (Pre-Norm)

- Equation

$$LN(z_{l-1}), \quad l = 1 \dots L$$



- Role in ViT

- ✓ Ensures **stable training** before Multi-Head Self-Attention (Pre-Norm).
- ✓ Normalizes input features so that **scale differences** across tokens do not disturb attention.
- ✓ Prevents exploding or vanishing gradients when stacking many Transformer layers.

- Motivation

- ✓ Transformers stack dozens of layers → without normalization, training diverges quickly.
- ✓ Pre-Norm improves **gradient flow** and makes optimization easier compared to Post-Norm.

Key Points

- Layer Norm operates **independently on each token vector**.
- Keeps feature statistics balanced across dimensions.
- A crucial step to stabilize very deep architectures like ViT.

Vision Transformer Encoder Overview

■ Transformer Encoder in Vision Transformer (ViT)

- Step 1: Layer Normalization (Pre-Norm)

- What is Layer Normalization?

- ✓ Definition

- A normalization technique that standardizes **across the hidden dimensions** of each token, not across the batch.

- ✓ Equation

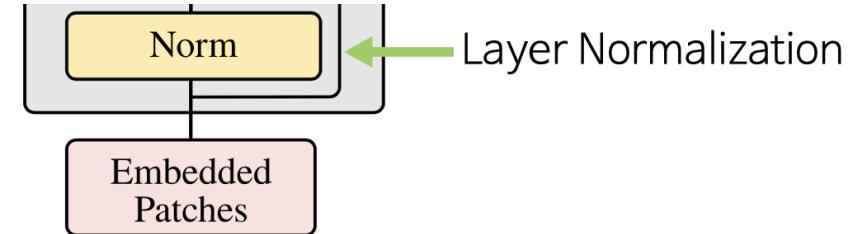
$$LN(x) = \frac{x - \mu}{\sigma} \cdot \gamma + \beta$$

- μ, σ : mean and standard deviation computed across the feature dimension.

- γ, β : learnable scale and shift parameters.

- ✓ Key Intuition

- Every token embedding is normalized **individually**, considering its feature vector.



Vision Transformer Encoder Overview

■ Transformer Encoder in Vision Transformer (ViT)

- Step 1: Layer Normalization (Pre-Norm)

- Layer Normalization vs Batch Normalization

- ✓ Batch Normalization (BN)

- Normalizes across the batch dimension.

- Computes mean and variance using all samples in the batch, per feature channel.

- Works well in CNNs, but **less stable in sequence models** (e.g., variable-length inputs) and **small batch-size**.

- ✓ Layer Normalization (LN)

- Normalizes within each sample across the hidden dimensions.

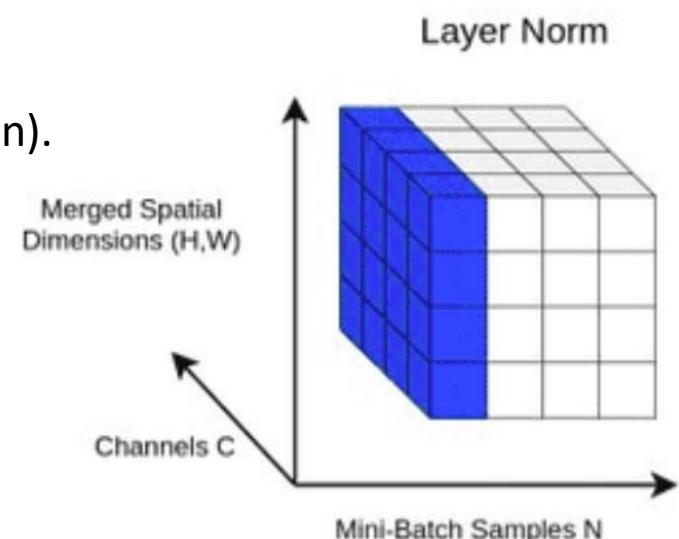
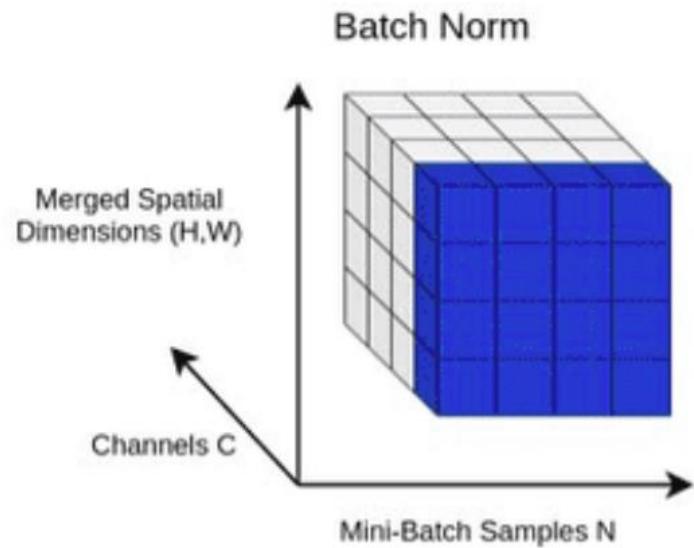
- Independent of batch size → works well for Transformers (NLP, Vision).

- Ensures consistent behavior even when **batch size is small** (important in ViT, where large images can limit batch size).

- ✓ Summary

- BN = normalize across batch (good for CNNs).

- LN = normalize across features (good for Transformers).



Vision Transformer Encoder Overview

■ Transformer Encoder in Vision Transformer (ViT)

- Step 2: Q, K, V Projection – How to make Q, K, and V

- Equation

- ✓ Input

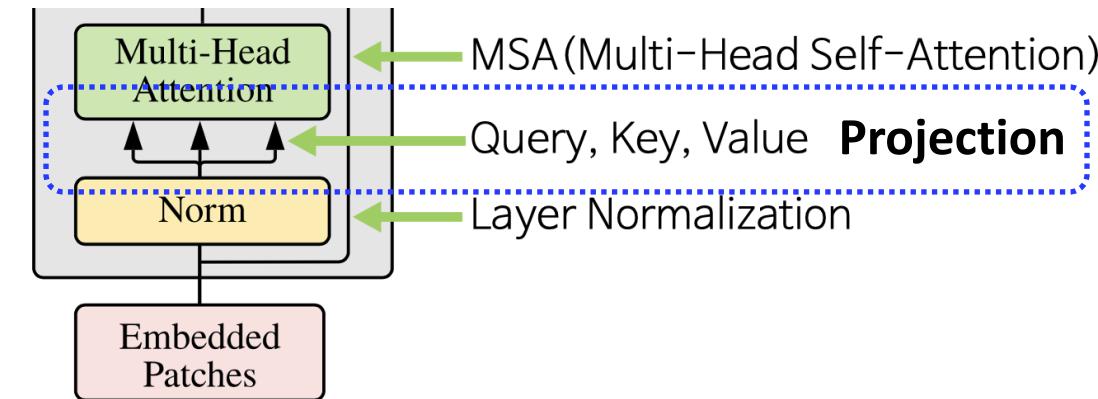
$$LN(Z) = \hat{Z} \in \mathbb{R}^{N \times D}$$

➤ where N = number of tokens (patches), D = embedding dimension.

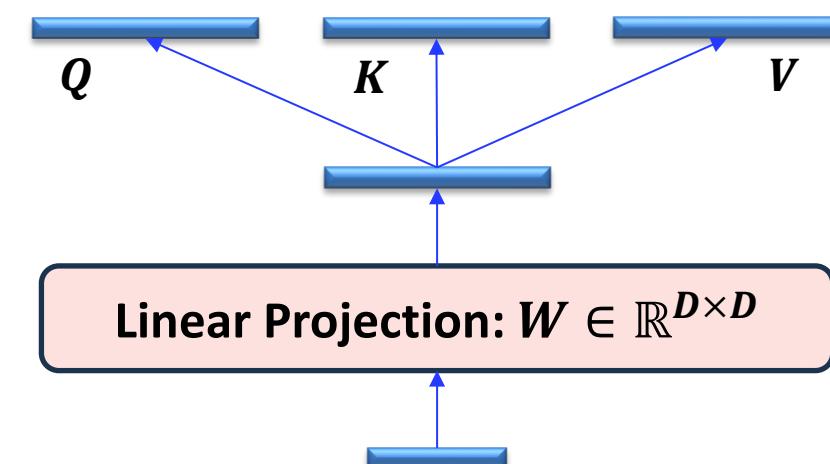
✓ 1. Projection

$$Q = \hat{Z}W_Q, \quad K = \hat{Z}W_K, \quad V = \hat{Z}W_V; \quad Q, K, V \in \mathbb{R}^{N \times D}$$

➤ where $W_Q, W_K, W_V \in \mathbb{R}^{D \times D}$ are learnable projection matrices.



$$Q, K, V \in \mathbb{R}^{N \times D}$$



$$LN(Z) = \hat{Z} \in \mathbb{R}^{N \times D}$$

Vision Transformer Encoder Overview

- Transformer Encoder in Vision Transformer (ViT)

- Step 2: Q, K, V Projection – How to make Q, K, and V

- Equation

- ✓ 2. Head-wise Split – For multi-head attention with h heads

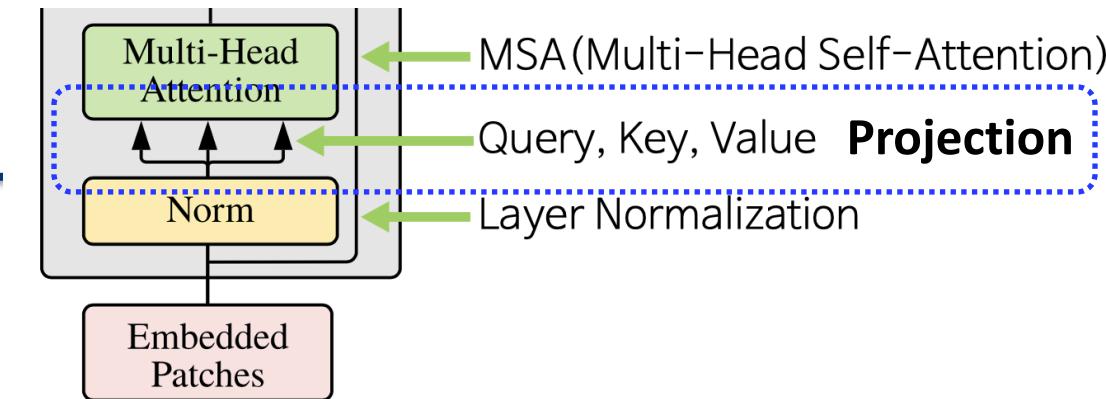
$$D_h = \frac{D}{h}$$

$$Q_i, K_i, V_i \in \mathbb{R}^{N \times D_h}; \quad \text{where } i = 1, \dots, h$$

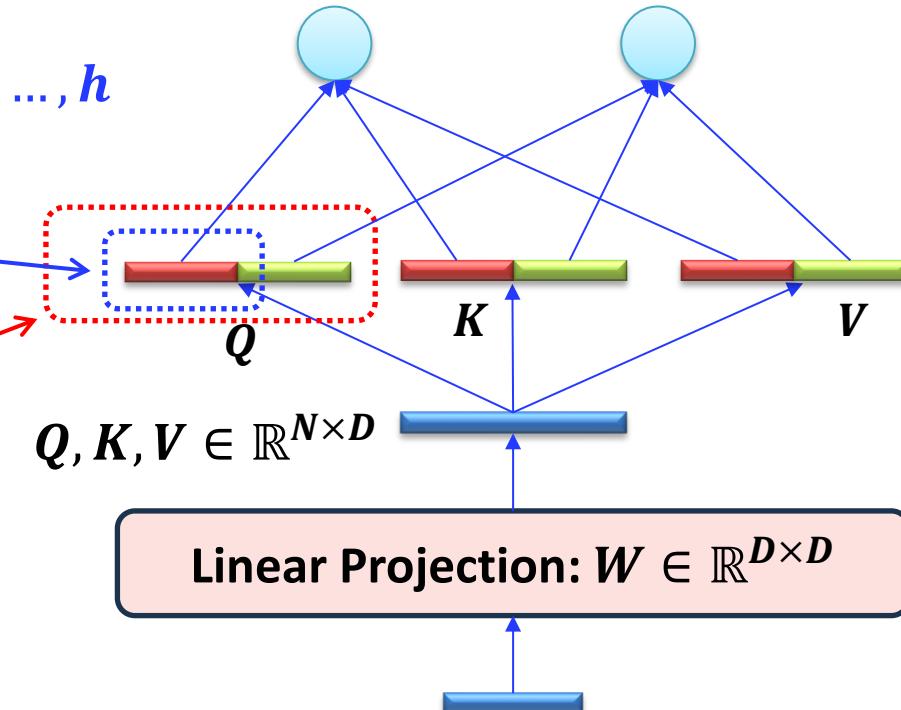
► Each projection is reshaped as

$$Q, K, V \in \mathbb{R}^{N \times D} \rightarrow [h \ N \ D_h]$$

$$Q, K, V \in \mathbb{R}^{h \times N \times D_h}$$



if $h = 2$



$$LN(Z) = \hat{Z} \in \mathbb{R}^{N \times D}$$

Vision Transformer Encoder Overview

- Transformer Encoder in Vision Transformer (ViT)

- Step 2: Multi-Head Self-Attention (MSA)

- Equation

- ✓ 3. Per-head attention – For each head i

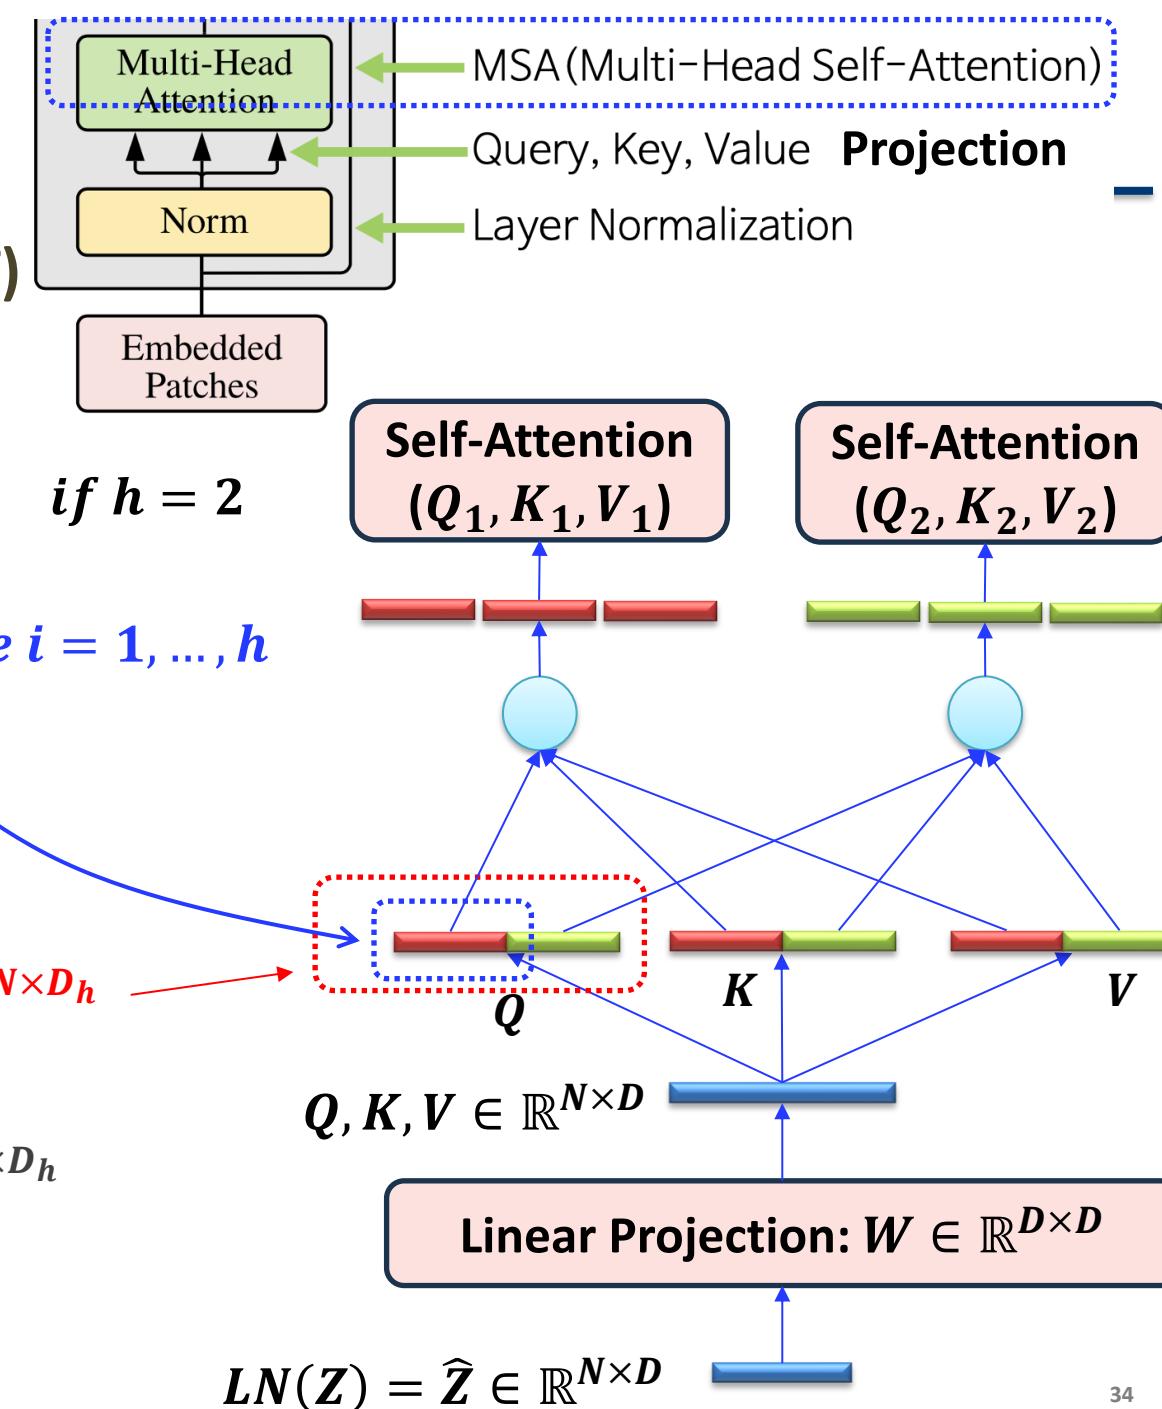
$$Q_i, K_i, V_i \in \mathbb{R}^{N \times D_h}; \quad \text{where } i = 1, \dots, h$$

$$Q_i, K_i, V_i \in \mathbb{R}^{N \times D_h}$$

$$A_i = \text{softmax}\left(\frac{Q_i K_i^\top}{\sqrt{D_h}}\right)$$

$$O_i = \text{Self-Attention}(Q_i, K_i, V_i) = A_i V_i \in \mathbb{R}^{N \times D_h}$$

where N : #tokens, D : embedding dim



Vision Transformer Encoder Overview

- Transformer Encoder in Vision Transformer (ViT)

- Step 2: Multi-Head Self-Attention (MSA)

- Equation

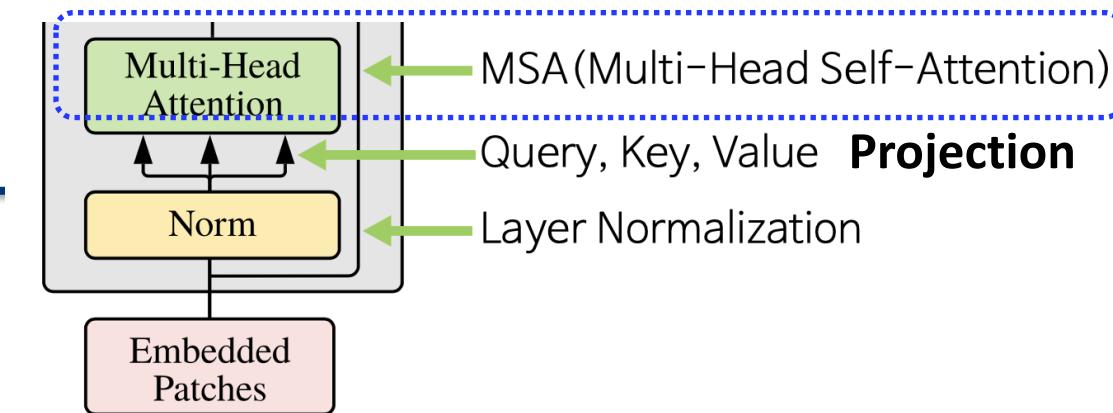
- ✓ 4. Concatenate heads (channel-wise)

$$\begin{aligned} \mathbf{o}_i &= \text{Self-Attention}(Q_i, K_i, V_i) \\ &= A_i V_i \in \mathbb{R}^{N \times D_h} \end{aligned}$$

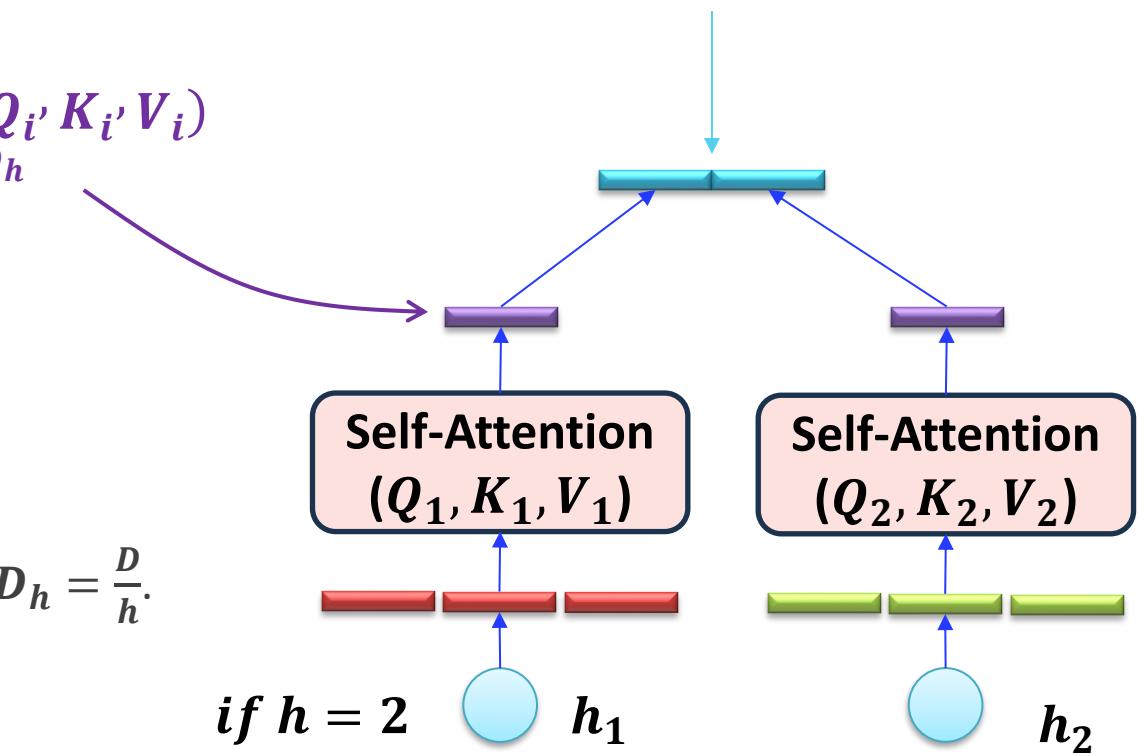
$$\mathbf{o} = \text{Concat}(\mathbf{o}_1, \dots, \mathbf{o}_h) \in \mathbb{R}^{N \times (hD_h)} = \mathbb{R}^{N \times D}$$

➤ Each head output is $[N, D_h]$;

after concat you get $[N, h \cdot D_h] = [N, D]$; where $D_h = \frac{D}{h}$.



$$\mathbf{o} = \text{Concat}(\mathbf{o}_1, \dots, \mathbf{o}_h) \in \mathbb{R}^{N \times (hD_h)} = \mathbb{R}^{N \times D}$$



Vision Transformer Encoder Overview

■ Transformer Encoder in Vision Transformer (ViT)

- Step 2: Multi-Head Self-Attention (MSA)

- Equation

- ✓ 5. Output projection (mix heads)

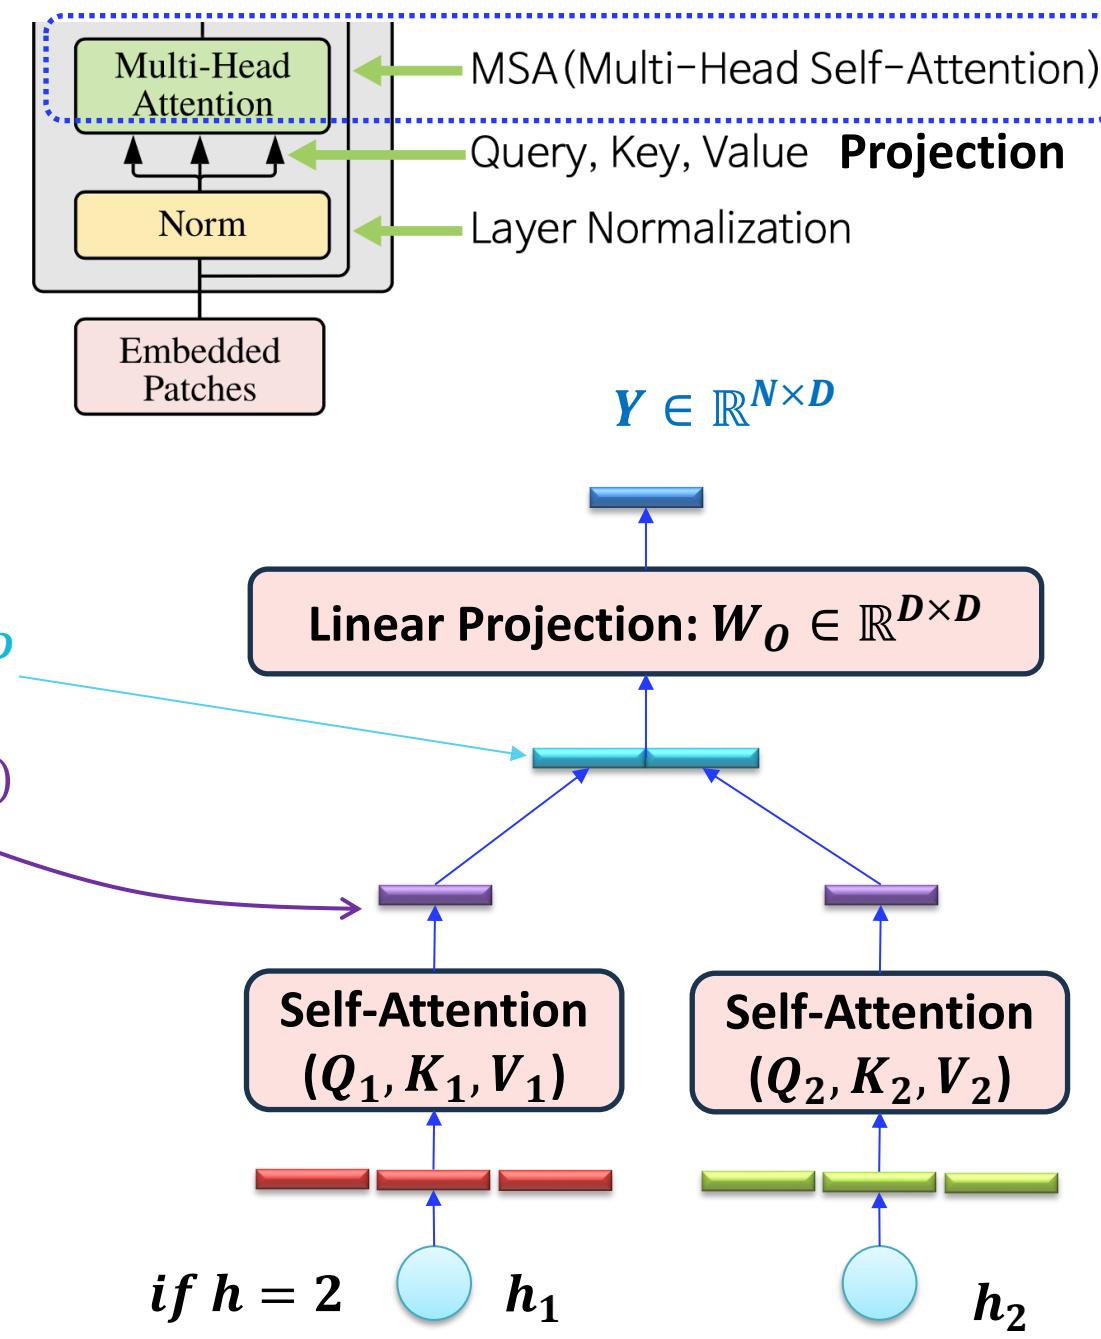
$$\mathbf{O} = \text{Concat}(\mathbf{O}_1, \dots, \mathbf{O}_h) \in \mathbb{R}^{N \times (hD_h)} = \mathbb{R}^{N \times D}$$

$$\begin{aligned}\mathbf{O}_i &= \text{Self-Attention}(\mathbf{Q}_i, \mathbf{K}_i, \mathbf{V}_i) \\ &= \mathbf{A}_i \mathbf{V}_i \in \mathbb{R}^{N \times D_h}\end{aligned}$$

$$Y = \mathbf{O} W_O, \quad W_O \in \mathbb{R}^{D \times D}, \quad Y \in \mathbb{R}^{N \times D}$$

➤ Purpose: **couple the heads** (a learned linear mix) and return to the model width D .

➤ Without W_O , heads would remain independent concatenations.



Vision Transformer Encoder Overview

■ Transformer Encoder in Vision Transformer (ViT)

• Step 2: Multi-Head Self-Attention (MSA)

- Why Do We Project Twice in Attention?

✓ Part 1 — Q, K, V Projection

➤ Why? To create three different views of the same embedding.

$$Q = ZW_Q, \quad K = ZW_K, \quad V = ZW_V$$

- **Q (Query)**: what this token is “asking for”
- **K (Key)**: how other tokens “describe themselves”
- **V (Value)**: the actual **content information**

➤ Reason

- If Q, K, V were the same, **attention would collapse**.
(NO distinction between asking, matching, and answering).
- Projection lets the model learn **different subspaces** for “matching” and “carrying info.”
- Each head then focuses on **different aspects** (local edges, global shape, texture, etc.).

Vision Transformer Encoder Overview

- Transformer Encoder in Vision Transformer (ViT)

- Step 2: Multi-Head Self-Attention (MSA)

- Why Do We Project Twice in Attention?

- ✓ Part 2 — Output Projection (After Concatenation)

- After per-head attention: Each head produces $O_i \in \mathbb{R}^{N \times D_h}$.

$$\text{Concatenation} \rightarrow O = \text{Concat}(O_1, \dots, O_h) \in \mathbb{R}^{N \times (hD_h)} = \mathbb{R}^{N \times D}$$

- Why another projection?

$$Y = OW_O, \quad W_O \in \mathbb{R}^{D \times D}$$

- To mix information across heads (otherwise heads remain independent).
- To bring the concatenated result **back into the model's fixed embedding dimension D** .
- Ensures the multi-head mechanism acts as diverse views + integration, not just parallel streams.

Vision Transformer Encoder Overview

- Transformer Encoder in Vision Transformer (ViT)

- Step 2: Multi-Head Self-Attention (MSA)

- Why Do We Project Twice in Attention?

- ✓ Key Points

- Projection (1st time) → separates roles into Q, K, V.

- Multi-head splitting → multiple independent subspaces, each learning a different relation.

- Projection (2nd time) → mixes those subspaces back together, so the model uses them jointly.

- Think of it as: asking the same question in many ways, then combining all the answers into one clear representation.

Vision Transformer Encoder Overview

■ Transformer Encoder in Vision Transformer (ViT)

- Step 3: Residual Connection (after MSA)

- Equation

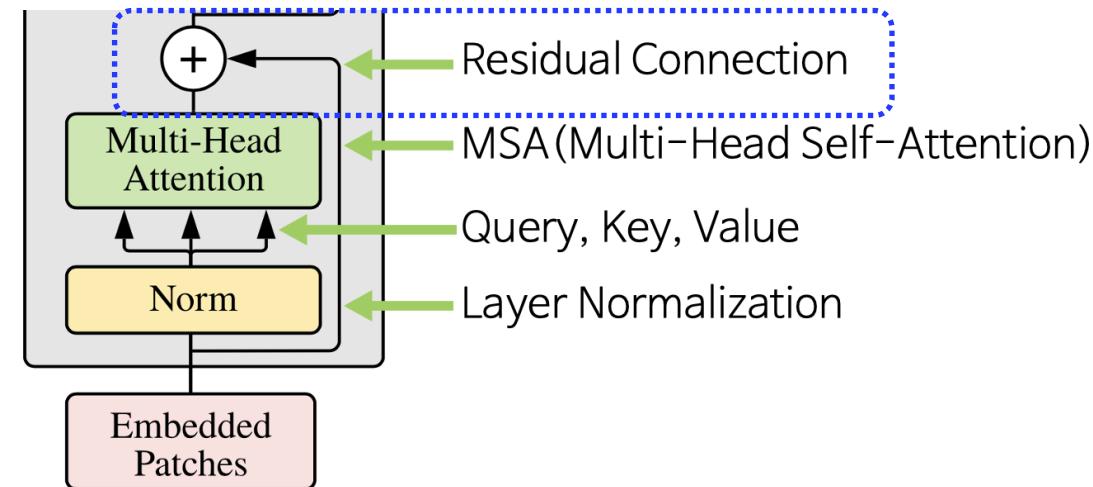
$$z'_\ell = MSA(LN(z_{\ell-1})) + z_{\ell-1}, \quad \ell = 1 \dots L$$

- Role & Purpose

- ✓ Add **shortcut connection** to stabilize training.
 - ✓ Prevent vanishing gradients by allowing gradient flow directly.
 - ✓ Ensure **information preservation**: even if MSA fails, the model can keep the original representation.

- Why Important in ViT?

- ✓ Vision tasks involve **long sequences (many patches)**, making deep training unstable.
 - ✓ Residual paths act as an “**information highway**”, supporting both learning new features and keeping old ones.



Vision Transformer Encoder Overview

■ Transformer Encoder in Vision Transformer (ViT)

- Step 4: Second Layer Normalization

- Equation

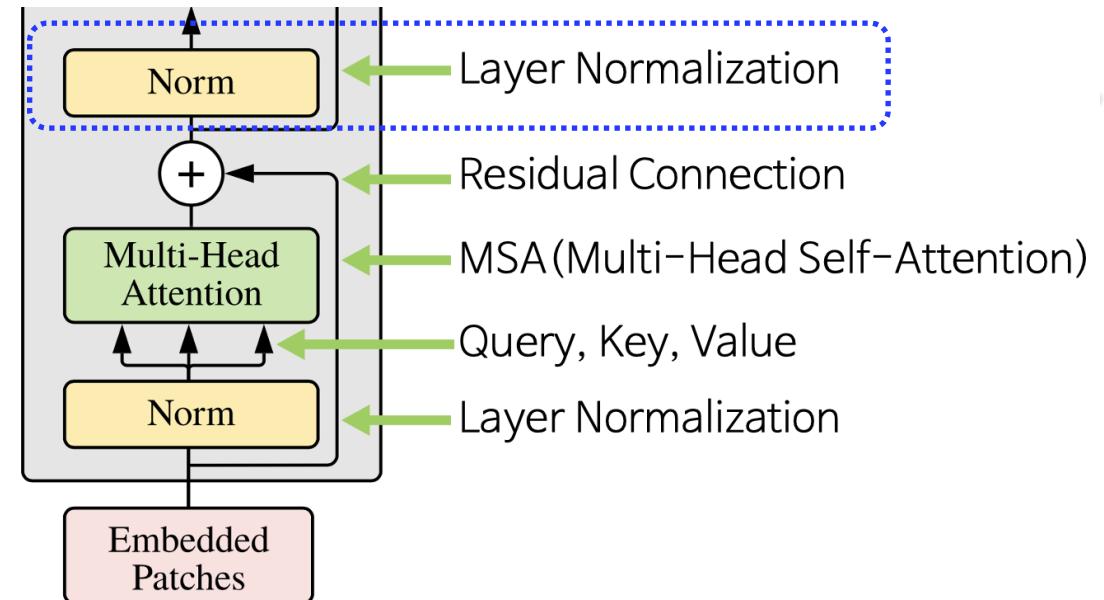
$$\hat{\mathbf{z}}'_\ell = LN(\mathbf{z}'_\ell), \quad \ell = 1 \dots L$$

- Why Another LN?

- ✓ After residual addition, activations may drift (scale/shift).
 - ✓ LN re-centers and re-scales token embeddings before feeding into the **Feed-Forward Network (MLP Block)**.
 - ✓ Stabilizes distribution → improves convergence.

- Key Point

- ✓ LN is applied **before every major sub-block** (Pre-Norm design).
 - ✓ Without this, deep ViTs would suffer from instability.



Vision Transformer Encoder Overview

■ Transformer Encoder in Vision Transformer (ViT)

• Step 5: Feed-Forward Network (MLP Block)

- Equation

$$\mathbf{z}_\ell = \mathbf{MLP}(\mathbf{LN}(\mathbf{z}'_\ell)) = \mathbf{MLP}(\hat{\mathbf{z}}'_\ell), \quad \ell = 1 \dots L$$

- ✓ Where

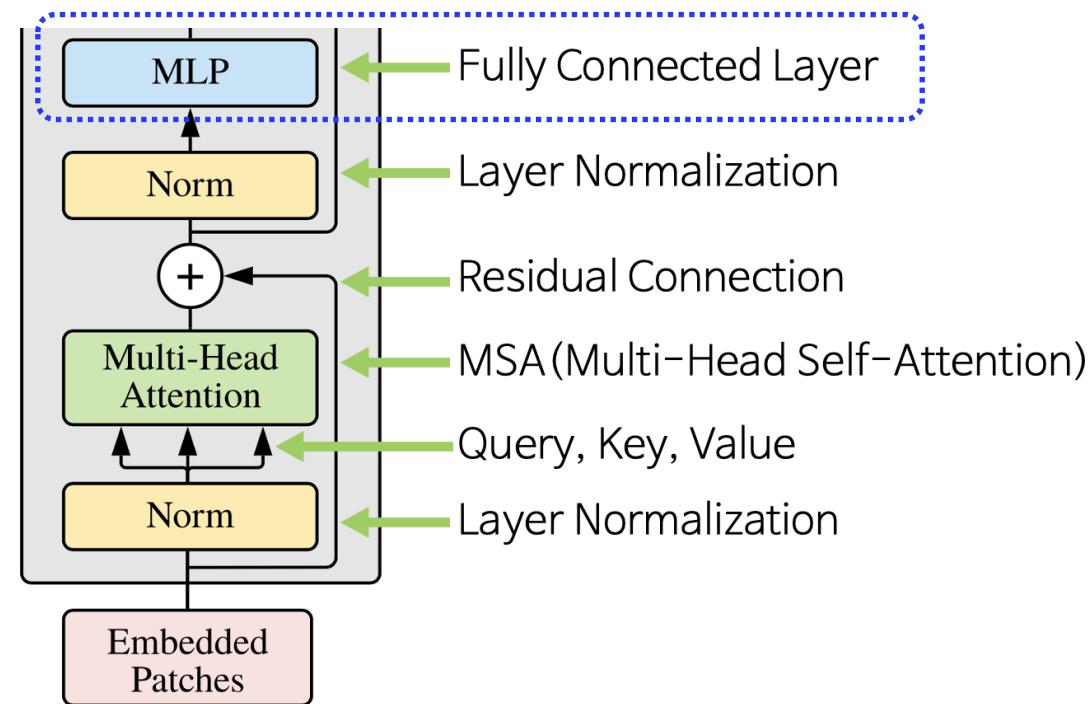
$$\mathbf{MLP}(x) = W_2 \cdot \mathbf{GELU}(W_1 x + b_1) + b_2$$

- Role

- ✓ Token-wise transformation: each patch embedding processed independently.
- ✓ Expands embedding dimension (e.g., $D \rightarrow 4D \rightarrow D$) to capture richer features.

- Why Fully-Connected Layers (Not CNN)?

- ✓ Unlike CNN, ViT has **no spatial inductive bias** — MLP just processes each token vector.
- ✓ Efficient for token-wise non-linear transformation.



Why GELU (Gaussian Error Linear Unit)?

- Captures **small variations** effectively (important for vision tasks).
- Empirically improves convergence and accuracy in Transformer models.

Vision Transformer Encoder Overview

■ Transformer Encoder in Vision Transformer (ViT)

- Step 6: Residual Connection (after MLP)

- Equation

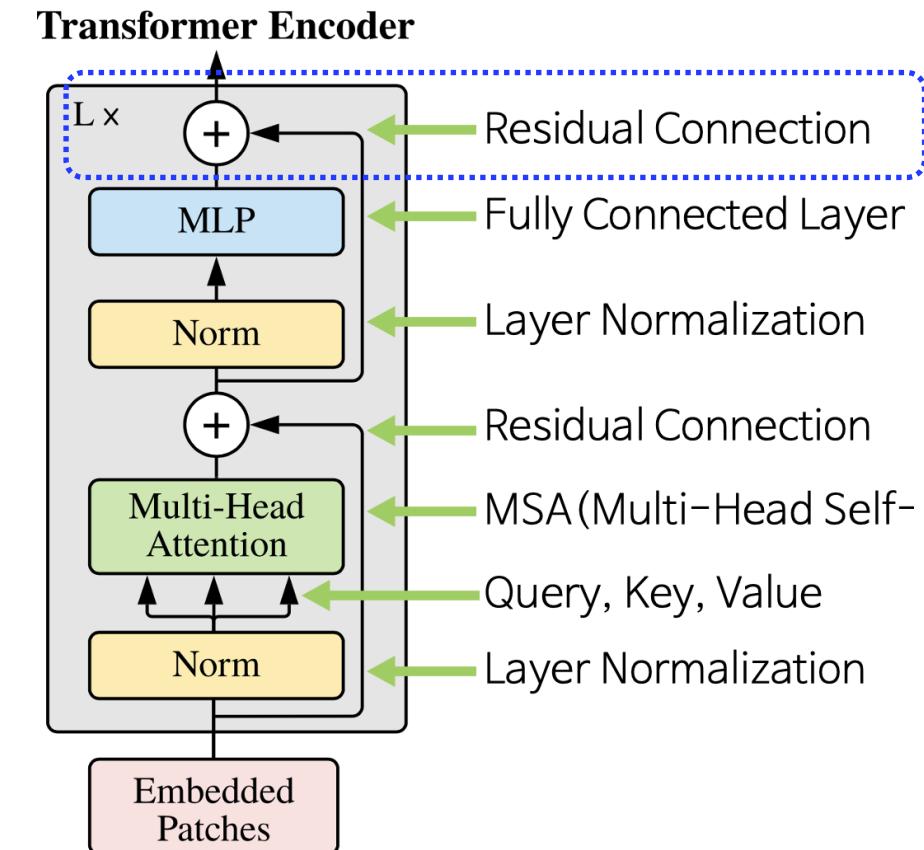
$$\mathbf{z}_\ell = \mathbf{MLP}(\mathbf{LN}(\mathbf{z}'_\ell)) + \mathbf{z}'_\ell = \mathbf{MLP}(\hat{\mathbf{z}}'_\ell) + \mathbf{z}'_\ell, \quad \ell = 1 \dots L$$

- Why Another Residual?

- ✓ Same motivation: stabilize deep training, prevent information loss.
 - ✓ After the **MLP transformation**, residual ensures the model can fallback on the input representation.
 - ✓ Enables stacking many Transformer layers (L) without degradation.

- Key Point

- ✓ Every Transformer block = **(LN → Sub-layer → Residual) × 2**.
 - First sub-layer = MSA
 - Second sub-layer = MLP



ViT Output and Classification Head

- Only the [CLS] token is used for final prediction

- Output of Vision Transformer (ViT)

- Equation

$$y = LN(z_L^0)$$

- ✓ z_L^0 : the [CLS] token representation after passing through L encoder layers.

- ✓ LN : final Layer Normalization applied to stabilize the representation.

- ✓ y : serves as the input to the MLP Head for classification.

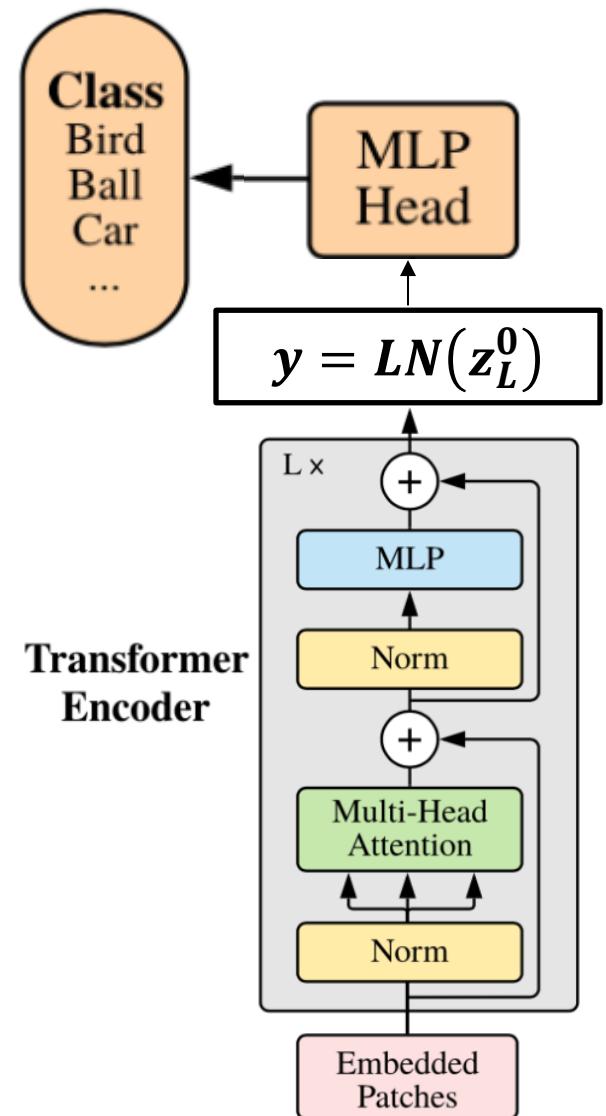
- Classification Head

- ✓ After the Transformer Encoder, **ONLY the [CLS] token is used.**

- ✓ This token has aggregated **global information** from all patches through multiple self-attention layers.

- ✓ y is passed through an MLP Head

(usually one or two fully connected layers with softmax at the end).



ViT Output and Classification Head

- Only the [CLS] token is used for final prediction

- Where does z_L^0 from exactly?

- Each Transformer block ℓ has two sub-layers

- ✓ 1st sub-layer: Multi-Head Self-Attention (MSA) + Residual

- ✓ 2nd sub-layer: MLP (Feed-Forward Network) + Residual

- z_L^0

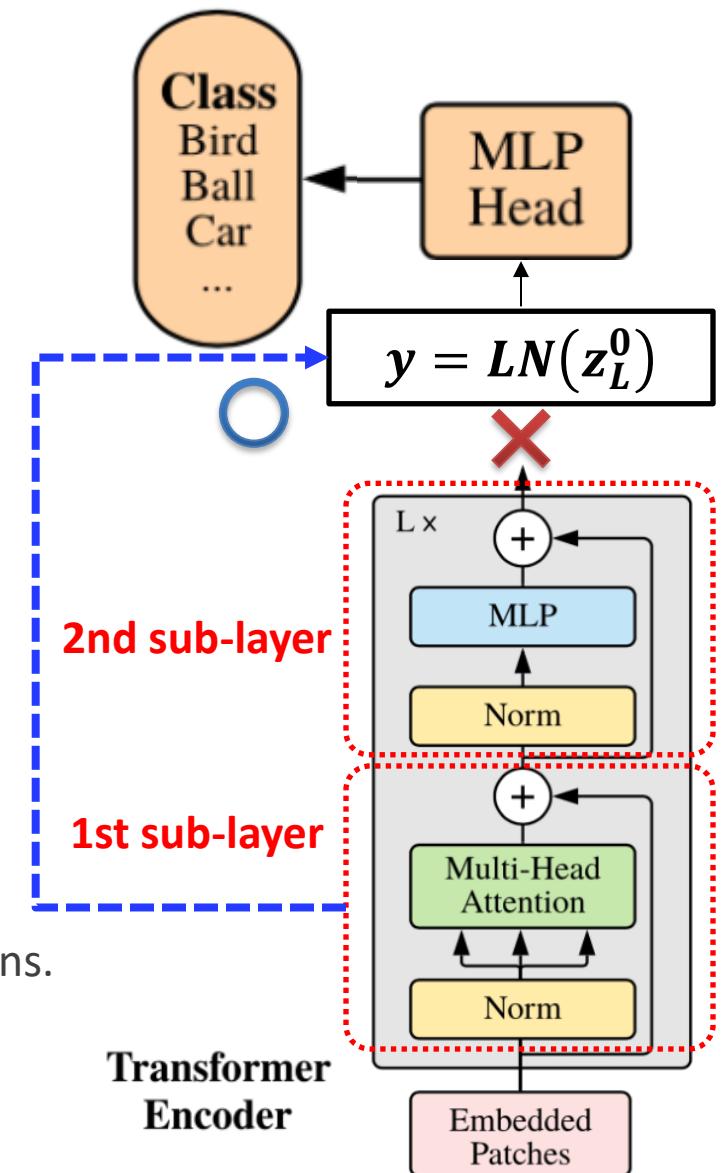
- ✓ The intermediate representation after the first sub-layer (MSA + Residual) of the last Transformer block (L).

- Why Only the [CLS] Token?

- It acts as a **summary representation** of the entire image.

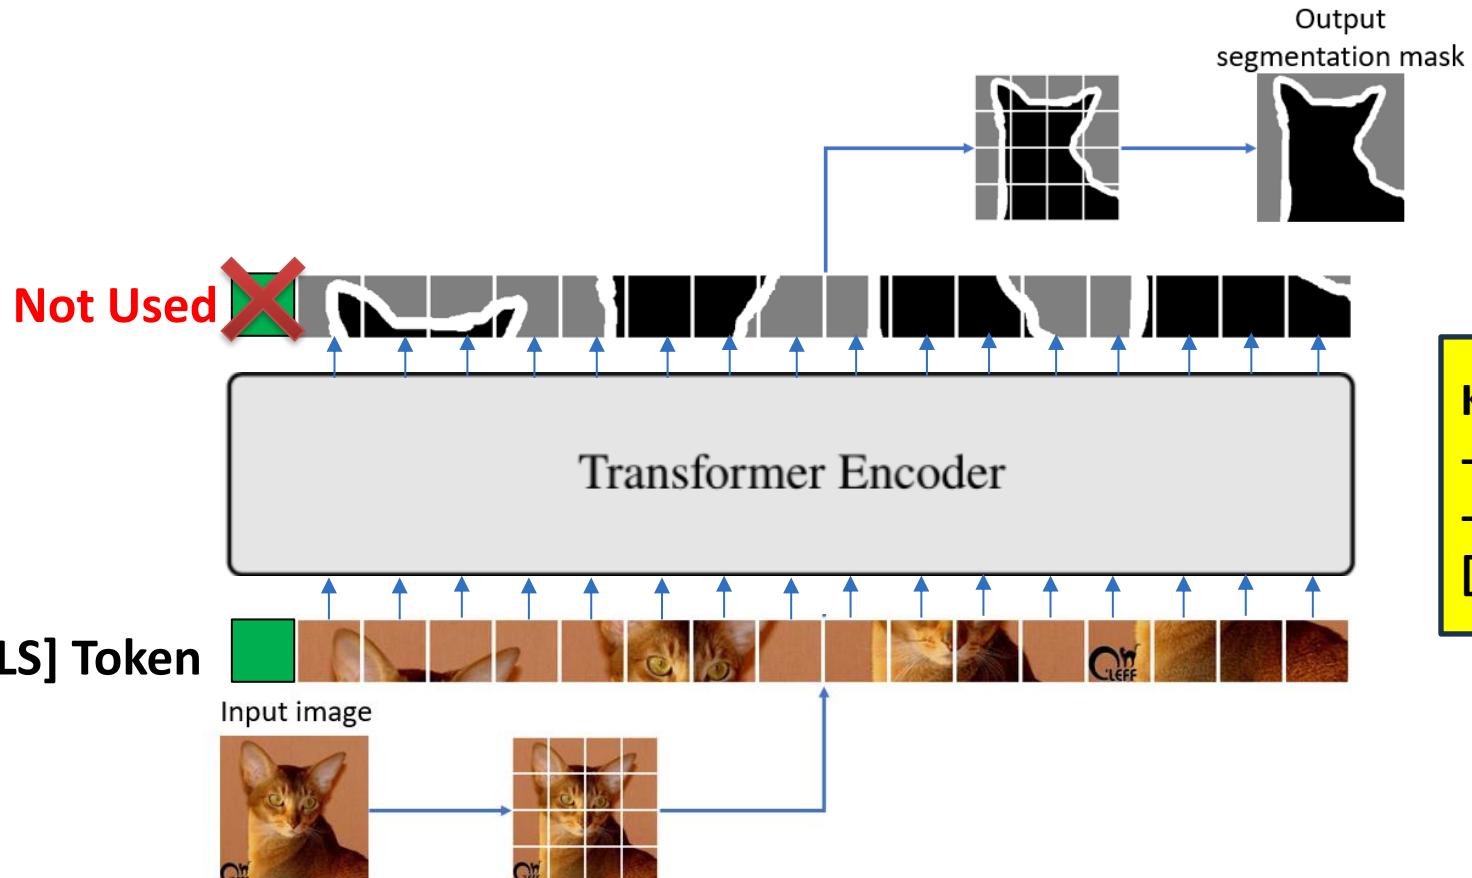
- Attention layers allow the [CLS] token to gather information from all other tokens.

- This makes it sufficient for **classification tasks**, while ignoring other tokens.



Beyond Classification: Other Tasks with ViT

- Segmentation and Detection need spatial outputs, not just [CLS]
 - In Classification, we only use the [CLS] token, because it already summarizes global information.



Key Points

- ViT is **NOT limited** to classification.
- Depending on the task, we either use **[CLS] (global)** or **all patches (spatial)**.

- But for other vision tasks like **Segmentation or Detection**, we need **spatial details**.
 - Instead of using only the [CLS] token, ViT uses **all patch embeddings**.

Summary: Vision Transformer (ViT)

■ Key Takeaways

- **From CNNs to Transformers**

- CNNs: strong inductive bias (locality, translation equivariance).
- ViT: weak inductive bias, relies on **data-driven global self-attention**.

- **Data Preparation**

- Split image into patches → Flatten → Linear Projection.
- Add [CLS] token (global representation).
- Add positional embeddings (spatial order).

- **Transformer Encoder**

- Repeated L blocks
 - ✓ LN → QKV Projection → Multi-Head Self-Attention → Residual.
 - ✓ LN → MLP (with GELU) → Residual.
- Residual & LN ensure stability in deep models.

Summary: Vision Transformer (ViT)

■ Key Takeaways

- **Output for Classification**

- Only the **[CLS] token** is used after the final block.
 - Passed through MLP Head → class probabilities.

- **Beyond Classification**

- For tasks like **Segmentation, Detection**
 - ✓ Use **all patch embeddings** instead of only [CLS].
 - ✓ Enables pixel-level or region-level predictions.